

(Ab)using compiler tools

(Ab)using compiler tools

Two binary analysis tools
Only one of them is strictly connected to compilers
No abuse, no fun

Bloaty McBloatface

- Josh Haberman, 2016
- <https://github.com/google/bloaty>
- <http://blog.reverberate.org/>
- What is taking up space in your binary?
- Supports ELF and Mach-O, PE is desired

```
1 #include <stdio.h>
2
3 int zero_init[10];
4
5 void do_stuff(int arg) {
6     int local = arg + 2;
7     int i;
8
9     for (i = 0; i < local; ++i) {
10         printf("i = %d\n", i);
11     }
12 }
13
14 int main() {
15     do_stuff(2);
16     return 0;
17 }
```

```
reka@ubuntu:~/binary-tools$ bloaty ./dummy
```

FILE SIZE		VM SIZE		
-----		-----		
16.1%	1.73Ki	0.0%	0	[Unmapped]
15.5%	1.66Ki	0.0%	0	.symtab
11.6%	1.24Ki	16.0%	396	[17 Others]
8.8%	964	0.0%	0	.debug_info
6.7%	737	0.0%	0	.debug_str
5.3%	579	23.4%	579	[LOAD #2 [RX]]
4.9%	542	0.0%	0	.strtab
4.8%	528	18.8%	464	.dynamic
4.5%	498	17.6%	434	.text
3.5%	387	0.0%	0	.shstrtab
3.4%	374	0.0%	0	.debug_abbrev
3.2%	352	11.7%	288	.eh_frame
2.7%	296	0.0%	0	.debug_line
1.5%	160	3.9%	96	.dynsym
1.2%	132	2.8%	68	.eh_frame_hdr
1.2%	128	0.0%	0	[ELF Headers]
1.2%	127	2.5%	63	.dynstr
1.0%	112	0.0%	0	.debug_aranges
1.0%	112	1.9%	48	.rela.dyn
1.0%	107	0.0%	0	.comment
0.9%	100	1.5%	36	.note.gnu.build-id
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

```
reka@ubuntu:~/binary-tools$ bloaty ./dummy
```

FILE SIZE		VM SIZE		
-----	-----	-----	-----	
16.1%	1.73Ki	0.0%	0	[Unmapped]
15.5%	1.66Ki	0.0%	0	.symtab
11.6%	1.24Ki	16.0%	396	[17 Others]
8.8%	964	0.0%	0	.debug_info
6.7%	737	0.0%	0	.debug_str
5.3%	579	23.4%	579	[LOAD #2 [RX]]
4.9%	542	0.0%	0	.strtab
4.8%	528	18.8%	464	.dynamic
4.5%	498	17.6%	434	.text
3.5%	387	0.0%	0	.shstrtab
3.4%	374	0.0%	0	.debug_abbrev
3.2%	352	11.7%	288	.eh_frame
2.7%	296	0.0%	0	.debug_line
1.5%	160	3.9%	96	.dynsym
1.2%	132	2.8%	68	.eh_frame_hdr
1.2%	128	0.0%	0	[ELF Headers]
1.2%	127	2.5%	63	.dynstr
1.0%	112	0.0%	0	.debug_aranges
1.0%	112	1.9%	48	.rela.dyn
1.0%	107	0.0%	0	.comment
0.9%	100	1.5%	36	.note.gnu.build-id
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

```
reka@ubuntu:~/binary-tools$ bloaty -n 0 ./dummy
```

FILE SIZE		VM SIZE		
16.1%	1.73Ki	0.0%	0	[Unmapped]
15.5%	1.66Ki	0.0%	0	.symtab
8.8%	964	0.0%	0	.debug_info
6.7%	737	0.0%	0	.debug_str
5.3%	579	23.4%	579	[LOAD #2 [RX]]
4.9%	542	0.0%	0	.strtab
4.8%	528	18.8%	464	.dynamic
4.5%	498	17.6%	434	.text
3.5%	387	0.0%	0	.shstrtab
3.4%	374	0.0%	0	.debug_abbrev
3.2%	352	11.7%	288	.eh_frame
2.7%	296	0.0%	0	.debug_line
1.5%	160	3.9%	96	.dynsym
1.2%	132	2.8%	68	.eh_frame_hdr
1.2%	128	0.0%	0	[ELF Headers]
1.2%	127	2.5%	63	.dynstr
1.0%	112	0.0%	0	.debug_aranges
1.0%	112	1.9%	48	.rela.dyn
1.0%	107	0.0%	0	.comment
0.9%	100	1.5%	36	.note.gnu.build-id
0.9%	96	1.3%	32	.gnu.version_r
0.9%	96	1.3%	32	.got.plt
0.9%	96	1.3%	32	.note.ABI-tag
0.9%	96	1.3%	32	.plt
0.8%	92	1.1%	28	.gnu.hash
0.8%	92	1.1%	28	.interp
0.8%	88	1.0%	24	.rela.plt
0.8%	87	0.9%	23	.init
0.7%	80	0.6%	16	.data
0.7%	80	0.6%	16	.got
0.7%	76	0.5%	12	.rodata
0.7%	73	0.4%	9	.fini
0.0%	0	2.9%	72	.bss
0.7%	72	0.3%	8	.fini_array
0.7%	72	0.3%	8	.gnu.version
0.7%	72	0.3%	8	.init_array
0.0%	0	0.6%	16	[LOAD #3 [RW]]
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

File size

Bytes on disk

```
reka@ubuntu:~/binary-tools$ bloaty -n 0 ./dummy
```

FILE SIZE		VM SIZE		
16.1%	1.73Ki	0.0%	0	[Unmapped]
15.5%	1.66Ki	0.0%	0	.symtab
8.8%	964	0.0%	0	.debug_info
6.7%	737	0.0%	0	.debug_str
5.3%	579	23.4%	579	[LOAD #2 [RX]]
4.9%	542	0.0%	0	.strtab
4.8%	528	18.8%	464	.dynamic
4.5%	498	17.6%	434	.text
3.5%	387	0.0%	0	.shstrtab
3.4%	374	0.0%	0	.debug_abbrev
3.2%	352	11.7%	288	.eh_frame
2.7%	296	0.0%	0	.debug_line
1.5%	160	3.9%	96	.dynsym
1.2%	132	2.8%	68	.eh_frame_hdr
1.2%	128	0.0%	0	[ELF Headers]
1.2%	127	2.5%	63	.dynstr
1.0%	112	0.0%	0	.debug_aranges
1.0%	112	1.9%	48	.rela.dyn
1.0%	107	0.0%	0	.comment
0.9%	100	1.5%	36	.note.gnu.build-id
0.9%	96	1.3%	32	.gnu.version_r
0.9%	96	1.3%	32	.got.plt
0.9%	96	1.3%	32	.note.ABI-tag
0.9%	96	1.3%	32	.plt
0.8%	92	1.1%	28	.gnu.hash
0.8%	92	1.1%	28	.interp
0.8%	88	1.0%	24	.rela.plt
0.8%	87	0.9%	23	.init
0.7%	80	0.6%	16	.data
0.7%	80	0.6%	16	.got
0.7%	76	0.5%	12	.rodata
0.7%	73	0.4%	9	.fini
0.0%	0	2.9%	72	.bss
0.7%	72	0.3%	8	.fini_array
0.7%	72	0.3%	8	.gnu.version
0.7%	72	0.3%	8	.init_array
0.0%	0	0.6%	16	[LOAD #3 [RW]]
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL


```
reka@ubuntu:~/binary-tools$ bloaty -n 0 ./dummy
```

FILE SIZE		VM SIZE		
16.1%	1.73Ki	0.0%	0	[Unmapped]
15.5%	1.66Ki	0.0%	0	.symtab
8.8%	964	0.0%	0	.debug_info
6.7%	737	0.0%	0	.debug_str
5.3%	579	23.4%	579	[LOAD #2 [RX]]
4.9%	542	0.0%	0	.strtab
4.8%	528	18.8%	464	.dynamic
4.5%	498	17.6%	434	.text
3.5%	387	0.0%	0	.shstrtab
3.4%	374	0.0%	0	.debug_abbrev
3.2%	352	11.7%	288	.eh_frame
2.7%	296	0.0%	0	.debug_line
1.5%	160	3.9%	96	.dynsym
1.2%	132	2.8%	68	.eh_frame_hdr
1.2%	128	0.0%	0	[ELF Headers]
1.2%	127	2.5%	63	.dynstr
1.0%	112	0.0%	0	.debug_aranges
1.0%	112	1.9%	48	.rela.dyn
1.0%	107	0.0%	0	.comment
0.9%	100	1.5%	36	.note.gnu.build-id
0.9%	96	1.3%	32	.gnu.version_r
0.9%	96	1.3%	32	.got.plt
0.9%	96	1.3%	32	.note.ABI-tag
0.9%	96	1.3%	32	.plt
0.8%	92	1.1%	28	.gnu.hash
0.8%	92	1.1%	28	.interp
0.8%	88	1.0%	24	.rela.plt
0.8%	87	0.9%	23	.init
0.7%	80	0.6%	16	.data
0.7%	80	0.6%	16	.got
0.7%	76	0.5%	12	.rodata
0.7%	73	0.4%	9	.fini
0.0%	0	2.9%	72	.bss
0.7%	72	0.3%	8	.fini_array
0.7%	72	0.3%	8	.gnu.version
0.7%	72	0.3%	8	.init_array
0.0%	0	0.6%	16	[LOAD #3 [RW]]
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

File size

Bytes on disk

VM size

Bytes in virtual
memory after the
executable is loaded

```
reka@ubuntu:~/binary-tools$ bloaty -n 0 ./dummy
```

FILE SIZE		VM SIZE		
16.1%	1.73Ki	0.0%	0	[Unmapped]
15.5%	1.66Ki	0.0%	0	.symtab
8.8%	964	0.0%	0	.debug_info
6.7%	737	0.0%	0	.debug_str
5.3%	579	23.4%	579	[LOAD #2 [RX]]
4.9%	542	0.0%	0	.strtab
4.8%	528	18.8%	464	.dynamic
4.5%	498	17.6%	434	.text
3.5%	387	0.0%	0	.shstrtab
3.4%	374	0.0%	0	.debug_abbrev
3.2%	352	11.7%	288	.eh_frame
2.7%	296	0.0%	0	.debug_line
1.5%	160	3.9%	96	.dynsym
1.2%	132	2.8%	68	.eh_frame_hdr
1.2%	128	0.0%	0	[ELF Headers]
1.2%	127	2.5%	63	.dynstr
1.0%	112	0.0%	0	.debug_aranges
1.0%	112	1.9%	48	.rela.dyn
1.0%	107	0.0%	0	.comment
0.9%	100	1.5%	36	.note.gnu.build-id
0.9%	96	1.3%	32	.gnu.version_r
0.9%	96	1.3%	32	.got.plt
0.9%	96	1.3%	32	.note.ABI-tag
0.9%	96	1.3%	32	.plt
0.8%	92	1.1%	28	.gnu.hash
0.8%	92	1.1%	28	.interp
0.8%	88	1.0%	24	.rela.plt
0.8%	87	0.9%	23	.init
0.7%	80	0.6%	16	.data
0.7%	80	0.6%	16	.got
0.7%	76	0.5%	12	.rodata
0.7%	73	0.4%	9	.fini
0.0%	0	2.9%	72	.bss
0.7%	72	0.3%	8	.fini_array
0.7%	72	0.3%	8	.gnu.version
0.7%	72	0.3%	8	.init_array
0.0%	0	0.6%	16	[LOAD #3 [RW]]
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

File size

Bytes on disk

VM size

Bytes in virtual
memory after the
executable is loaded

```
reka@ubuntu:~/binary-tools$ bloaty -n 0 ./dummy
```

FILE SIZE		VM SIZE		
16.1%	1.73Ki	0.0%	0	[Unmapped]
15.5%	1.66Ki	0.0%	0	.symtab
8.8%	964	0.0%	0	.debug_info
6.7%	737	0.0%	0	.debug_str
5.3%	579	23.4%	579	[LOAD #2 [RX]]
4.9%	542	0.0%	0	.strtab
4.8%	528	18.8%	464	.dynamic
4.5%	498	17.6%	434	.text
3.5%	387	0.0%	0	.shstrtab
3.4%	374	0.0%	0	.debug_abbrev
3.2%	352	11.7%	288	.eh_frame
2.7%	296	0.0%	0	.debug_line
1.5%	160	3.9%	96	.dynsym
1.2%	132	2.8%	68	.eh_frame_hdr
1.2%	128	0.0%	0	[ELF Headers]
1.2%	127	2.5%	63	.dynstr
1.0%	112	0.0%	0	.debug_aranges
1.0%	112	1.9%	48	.rela.dyn
1.0%	107	0.0%	0	.comment
0.9%	100	1.5%	36	.note.gnu.build-id
0.9%	96	1.3%	32	.gnu.version_r
0.9%	96	1.3%	32	.got.plt
0.9%	96	1.3%	32	.note.ABI-tag
0.9%	96	1.3%	32	.plt
0.8%	92	1.1%	28	.gnu.hash
0.8%	92	1.1%	28	.interp
0.8%	88	1.0%	24	.rela.plt
0.8%	87	0.9%	23	.init
0.7%	80	0.6%	16	.data
0.7%	80	0.6%	16	.got
0.7%	76	0.5%	12	.rodata
0.7%	73	0.4%	9	.fini
0.0%	0	2.9%	72	.bss
0.7%	72	0.3%	8	.fini_array
0.7%	72	0.3%	8	.gnu.version
0.7%	72	0.3%	8	.init_array
0.0%	0	0.6%	16	[LOAD #3 [RW]]
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

File size

Bytes on disk

VM size

Bytes in virtual
memory after the
executable is loaded

ELF (Executable and Linkage Format)

File type for binaries on Linux

3 flavors:

- **Relocatable** – needs to be processed by the linker
- **Executable** – all symbols resolved (except for shared lib symbols), runnable
- **Shared object** – has both symbol info for the linker and runnable code

Linkable
sections

Executable
segments

(Optional, ignored)

Sections

Describes sections

ELF header

Program header table


Section header table

Describes segments

Segments

(Optional, ignored)

John Levine: Linkers and Loaders



Segments and sections

```
reka@ubuntu:~/binary-tools$ bloaty -d segments ./dummy
```

FILE SIZE		VM SIZE		
-----		-----		
58.4%	6.26Ki	0.0%	0	[Unmapped]
19.8%	2.12Ki	0.0%	0	[ELF Headers]
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX]
5.0%	544	25.6%	632	LOAD #3 [RW]
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

Segments and sections

```
reka@ubuntu:~/binary-tools$ bloaty -d segments ./dummy
```

FILE SIZE		VM SIZE		
-----	-----	-----	-----	
58.4%	6.26Ki	0.0%	0	[Unmapped]
19.8%	2.12Ki	0.0%	0	[ELF Headers]
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX]
5.0%	544	25.6%	632	LOAD #3 [RW]
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

```
reka@ubuntu:~/binary-tools$ bloaty -d segments,sections dummy
```

```
reka@ubuntu:~/binary-tools$ bloaty -d segments,sections dummy
```

FILE SIZE		VM SIZE		
-----	-----	-----	-----	
58.4%	6.26Ki	0.0%	0	[Unmapped]
27.6%	1.73Ki	NAN%	0	[Unmapped]
25.5%	1.59Ki	NAN%	0	.symtab
14.0%	900	NAN%	0	.debug_info
10.5%	673	NAN%	0	.debug_str
7.5%	478	NAN%	0	.strtab
5.0%	323	NAN%	0	.shstrtab
4.8%	310	NAN%	0	.debug_abbrev
3.6%	232	NAN%	0	.debug_line
0.7%	48	NAN%	0	.debug_aranges
0.7%	43	NAN%	0	.comment
19.8%	2.12Ki	0.0%	0	[ELF Headers]
38.2%	832	NAN%	0	[13 Others]
5.9%	128	NAN%	0	[ELF Headers]
2.9%	64	NAN%	0	.comment
2.9%	64	NAN%	0	.data
2.9%	64	NAN%	0	.debug_abbrev
2.9%	64	NAN%	0	.debug_aranges
2.9%	64	NAN%	0	.debug_info
2.9%	64	NAN%	0	.debug_line
2.9%	64	NAN%	0	.debug_str
2.9%	64	NAN%	0	.dynamic
2.9%	64	NAN%	0	.dynstr
2.9%	64	NAN%	0	.dynsym
2.9%	64	NAN%	0	.eh_frame
2.9%	64	NAN%	0	.eh_frame_hdr
2.9%	64	NAN%	0	.fini
2.9%	64	NAN%	0	.fini_array
2.9%	64	NAN%	0	.gnu.hash
2.9%	64	NAN%	0	.gnu.version
2.9%	64	NAN%	0	.gnu.version_r
2.9%	64	NAN%	0	.got
2.9%	64	NAN%	0	.got.plt
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX]
31.5%	579	31.5%	579	[LOAD #2 [RX]]
23.6%	434	23.6%	434	.text
15.7%	288	15.7%	288	.eh_frame
5.2%	96	5.2%	96	.dynsym
3.7%	68	3.7%	68	.eh_frame_hdr
3.4%	63	3.4%	63	.dynstr
2.6%	48	2.6%	48	.rela.dyn
2.0%	36	2.0%	36	.note.gnu.build-id
1.7%	32	1.7%	32	.gnu.version_r
1.7%	32	1.7%	32	.note.ABI-tag
1.7%	32	1.7%	32	.plt
1.5%	28	1.5%	28	.gnu.hash
1.5%	28	1.5%	28	.interp
1.3%	24	1.3%	24	.rela.plt
1.2%	23	1.2%	23	.init
0.7%	12	0.7%	12	.rodata
0.5%	9	0.5%	9	.fini
0.4%	8	0.4%	8	.gnu.version
5.0%	544	25.6%	632	LOAD #3 [RW]
85.3%	464	73.4%	464	.dynamic
0.0%	0	11.4%	72	.bss
5.9%	32	5.1%	32	.got.plt
2.9%	16	2.5%	16	.data
2.9%	16	2.5%	16	.got
0.0%	0	2.5%	16	[LOAD #3 [RW]]
1.5%	8	1.3%	8	.fini_array
1.5%	8	1.3%	8	.init_array
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

Segments and sections

```
reka@ubuntu:~/binary-tools$ bloaty -d segments ./dummy
```

FILE SIZE		VM SIZE		
-----	-----	-----	-----	
58.4%	6.26Ki	0.0%	0	[Unmapped]
19.8%	2.12Ki	0.0%	0	[ELF Headers]
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX]
5.0%	544	25.6%	632	LOAD #3 [RW]
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

```
reka@ubuntu:~/binary-tools$ bloaty -d segments,sections dummy
```

```
reka@ubuntu:~/binary-tools$ bloaty -d segments,sections dummy
```

FILE SIZE		VM SIZE		
-----	-----	-----	-----	
58.4%	6.26Ki	0.0%	0	[Unmapped]
27.6%	1.73Ki	NAN%	0	[Unmapped]
25.5%	1.59Ki	NAN%	0	.symtab
14.0%	900	NAN%	0	.debug_info
10.5%	673	NAN%	0	.debug_str
7.5%	478	NAN%	0	.strtab
5.0%	323	NAN%	0	.shstrtab
4.8%	310	NAN%	0	.debug_abbrev
3.6%	232	NAN%	0	.debug_line
0.7%	48	NAN%	0	.debug_aranges
0.7%	43	NAN%	0	.comment
19.8%	2.12Ki	0.0%	0	[ELF Headers]
38.2%	832	NAN%	0	[13 Others]
5.9%	128	NAN%	0	[ELF Headers]
2.9%	64	NAN%	0	.comment
2.9%	64	NAN%	0	.data
2.9%	64	NAN%	0	.debug_abbrev
2.9%	64	NAN%	0	.debug_aranges
2.9%	64	NAN%	0	.debug_info
2.9%	64	NAN%	0	.debug_line
2.9%	64	NAN%	0	.debug_str
2.9%	64	NAN%	0	.dynamic
2.9%	64	NAN%	0	.dynstr
2.9%	64	NAN%	0	.dynsym
2.9%	64	NAN%	0	.eh_frame
2.9%	64	NAN%	0	.eh_frame_hdr
2.9%	64	NAN%	0	.fini
2.9%	64	NAN%	0	.fini_array
2.9%	64	NAN%	0	.gnu.hash
2.9%	64	NAN%	0	.gnu.version
2.9%	64	NAN%	0	.gnu.version_r
2.9%	64	NAN%	0	.got
2.9%	64	NAN%	0	.got.plt
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX]
31.5%	579	31.5%	579	[LOAD #2 [RX]]
23.6%	434	23.6%	434	.text
15.7%	288	15.7%	288	.eh_frame
5.2%	96	5.2%	96	.dynsym
3.7%	68	3.7%	68	.eh_frame_hdr
3.4%	63	3.4%	63	.dynstr
2.6%	48	2.6%	48	.rela.dyn
2.0%	36	2.0%	36	.note.gnu.build-id
1.7%	32	1.7%	32	.gnu.version_r
1.7%	32	1.7%	32	.note.ABI-tag
1.7%	32	1.7%	32	.plt
1.5%	28	1.5%	28	.gnu.hash
1.5%	28	1.5%	28	.interp
1.3%	24	1.3%	24	.rela.plt
1.2%	23	1.2%	23	.init
0.7%	12	0.7%	12	.rodata
0.5%	9	0.5%	9	.fini
0.4%	8	0.4%	8	.gnu.version
5.0%	544	25.6%	632	LOAD #3 [RW]
85.3%	464	73.4%	464	.dynamic
0.0%	0	11.4%	72	.bss
5.9%	32	5.1%	32	.got.plt
2.9%	16	2.5%	16	.data
2.9%	16	2.5%	16	.got
0.0%	0	2.5%	16	[LOAD #3 [RW]]
1.5%	8	1.3%	8	.fini_array
1.5%	8	1.3%	8	.init_array
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

Segments and sections

```
reka@ubuntu:~/binary-tools$ bloaty -d segments ./dummy
```

FILE SIZE		VM SIZE		
-----		-----		
58.4%	6.26Ki	0.0%	0	[Unmapped]
19.8%	2.12Ki	0.0%	0	[ELF Headers]
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX]
5.0%	544	25.6%	632	LOAD #3 [RW]
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

```
reka@ubuntu:~/binary-tools$ bloaty -d segments,sections dummy
```

FILE SIZE		VM SIZE		
-----		-----		
58.4%	6.26Ki	0.0%	0	[Unmapped]
27.6%	1.73Ki	NAN%	0	[Unmapped]
25.5%	1.59Ki	NAN%	0	.symtab
14.0%	900	NAN%	0	.debug_info
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX]
31.5%	579	31.5%	579	[LOAD #2 [RX]]
23.6%	434	23.6%	434	.text
15.7%	288	15.7%	288	.eh_frame
5.2%	96	5.2%	96	.dynsym
3.7%	68	3.7%	68	.eh_frame_hdr
3.4%	63	3.4%	63	.dynstr
2.6%	48	2.6%	48	.rela.dyn
2.0%	36	2.0%	36	.note.gnu.build-id
1.7%	32	1.7%	32	.gnu.version_r
1.7%	32	1.7%	32	.note.ABI-tag
1.7%	32	1.7%	32	.plt
1.5%	28	1.5%	28	.gnu.hash
1.5%	28	1.5%	28	.interp
1.3%	24	1.3%	24	.rela.plt
1.2%	23	1.2%	23	.init
0.7%	12	0.7%	12	.rodata
0.5%	9	0.5%	9	.fini
0.4%	8	0.4%	8	.gnu.version
5.0%	544	25.6%	632	LOAD #3 [RW]
85.3%	464	73.4%	464	.dynamic
0.0%	0	11.4%	72	.bss
5.9%	32	5.1%	32	.got.plt
2.9%	16	2.5%	16	.data
2.9%	16	2.5%	16	.got
0.0%	0	2.5%	16	[LOAD #3 [RW]]
1.5%	8	1.3%	8	.fini_array
1.5%	8	1.3%	8	.init_array
2.9%	16	2.5%	16	.data
2.9%	16	2.5%	16	.got
0.0%	0	2.5%	16	[LOAD #3 [RW]]
1.5%	8	1.3%	8	.fini_array
1.5%	8	1.3%	8	.init_array
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

Segments and sections

```
reka@ubuntu:~/binary-tools$ bloaty -d segments ./dummy
```

FILE SIZE		VM SIZE		
-----		-----		
58.4%	6.26Ki	0.0%	0	[Unmapped]
19.8%	2.12Ki	0.0%	0	[ELF Headers]
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX] ←
5.0%	544	25.6%	632	LOAD #3 [RW]
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

```
reka@ubuntu:~/binary-tools$ bloaty -d segments,sections dummy
```

FILE SIZE		VM SIZE		
-----		-----		
58.4%	6.26Ki	0.0%	0	[Unmapped]
27.6%	1.73Ki	NAN%	0	[Unmapped]
25.5%	1.59Ki	NAN%	0	.symtab
14.0%	900	NAN%	0	.debug_info
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX]
31.5%	579	31.5%	579	[LOAD #2 [RX]]
23.6%	434	23.6%	434	.text ←
15.7%	288	15.7%	288	.eh_frame
5.2%	96	5.2%	96	.dynsym
3.7%	68	3.7%	68	.eh_frame_hdr
3.4%	63	3.4%	63	.dynstr
2.6%	48	2.6%	48	.rela.dyn
2.0%	36	2.0%	36	.note.gnu.build-id
1.7%	32	1.7%	32	.gnu.version_r
1.7%	32	1.7%	32	.note.ABI-tag
1.7%	32	1.7%	32	.plt
1.5%	28	1.5%	28	.gnu.hash
1.5%	28	1.5%	28	.interp
1.3%	24	1.3%	24	.rela.plt
1.2%	23	1.2%	23	.init
0.7%	12	0.7%	12	.rodata ←
0.5%	9	0.5%	9	.fini
0.4%	8	0.4%	8	.gnu.version
5.0%	544	25.6%	632	LOAD #3 [RW]
85.3%	464	73.4%	464	.dynamic
0.0%	0	11.4%	72	.bss
5.9%	32	5.1%	32	.got.plt
2.9%	16	2.5%	16	.data
2.9%	16	2.5%	16	.got
0.0%	0	2.5%	16	[LOAD #3 [RW]]
1.5%	8	1.3%	8	.fini_array
1.5%	8	1.3%	8	.init_array
2.9%	16	2.5%	16	.data
2.9%	16	2.5%	16	.got
0.0%	0	2.5%	16	[LOAD #3 [RW]]
1.5%	8	1.3%	8	.fini_array
1.5%	8	1.3%	8	.init_array
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

Segments and sections

```
reka@ubuntu:~/binary-tools$ bloaty -d segments ./dummy
```

FILE SIZE		VM SIZE		
-----		-----		
58.4%	6.26Ki	0.0%	0	[Unmapped]
19.8%	2.12Ki	0.0%	0	[ELF Headers]
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX]
5.0%	544	25.6%	632	LOAD #3 [RW] ←
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

```
reka@ubuntu:~/binary-tools$ bloaty -d segments,sections dummy
```

FILE SIZE		VM SIZE		
-----		-----		
58.4%	6.26Ki	0.0%	0	[Unmapped]
27.6%	1.73Ki	NAN%	0	[Unmapped]
25.5%	1.59Ki	NAN%	0	.symtab
14.0%	900	NAN%	0	.debug_info
16.8%	1.80Ki	74.4%	1.80Ki	LOAD #2 [RX]
31.5%	579	31.5%	579	[LOAD #2 [RX]]
23.6%	434	23.6%	434	.text
15.7%	288	15.7%	288	.eh_frame
5.2%	96	5.2%	96	.dynsym
3.7%	68	3.7%	68	.eh_frame_hdr
3.4%	63	3.4%	63	.dynstr
2.6%	48	2.6%	48	.rela.dyn
2.0%	36	2.0%	36	.note.gnu.build-id
1.7%	32	1.7%	32	.gnu.version_r
1.7%	32	1.7%	32	.note.ABI-tag
1.7%	32	1.7%	32	.plt
1.5%	28	1.5%	28	.gnu.hash
1.5%	28	1.5%	28	.interp
1.3%	24	1.3%	24	.rela.plt
1.2%	23	1.2%	23	.init
0.7%	12	0.7%	12	.rodata
0.5%	9	0.5%	9	.fini
0.4%	8	0.4%	8	.gnu.version
5.0%	544	25.6%	632	LOAD #3 [RW]
85.3%	464	73.4%	464	.dynamic
0.0%	0	11.4%	72	.bss ←
5.9%	32	5.1%	32	.got.plt
2.9%	16	2.5%	16	.data ←
2.9%	16	2.5%	16	.got
0.0%	0	2.5%	16	[LOAD #3 [RW]]
1.5%	8	1.3%	8	.fini_array
1.5%	8	1.3%	8	.init_array
2.9%	16	2.5%	16	.data
2.9%	16	2.5%	16	.got
0.0%	0	2.5%	16	[LOAD #3 [RW]]
1.5%	8	1.3%	8	.fini_array
1.5%	8	1.3%	8	.init_array
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

Track changes

```

1 #include <stdio.h>
2
3 int zero_init[10];
4
5 void do_stuff(int arg) {
6     int local = arg + 2;
7     int i;
8
9     for (i = 0; i < local; ++i) {
10         printf("i = %d\n", i);
11     }
12 }
13
14 int main() {
15     do_stuff(2);
16     return 0;
17 }

```

```

reka@ubuntu:~/binary-tools$ bloaty -d compileunits,sections,symbols ./dummy

```

FILE SIZE		VM SIZE		
-----		-----		
19.3%	2.06Ki	9.7%	239	dummy.c
45.6%	964	0.0%	0	.debug_info
93.4%	900	NAN%	0	[section .debug_info]
6.6%	64	NAN%	0	[ELF Headers]
24.4%	515	0.0%	0	.debug_str
14.0%	296	0.0%	0	.debug_line
78.4%	232	NAN%	0	[section .debug_line]
21.6%	64	NAN%	0	[ELF Headers]
4.5%	96	0.0%	0	.symtab
50.0%	48	NAN%	0	do_stuff
25.0%	24	NAN%	0	main
25.0%	24	NAN%	0	zero_init
4.1%	87	36.4%	87	.text
75.9%	66	75.9%	66	do_stuff
24.1%	21	24.1%	21	main
3.0%	64	26.8%	64	.eh_frame
50.0%	32	50.0%	32	do_stuff
50.0%	32	50.0%	32	main
2.0%	43	0.0%	0	.strtab
65.1%	28	NAN%	0	do_stuff
23.3%	10	NAN%	0	zero_init
11.6%	5	NAN%	0	main
0.0%	0	16.7%	40	.bss
NAN%	0	100.0%	40	zero_init
1.1%	24	10.0%	24	.rela.dyn
100.0%	24	100.0%	24	_start
0.8%	16	6.7%	16	.eh_frame_hdr
50.0%	8	50.0%	8	do_stuff
50.0%	8	50.0%	8	main
0.4%	8	3.3%	8	.rodata
100.0%	8	100.0%	8	do_stuff
18.7%	2.00Ki	0.0%	0	[ELF Headers]

Track changes

```

1 #include <stdio.h>
2
3 int zero_init[10]; ←
4
5 void do_stuff(int arg) {
6     int local = arg + 2;
7     int i;
8
9     for (i = 0; i < local; ++i) {
10         printf("i = %d\n", i);
11     }
12 }
13
14 int main() {
15     do_stuff(2);
16     return 0;
17 }

```

```

reka@ubuntu:~/binary-tools$ bloaty -d compileunits,sections,symbols ./dummy

```

FILE SIZE		VM SIZE		
-----	-----	-----	-----	
19.3%	2.06Ki	9.7%	239	dummy.c
45.6%	964	0.0%	0	.debug_info
93.4%	900	NAN%	0	[section .debug_info]
6.6%	64	NAN%	0	[ELF Headers]
24.4%	515	0.0%	0	.debug_str
14.0%	296	0.0%	0	.debug_line
78.4%	232	NAN%	0	[section .debug_line]
21.6%	64	NAN%	0	[ELF Headers]
4.5%	96	0.0%	0	.symtab
50.0%	48	NAN%	0	do_stuff
25.0%	24	NAN%	0	main
25.0%	24	NAN%	0	zero_init
4.1%	87	36.4%	87	.text
75.9%	66	75.9%	66	do_stuff
24.1%	21	24.1%	21	main
3.0%	64	26.8%	64	.eh_frame
50.0%	32	50.0%	32	do_stuff
50.0%	32	50.0%	32	main
2.0%	43	0.0%	0	.strtab
65.1%	28	NAN%	0	do_stuff
23.3%	10	NAN%	0	zero_init
11.6%	5	NAN%	0	main
0.0%	0	16.7%	40	.bss
NAN%	0	100.0%	40	zero_init
1.1%	24	10.0%	24	.rela.dyn
100.0%	24	100.0%	24	_start
0.8%	16	6.7%	16	.eh_frame_hdr
50.0%	8	50.0%	8	do_stuff
50.0%	8	50.0%	8	main
0.4%	8	3.3%	8	.rodata
100.0%	8	100.0%	8	do_stuff
18.7%	2.00Ki	0.0%	0	[ELF Headers]

Track changes

```

1 #include <stdio.h>
2
3 int zero_init[10] = {1}; ←
4
5 void do_stuff(int arg) {
6     int local = arg + 2;
7     int i;
8
9     for (i = 0; i < local; ++i) {
10         printf("i = %d\n", i);
11     }
12 }
13
14 int main() {
15     do_stuff(2);
16     return 0;
17 }

```

```

reka@ubuntu:~/binary-tools$ bloaty -d compileunits,sections,symbols ./dummy

```

FILE SIZE		VM SIZE		
19.5%	2.10Ki	9.8%	239	dummy.c
44.8%	964	0.0%	0	.debug_info
93.4%	900	NAN%	0	[section .debug_info]
6.6%	64	NAN%	0	[ELF Headers]
23.9%	515	0.0%	0	.debug_str
13.7%	296	0.0%	0	.debug_line
78.4%	232	NAN%	0	[section .debug_line]
21.6%	64	NAN%	0	[ELF Headers]
4.5%	96	0.0%	0	.symtab
50.0%	48	NAN%	0	do_stuff
25.0%	24	NAN%	0	main
25.0%	24	NAN%	0	zero_init
4.0%	87	36.4%	87	.text
75.9%	66	75.9%	66	do_stuff
24.1%	21	24.1%	21	main
3.0%	64	26.8%	64	.eh_frame
50.0%	32	50.0%	32	do_stuff
50.0%	32	50.0%	32	main
2.0%	43	0.0%	0	.strtab
65.1%	28	NAN%	0	do_stuff
23.3%	10	NAN%	0	zero_init
11.6%	5	NAN%	0	main
1.9%	40	16.7%	40	.data
100.0%	40	100.0%	40	zero_init
1.1%	24	10.0%	24	.rela.dyn
100.0%	24	100.0%	24	_start
0.7%	16	6.7%	16	.eh_frame_hdr
50.0%	8	50.0%	8	do_stuff
50.0%	8	50.0%	8	main
0.4%	8	3.3%	8	.rodata
100.0%	8	100.0%	8	do_stuff
18.6%	2.00Ki	0.0%	0	[ELF Headers]

Diff mode

```
$ vim olddummy.c
```

```
1 #include <stdio.h>
2
3 void do_stuff() {
4     printf("stuff\n");
5 }
6
7 int main() {
8     do_stuff();
9     return 0;
10 }
```

```
$ vim dummy.c
```

```
1 #include <stdio.h>
2
3 int zero_init[10];
4
5 void do_stuff(int arg) {
6     int local = arg + 2;
7     int i;
8
9     for (i = 0; i < local; ++i) {
10         printf("i = %d\n", i);
11     }
12 }
13
14 int main() {
15     do_stuff(2);
16     return 0;
17 }
```

Diff mode

```
$ vim olddummy.c
```

```
1 #include <stdio.h>
2
3 void do_stuff() {
4     printf("stuff\n");
5 }
6
7 int main() {
8     do_stuff();
9     return 0;
10 }
```

```
$ vim dummy.c
```

```
1 #include <stdio.h>
2
3 int zero_init[10];
4
5 void do_stuff(int arg) {
6     int local = arg + 2;
7     int i;
8
9     for (i = 0; i < local; ++i) {
10         printf("i = %d\n", i);
11     }
12 }
13
14 int main() {
15     do_stuff(2);
16     return 0;
17 }
```

```
reka@ubuntu:~/binary-tools$ bloaty -d symbols,sections dummy -- olddummy
```

FILE SIZE		VM SIZE		
-----		-----		
+30%	+103	[=]	0	_start
+60%	+72	[=]	0	.symtab
+65%	+31	[=]	0	.strtab
+9.5%	+78	[=]	0	[section .debug_info]
+26%	+64	[=]	0	[section .debug_abbrev]
+35%	+49	+75%	+49	do_stuff
+247%	+47	+247%	+47	.text
+33%	+2	+33%	+2	.rodata
[NEW]	+34	[NEW]	+40	zero_init
[=]	0	[NEW]	+40	.bss
[NEW]	+24	[=]	0	.symtab
[NEW]	+10	[=]	0	.strtab
[=]	0	+343%	+24	[section .bss]
[=]	0	[NEW]	+16	[LOAD #3 [RW]]
+7.4%	+16	[=]	0	[section .debug_line]
+2.0%	+13	[=]	0	[section .debug_str]
+3.3%	+2	+3.3%	+2	[section .dynstr]
+0.5%	+1	+0.5%	+1	[section .text]
-0.6%	-1	[=]	0	[section .strtab]
-0.7%	-4	-0.7%	-4	[LOAD #2 [RX]]
-2.4%	-44	[=]	0	[Unmapped]
-72.5%	-103	[=]	0	completed.7697
-67.4%	-31	[=]	0	.strtab
-75.0%	-72	[=]	0	.symtab
+1.9%	+208	+5.5%	+128	TOTAL

Diff mode

```
$ vim olddummy.c
```

```
1 #include <stdio.h>
2
3 void do_stuff() {
4     printf("stuff\n");
5 }
6
7 int main() {
8     do_stuff();
9     return 0;
10 }
```

```
$ vim dummy.c
```

```
1 #include <stdio.h>
2
3 int zero_init[10];
4
5 void do_stuff(int arg) {
6     int local = arg + 2;
7     int i;
8
9     for (i = 0; i < local; ++i) {
10         printf("i = %d\n", i);
11     }
12 }
13
14 int main() {
15     do_stuff(2);
16     return 0;
17 }
```

```
reka@ubuntu:~/binary-tools$ bloaty -d symbols,sections dummy -- olddummy
```

FILE SIZE		VM SIZE		

+30%	+103	[=]	0	_start
+60%	+72	[=]	0	.symtab
+65%	+31	[=]	0	.strtab
+9.5%	+78	[=]	0	[section .debug_info]
+26%	+64	[=]	0	[section .debug_abbrev]
+35%	+49	+75%	+49	do_stuff
+247%	+47	+247%	+47	.text
+33%	+2	+33%	+2	.rodata
[NEW]	+34	[NEW]	+40	zero_init
[=]	0	[NEW]	+40	.bss
[NEW]	+24	[=]	0	.symtab
[NEW]	+10	[=]	0	.strtab
[=]	0	+343%	+24	[section .bss]
[=]	0	[NEW]	+16	[LOAD #3 [RW]]
+7.4%	+16	[=]	0	[section .debug_line]
+2.0%	+13	[=]	0	[section .debug_str]
+3.3%	+2	+3.3%	+2	[section .dynstr]
+0.5%	+1	+0.5%	+1	[section .text]
-0.6%	-1	[=]	0	[section .strtab]
-0.7%	-4	-0.7%	-4	[LOAD #2 [RX]]
-2.4%	-44	[=]	0	[Unmapped]
-72.5%	-103	[=]	0	completed.7697
-67.4%	-31	[=]	0	.strtab
-75.0%	-72	[=]	0	.symtab
+1.9%	+208	+5.5%	+128	TOTAL

Diff mode

\$ vim olddummy.c

```
1 #include <stdio.h>
2
3 void do_stuff() {
4     printf("stuff\n");
5 }
6
7 int main() {
8     do_stuff();
9     return 0;
10 }
```

\$ vim dummy.c

```
1 #include <stdio.h>
2
3 int zero_init[10];
4
5 void do_stuff(int arg) {
6     int local = arg + 2;
7     int i;
8
9     for (i = 0; i < local; ++i) {
10         printf("i = %d\n", i);
11     }
12 }
13
14 int main() {
15     do_stuff(2);
16     return 0;
17 }
```

reka@ubuntu:~/binary-tools\$ bloaty -d symbols,sections dummy -- olddummy

FILE SIZE		VM SIZE		
-----		-----		
+30%	+103	[=]	0	_start
+60%	+72	[=]	0	.symtab
+65%	+31	[=]	0	.strtab
+9.5%	+78	[=]	0	[section .debug_info]
+26%	+64	[=]	0	[section .debug_abbrev]
+35%	+49	+75%	+49	do_stuff
+247%	+47	+247%	+47	.text
+33%	+2	+33%	+2	.rodata
[NEW]	+34	[NEW]	+40	zero_init
[=]	0	[NEW]	+40	.bss
[NEW]	+24	[=]	0	.symtab
[NEW]	+10	[=]	0	.strtab
[=]	0	+343%	+24	[section .bss]
[=]	0	[NEW]	+16	[LOAD #3 [RW]]
+7.4%	+16	[=]	0	[section .debug_line]
+2.0%	+13	[=]	0	[section .debug_str]
+3.3%	+2	+3.3%	+2	[section .dynstr]
+0.5%	+1	+0.5%	+1	[section .text]
-0.6%	-1	[=]	0	[section .strtab]
-0.7%	-4	-0.7%	-4	[LOAD #2 [RX]]
-2.4%	-44	[=]	0	[Unmapped]
-72.5%	-103	[=]	0	completed.7697
-67.4%	-31	[=]	0	.strtab
-75.0%	-72	[=]	0	.symtab
+1.9%	+208	+5.5%	+128	TOTAL

Data Sources

```
reka@ubuntu:~/binary-tools$ bloaty -d symbols,sections ./dummy
```

1.8%	199	8.1%	199	[section .text]
1.7%	190	4.6%	114	do_stuff
34.7%	66	57.9%	66	.text
25.3%	48	0.0%	0	.symtab
16.8%	32	28.1%	32	.eh_frame
14.7%	28	0.0%	0	.strtab
4.2%	8	7.0%	8	.eh_frame_hdr
4.2%	8	7.0%	8	.rodata
1.4%	155	0.0%	0	[section .strtab]
0.9%	96	3.9%	96	[section .dynsym]
0.8%	92	3.7%	92	[section .eh_frame]
0.8%	90	2.5%	61	main
35.6%	32	52.5%	32	.eh_frame
26.7%	24	0.0%	0	.symtab
23.3%	21	34.4%	21	.text
8.9%	8	13.1%	8	.eh_frame_hdr
5.6%	5	0.0%	0	.strtab
0.7%	78	1.2%	30	dl_relocate_static_pie

```
reka@ubuntu:~/binary-tools$ bloaty -d symbols,sections ./dummy
```

FILE SIZE		VM SIZE		
18.1%	1.94Ki	0.0%	0	[ELF Headers]
32.3%	640	NAN%	0	[10 Others]
6.5%	128	NAN%	0	[ELF Headers]
3.2%	64	NAN%	0	.comment
3.2%	64	NAN%	0	.data
3.2%	64	NAN%	0	.debug_abbrev
3.2%	64	NAN%	0	.debug_aranges
3.2%	64	NAN%	0	.debug_info
3.2%	64	NAN%	0	.debug_line
3.2%	64	NAN%	0	.debug_str
3.2%	64	NAN%	0	.dynamic
3.2%	64	NAN%	0	.dynstr
3.2%	64	NAN%	0	.dynsym
3.2%	64	NAN%	0	.eh_frame
3.2%	64	NAN%	0	.eh_frame_hdr
3.2%	64	NAN%	0	.fini
3.2%	64	NAN%	0	.gnu.hash
3.2%	64	NAN%	0	.gnu.version
3.2%	64	NAN%	0	.gnu.version_r
3.2%	64	NAN%	0	.got.plt
3.2%	64	NAN%	0	.init
3.2%	64	NAN%	0	.interp
16.1%	1.73Ki	0.0%	0	[Unmapped]
10.1%	1.08Ki	0.0%	0	[section .symtab]
8.2%	900	0.0%	0	[section .debug_info]
6.1%	673	0.0%	0	[section .debug_str]
5.6%	614	20.2%	499	[22 Others]
5.5%	600	8.0%	197	__libc_csu_init
24.0%	144	0.0%	0	.symtab
21.8%	131	0.0%	0	.strtab
16.8%	101	51.3%	101	.text
12.0%	72	36.5%	72	.eh_frame
12.0%	72	4.1%	8	.fini_array
12.0%	72	4.1%	8	.init_array
1.3%	8	4.1%	8	.eh_frame_hdr
5.3%	579	23.4%	579	[LOAD #2 [RX]]
4.2%	464	10.8%	464	[section .dynamic]
4.1%	446	4.5%	111	.start
43.0%	192	0.0%	0	.symtab
17.9%	80	14.4%	16	.got
17.7%	79	0.0%	0	.strtab
9.6%	43	38.7%	43	.text
5.4%	24	21.6%	24	.rela.dyn
4.5%	20	18.0%	20	.eh_frame
1.8%	8	7.2%	8	.eh_frame_hdr
2.9%	323	0.0%	0	[section .shstrtab]
2.8%	310	0.0%	0	[section .debug_abbrev]
2.1%	232	0.0%	0	[section .debug_line]
1.0%	199	0.1%	199	[section .text]
1.7%	190	4.6%	114	do_stuff
34.7%	66	57.9%	66	.text
25.3%	48	0.0%	0	.symtab
16.8%	32	28.1%	32	.eh_frame
14.7%	28	0.0%	0	.strtab
4.2%	8	7.0%	8	.eh_frame_hdr
4.2%	8	7.0%	8	.rodata
1.4%	155	0.0%	0	[section .strtab]
0.9%	96	3.9%	96	[section .dynsym]
0.8%	92	3.7%	92	[section .eh_frame]
0.8%	90	2.5%	61	main
35.6%	32	52.5%	32	.eh_frame
26.7%	24	0.0%	0	.symtab
23.3%	21	34.4%	21	.text
8.9%	8	13.1%	8	.eh_frame_hdr
5.6%	5	0.0%	0	.strtab
0.7%	78	1.2%	30	dl_relocate_static_pie
30.8%	24	0.0%	0	.strtab
30.8%	24	0.0%	0	.symtab
25.6%	20	66.7%	20	.eh_frame
10.3%	8	26.7%	8	.eh_frame_hdr
2.6%	2	6.7%	2	.text
0.6%	70	1.2%	30	__libc_csu_fini
34.3%	24	0.0%	0	.symtab
28.6%	20	66.7%	20	.eh_frame
22.9%	16	0.0%	0	.strtab
11.4%	8	26.7%	8	.eh_frame_hdr
2.9%	2	6.7%	2	.text
100.0%	10.7Ki	100.0%	2.41Ki	TOTAL

Data Sources

```
reka@ubuntu:~/binary-tools$ bloaty -d symbols,sections ./dummy
```

1.8%	199	8.1%	199	[section .text]
1.7%	190	4.6%	114	do_stuff
34.7%	66	57.9%	66	.text
25.3%	48	0.0%	0	.symtab
16.8%	32	28.1%	32	.eh_frame ←
14.7%	28	0.0%	0	.strtab
4.2%	8	7.0%	8	.eh_frame_hdr ←
4.2%	8	7.0%	8	.rodata
1.4%	155	0.0%	0	[section .strtab]
0.9%	96	3.9%	96	[section .dynsym]
0.8%	92	3.7%	92	[section .eh_frame]
0.8%	90	2.5%	61	main
35.6%	32	52.5%	32	.eh_frame ←
26.7%	24	0.0%	0	.symtab
23.3%	21	34.4%	21	.text
8.9%	8	13.1%	8	.eh_frame_hdr ←
5.6%	5	0.0%	0	.strtab
0.7%	78	1.2%	30	dl_relocate_static_pie

The **symbol table** only refers to the machine code of the function

Other artifacts a function emits into the binary:

- **Unwind info**

Data Sources

```
reka@ubuntu:~/binary-tools$ bloaty -d symbols,sections ./dummy
```

1.8%	199	8.1%	199	[section .text]
1.7%	190	4.6%	114	do_stuff
34.7%	66	57.9%	66	.text
25.3%	48	0.0%	0	.symtab
16.8%	32	28.1%	32	.eh_frame
14.7%	28	0.0%	0	.strtab
4.2%	8	7.0%	8	.eh_frame_hdr
4.2%	8	7.0%	8	.rodata
1.4%	155	0.0%	0	[section .strtab]
0.9%	96	3.9%	96	[section .dynsym]
0.8%	92	3.7%	92	[section .eh_frame]
0.8%	90	2.5%	61	main
35.6%	32	52.5%	32	.eh_frame
26.7%	24	0.0%	0	.symtab
23.3%	21	34.4%	21	.text
8.9%	8	13.1%	8	.eh_frame_hdr
5.6%	5	0.0%	0	.strtab
0.7%	78	1.2%	30	dl_relocate_static_pie

The **symbol table** only refers to the machine code of the function

Other artifacts a function emits into the binary:

- **Unwind info**
- **Relocation info**
- **Debug info**

Data Sources

```
reka@ubuntu:~/binary-tools$ bloaty -d symbols,sections ./dummy
```

1.8%	199	8.1%	199	[section .text]
1.7%	190	4.6%	114	do_stuff
34.7%	66	57.9%	66	.text
25.3%	48	0.0%	0	.symtab ←
16.8%	32	28.1%	32	.eh_frame
14.7%	28	0.0%	0	.strtab
4.2%	8	7.0%	8	.eh_frame_hdr
4.2%	8	7.0%	8	.rodata
1.4%	155	0.0%	0	[section .strtab]
0.9%	96	3.9%	96	[section .dynsym]
0.8%	92	3.7%	92	[section .eh_frame]
0.8%	90	2.5%	61	main
35.6%	32	52.5%	32	.eh_frame
26.7%	24	0.0%	0	.symtab ←
23.3%	21	34.4%	21	.text
8.9%	8	13.1%	8	.eh_frame_hdr
5.6%	5	0.0%	0	.strtab
0.7%	78	1.2%	30	dl_relocate_static_pie

The **symbol table** only refers to the machine code of the function

Other artifacts a function emits into the binary:

- **Unwind info**
- **Relocation info**
- **Debug info**
- **Its entry in the symbol table itself**

Data Sources

```
reka@ubuntu:~/binary-tools$ bloaty -d symbols,sections ./dummy
```

1.8%	199	8.1%	199	[section .text]
1.7%	190	4.6%	114	do_stuff
34.7%	66	57.9%	66	.text
25.3%	48	0.0%	0	.symtab
16.8%	32	28.1%	32	.eh_frame
14.7%	28	0.0%	0	.strtab
4.2%	8	7.0%	8	.eh_frame_hdr
4.2%	8	7.0%	8	.rodata
1.4%	155	0.0%	0	[section .strtab]
0.9%	96	3.9%	96	[section .dynsym]
0.8%	92	3.7%	92	[section .eh_frame]
0.8%	90	2.5%	61	main
35.6%	32	52.5%	32	.eh_frame
26.7%	24	0.0%	0	.symtab
23.3%	21	34.4%	21	.text
8.9%	8	13.1%	8	.eh_frame_hdr
5.6%	5	0.0%	0	.strtab
0.7%	78	1.2%	30	dl_relocate_static_pie

The **symbol table** only refers to the machine code of the function

Other artifacts a function emits into the binary:


- **Unwind info**
- **Relocation info**
- **Debug info**
- **Its entry in the symbol table itself**

Bloaty attributes these to the function to get a faithful total size


In practice, there is no 100% coverage
- **(Unmapped)**

BOLT

- Maksim Panchenko et al., 2018
- <https://github.com/facebookincubator/BOLT>
- <https://arxiv.org/abs/1807.06735>
- Binary layout optimization
- Supports ELF x86-64, AArch64, uses LLVM libs & perf



(a) without BOLT




(b) with BOLT

Heat maps for instruction memory accesses of the HHVM binary, without and with BOLT. Heat is in a log scale.

HHVM became **8% faster** on top of compiler opts + link-time function layout tool
Clang & GCC binaries **up to 20% faster** with PGO+LTO+BOLT than with PGO+LTO

PGO (Profile-Guided Optimization)

Also POGO or FDO (Feedback-Directed Optimization)



- Inserting profile data early means that many optimizations can benefit from them
- CPU and memory overhead

LTO (Link-Time Optimization)

Also LTCG (Link-Time Code Generation)

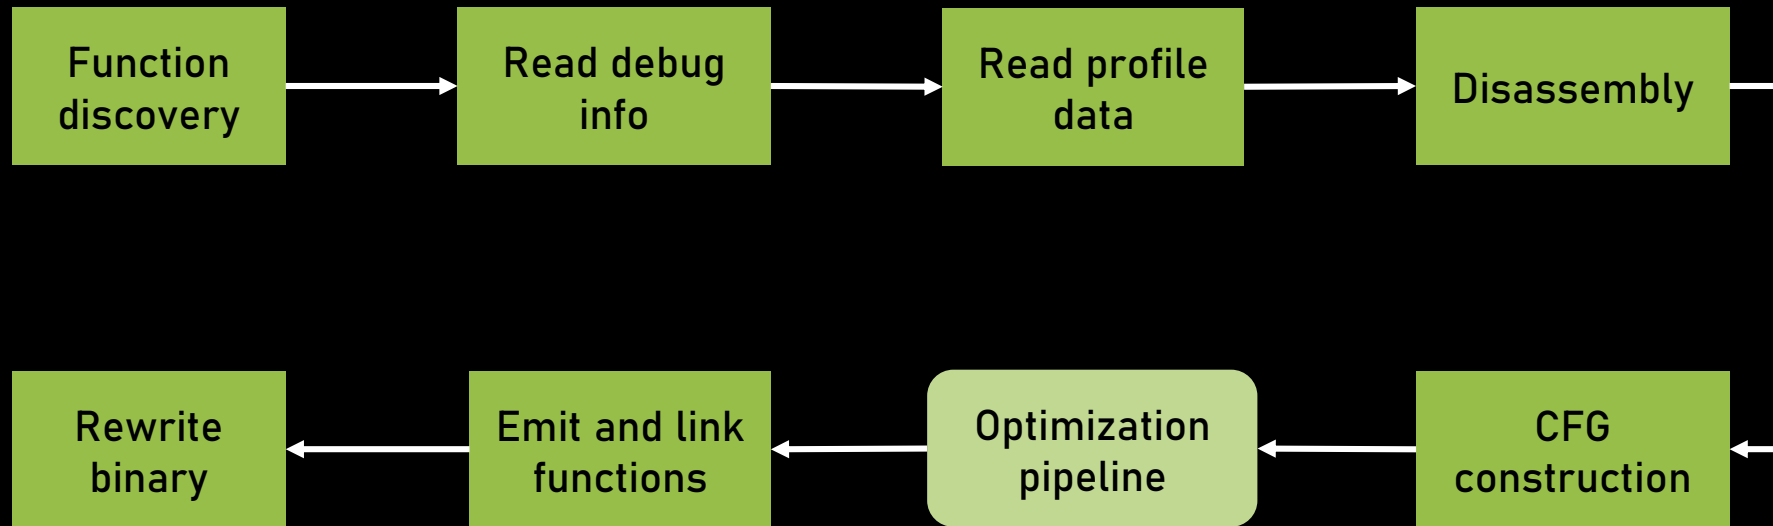
- Whole program analysis, cross-module optimization

BOLT

- Can optimize 3rd party libraries and assembly code
- Designed to work with the output of different compilers
- Uses sample-based profiling
 - Hardware profile counters e.g. Last Branch Records (Intel)
 - Negligible overhead
 - No special build required
- Binary-level profile data is applied on a binary level
 - No retrofitting needed, accurate

BOLT

Binary rewriting pipeline



BOLT

Optimization pipeline

Pass Name	Description
1. strip-rep-ret	Strip <code>repz</code> from <code>repz retq</code> instructions used for legacy AMD processors
2. icf	Identical code folding
3. icp	Indirect call promotion
4. peepholes	Simple peephole optimizations
5. inline-small	Inline small functions
6. simplify-ro-load	Fetch constant data in <code>.rodata</code> whose address is known statically and mutate a load into a <code>mov</code>
7. icf	Identical code folding (second run)
8. plt	Remove indirection from PLT calls
9. reorder-bbs	Reorder basic blocks and split hot/cold blocks into separate sections (layout optimization)
10. peepholes	Simple peephole optimizations (second run)
11. uce	Eliminate unreachable basic blocks
12. fixup-branches	Fix basic block terminator instructions to match the CFG and the current layout (redone by <code>reorder-bbs</code>)
13. reorder-functions	Apply HFSort [25] to reorder functions (layout optimization)
14. sctc	Simplify conditional tail calls
15. frame-opts	Removes unnecessary caller-saved register spilling
16. shrink-wrapping	Moves callee-saved register spills closer to where they are needed, if profiling data shows it is better to do so

BOLT

Optimization pipeline

Optimize functions compiled
without `-ffunction-sections`

3% code size reduction for
HHVM

Pass Name	Description
1. strip-rep-ret	Strip <code>repz</code> from <code>repz retq</code> instructions used for legacy AMD processors
2. icf	Identical code folding
3. icp	Indirect call promotion
4. peepholes	Simple peephole optimizations
5. inline-small	Inline small functions
6. simplify-ro-load	Fetch constant data in <code>.rodata</code> whose address is known statically and mutate a load into a <code>mov</code>
7. icf	Identical code folding (second run)
8. plt	Remove indirection from PLT calls
9. reorder-bbs	Reorder basic blocks and split hot/cold blocks into separate sections (layout optimization)
10. peepholes	Simple peephole optimizations (second run)
11. uce	Eliminate unreachable basic blocks
12. fixup-branches	Fix basic block terminator instructions to match the CFG and the current layout (redone by <code>reorder-bbs</code>)
13. reorder-functions	Apply HFSort [25] to reorder functions (layout optimization)
14. sctc	Simplify conditional tail calls
15. frame-opts	Removes unnecessary caller-saved register spilling
16. shrink-wrapping	Moves callee-saved register spills closer to where they are needed, if profiling data shows it is better to do so

BOLT

Optimization pipeline

Use **call frequency** information
to eliminate or change calls

Pass Name	Description
1. strip-rep-ret	Strip <code>repz</code> from <code>repz retq</code> instructions used for legacy AMD processors
2. icf	Identical code folding
3. icp	Indirect call promotion
4. peepholes	Simple peephole optimizations
5. inline-small	Inline small functions
6. simplify-ro-load	Fetch constant data in <code>.rodata</code> whose address is known statically and mutate a load into a <code>mov</code>
7. icf	Identical code folding (second run)
8. plt	Remove indirection from PLT calls
9. reorder-bbs	Reorder basic blocks and split hot/cold blocks into separate sections (layout optimization)
10. peepholes	Simple peephole optimizations (second run)
11. uce	Eliminate unreachable basic blocks
12. fixup-branches	Fix basic block terminator instructions to match the CFG and the current layout (redone by <code>reorder-bbs</code>)
13. reorder-functions	Apply HFSort [25] to reorder functions (layout optimization)
14. sctc	Simplify conditional tail calls
15. frame-opts	Removes unnecessary caller-saved register spilling
16. shrink-wrapping	Moves callee-saved register spills closer to where they are needed, if profiling data shows it is better to do so

BOLT

Optimization pipeline

Reduce the number of taken branches by making the hottest successor most likely a fall-through

Pass Name	Description
1. strip-rep-ret	Strip <code>repz</code> from <code>repz retq</code> instructions used for legacy AMD processors
2. icf	Identical code folding
3. icp	Indirect call promotion
4. peepholes	Simple peephole optimizations
5. inline-small	Inline small functions
6. simplify-ro-loads	Fetch constant data in <code>.rodata</code> whose address is known statically and mutate a load into a <code>mov</code>
7. icf	Identical code folding (second run)
8. plt	Remove indirection from PLT calls
9. reorder-bbs	Reorder basic blocks and split hot/cold blocks into separate sections (layout optimization)
10. peepholes	Simple peephole optimizations (second run)
11. uce	Eliminate unreachable basic blocks
12. fixup-branches	Fix basic block terminator instructions to match the CFG and the current layout (redone by <code>reorder-bbs</code>)
13. reorder-functions	Apply HFSort [25] to reorder functions (layout optimization)
14. sctc	Simplify conditional tail calls
15. frame-opts	Removes unnecessary caller-saved register spilling
16. shrink-wrapping	Moves callee-saved register spills closer to where they are needed, if profiling data shows it is better to do so

BOLT

Optimization pipeline

Uses a **weighted call graph** to do the reordering

Pass Name	Description
1. strip-rep-ret	Strip <code>repz</code> from <code>repz retq</code> instructions used for legacy AMD processors
2. icf	Identical code folding
3. icp	Indirect call promotion
4. peepholes	Simple peephole optimizations
5. inline-small	Inline small functions
6. simplify-ro-load	Fetch constant data in <code>.rodata</code> whose address is known statically and mutate a load into a <code>mov</code>
7. icf	Identical code folding (second run)
8. plt	Remove indirection from PLT calls
9. reorder-bbs	Reorder basic blocks and split hot/cold blocks into separate sections (layout optimization)
10. peepholes	Simple peephole optimizations (second run)
11. uce	Eliminate unreachable basic blocks
12. fixup-branches	Fix basic block terminator instructions to match the CFG and the current layout (redone by <code>reorder-bbs</code>)
13. reorder-functions	Apply HFSort [25] to reorder functions (layout optimization)
14. setc	Simplify conditional tail calls
15. frame-opts	Removes unnecessary caller-saved register spilling
16. shrink-wrapping	Moves callee-saved register spills closer to where they are needed, if profiling data shows it is better to do so

Workflow

- 0) Have a binary that does not fit into the instruction cache
- 1) Build a copy of it with `-Wl,-q` (or pass `--emit-relocs` to the linker)
- 2) Collect execution profiles with `perf`
- 3) Run `llvm-bolt` with the output of `perf`
- 4) Diff running times
- 5) See reduced number of instruction cache misses with `perf`

Clang binaries **up to 20% faster** with PGO+LTO+**BOLT** than with PGO+LTO

Metric	Over Baseline	Over PGO+LTO
executed forward branches	-1.6%	-1.0%
taken forward branches	-83.9%	-61.1%
executed backward branches	+9.6%	+6.0%
taken backward branches	-9.2%	-21.8%
executed unconditional branches	-66.6%	-36.3%
executed instructions	-1.2%	-0.7%
total branches	-7.3%	-2.2%
taken branches	-69.8%	-44.3%
non-taken conditional branches	+60.0%	+13.7%
taken conditional branches	-70.6%	-46.6%

Stats reported by **BOLT** when applied to Clang's baseline and PGO+LTO binaries

Clang binaries **up to 20% faster** with PGO+LTO+BOLT than with PGO+LTO

Metric	Over Baseline	Over PGO+LTO
executed forward branches	-1.6%	-1.0%
taken forward branches	-83.9%	-61.1%
executed backward branches	+9.6%	+6.0%
taken backward branches	-9.2%	-21.8%
executed unconditional branches	-66.6%	-36.3%
executed instructions	-1.2%	-0.7%
total branches	-7.3%	-2.2%
taken branches	-69.8%	-44.3%
non-taken conditional branches	+60.0%	+13.7%
taken conditional branches	-70.6%	-46.6%

Stats reported by BOLT when applied to Clang's baseline and PGO+LTO binaries

Thanks!