

# The C++20 Standard Library - Beyond Ranges

Jeff Garland

Created: 2019-09-18 Wed 06:36

intro

# talk goals

- io – formatted output [`~15 min`]
  - `output format {fmt}`
  - `sync_buf` and `osyncstream`
  - efficient `stringbuf` and `stringstream` access
- container and algorithm updates [`~12 min`]
  - `span`
  - associative containers `contains`
  - uniform container erase algorithms
  - `to_array` helper (unlikely)
  - `erase` returns on `list`
  - `shift` algorithms (unlikely)

# talk goals II

- chrono [~8 min]
- concurrency: threading and atomics [~10 min]
  - `jthread` and `stop_token`
  - `atomic` for `shared_ptr` ~`weak_ptr`
  - `atomic_ref<T>`

# talk goals III (really fast)

- misc [~5 min]
  - `string starts_with ends_with`
  - `shared_ptr, unique_ptr` initialization
  - `math`
  - `[ [nodiscard] ]` on functions
- bit manipulation [~5 min]
  - rotation and count
  - power functions
  - `bit_cast`
  - endian

# c++20 library changes not discussed today

- ranges is the big one
- concept support for ranges
- 3way compare support (aka spaceship support)
- concurrency
  - coordination types: `latch` and `barrier`
  - `counting_semaphore` and `binary_semaphore`
- `char8_t`, `u8string` and `u8string_view`

# c++20 library changes not discussed II

- `constexpr` on algorithms and containers
- various `noexcept` additions
- `variant` changes
  - sane `variant` converting constructor  
<http://wg21.link/P0608>
  - `visit<R>` explicit return type for `visit`  
<http://wg21.link/P0655>
- algorithm vectorization policies  
<http://wg21.link/P1001>

# things I will and won't do in talk

- **will:** take questions as we go - until we get behind
- **will:** defer questions I can't answer immediately
- **will:** no doubt get something wrong



# things I will and won't do in talk

- will: show you lots of code
- will: shorten namespaces (`std::ranges`) and leave out `#includes` on slides
- wont: show you an example that hasn't compiled
  - in some universe with some set of libraries/compilers (unless noted)
  - the environment: Linux g++ 8.2, 8.3, 9
    - v9 - trunk as of mid-Mar 2019
    - typically with `-fconcepts` and `-stdc++20`
    - various referenced library implementations

# status of c++20 implementations

*“Patience you must have my young padawan”*

- *Yoda*

# c++20 Library implementation status

- c++20 is feature complete
- g++, clang, Visual C++

[https://en.cppreference.com/w/cpp/compiler\\_support](https://en.cppreference.com/w/cpp/compiler_support)

- current working paper <http://wg21.link/n4830>

# input-output

- output format `{fmt}`
- `syncbuf` and `ostream`
- efficient `stringbuf` and `stringstream` access

# output string format `std::format`

- python style formatting to c++
- positional parameters
- custom type formats
- faster and less overhead then iostreams
- integrated with chrono (output only)
- in header `<format>`

# hello world

```
#include <fmt/core.h> //will be <format>

{
    //hello world
    string s = fmt::format("{} there {}\n", "hello", "world");
    cout << s;    //hello there world
}
```

- format function replaces the {} with corresponding parameters
- parameters can be any type
  - all built in numbers and strings support out of box
  - custom types can be supported

# indexed parameters

```
{  
    //indexed parameters -- parm[0] == 42, parm[1] == "hello"  
    //format reorders in string  
    string s = fmt::format("{1} and {0}\n", 42, "hello");  
    cout << s;    //hello and 42  
}
```

- cannot mix indexed and non-indexed

```
string s = fmt::format("{} and {0}\n", 42, "hello"); //error
```

# simple escaping

```
{  
    //escaping '{' and '  
    string s = fmt::format("{ { untouched } } hello {} \n", 42);  
    cout << s;    //{ untouched } hello 42  
}
```



# crazy town

```
{
    //escaping '{' and '}'
    string s = fmt::format("{ { in brace {} } }\n", 42);
    cout << s;    //{ in brace 42 }
}
{
    //braceapoloza
    string s = fmt::format("{{{{{{{{}}}}}}\n", 42);
    cout << s;    //{{42}}
}
```

# floating point format

```
{  
    int width = 10;  
    int precision = 3;  
    // format like {0:10.3f}  
    auto s = fmt::format("{0:{1}.{2}f}", 12.345678, width, precision);  
    // s == "      12.346"  
}
```

# rich language for fill, alignment, conversion

```
string s0 = format("{:6}", 42);    // s0 is "    42"
string s1 = format("{:6}", 'x');    // s1 is "x      "
string s2 = format("{:*<6}", 'x'); // s2 is "x*****"
string s3 = format("{:*>6}", 'x'); // s3 is "*****x"
string s4 = format("{:*^6}", 'x'); // s4 is "**x**"
string s5 = format("{:6}", true);  // s5 is "true  "

string s6 = format("{0:b} {0:d} {0:o}", 42); // s6 is "101010 42 52"
string s7 = format("{0:#x} {0:#X}", 42);    // s7 is "0x2a 0X2A"
```

# custom type usage

```
//annotated version  
enum class my_color {red, green, blue, orange};  
  
my_color c = my_color::red;  
string s = fmt::format("{}\n", c); //please work!
```

# how to extend for custom type

```
//annotated version
enum class my_color {red, green, blue, orange};

namespace fmt { // <- template specialization in fmt namespace
    template <> // <- syntax for specialization
    struct formatter<my_color>: formatter<string_view> { //inherit from f
        template <typename FormatContext>
        auto format(my_color c, FormatContext &ctx) {
            string_view name = "unknown color"; // <- some default text
            switch (c) {
                case my_color::red:    name = "red"; break;
                case my_color::green:  name = "green"; break;
                case my_color::blue:   name = "blue"; break;
            } //should there be a fallthru instead?
            return formatter<string_view>::format(name, ctx);
        }
    }; //struct formatter<my_color>
} //namespace fmt
```

## fmt diagnostics/safety

- currently runtime diagnostic (`std::exception` derivative)
- cases should be checkable in future
- what causes exceptions?
  - wrong parameter count
  - mismatched parameter types
- advice
  - dont use in catch blocks
  - dont use in destructors
  - dont use in non-exception friendly locations

## mis-matched parameter count

```
//2 format parameters only one actual parameter - throws  
string s = fmt::format("{1} hello {0}\n", "world");
```

# wrong format typing causes error, even when convertible

```
{
//no automatic conversion like python
//print("Sammy ate {0:f} percent of a {1}!".format(75, "pizza"))
//Sammy ate 75.000000 percent of a pizza!

    string s = fmt::format("float and decimal: {1:f} and {0:d}\n",
                           42, 42);

    cout << s;

//terminate called after throwing an instance of 'fmt::v5::format_error'
//what():  invalid type specifier
//formatting numbers
}
```

- <https://www.digitalocean.com/community/tutorials/how-to-use-string-formatters-in-python-3>



## `fmt` references and status

- moved by LWG in Cologne
- <http://wg21.link/p0645>
- <https://github.com/fmtlib/fmt>

# syncbuf and ostream

## overview

- streams in multi-threaded application can produce garbled output
- ostream used as an automatic-duration variable
- buffers output operations for a wrapped stream
- transfers sync buffer stream to underlying on destruction
- header <ostream>

# osyncstream simple example

```
#include <syncstream> //new header
//...
{
    osyncstream buffered_out(cout);
    buffered_output << " the answer might be 42" << endl;
} //on destruction buffered_out now calls emit()
//emit() flushes the buffer to cout
{
    std::ofstream out_file("my_file");
    osyncstream buffered_out(out_file);
    std::emit_on_flush(buffered_out);
    buffered_out << "hello world" << endl; //calls emit()
}
```

# osyncstream thread example

```
void do_output(std::ostream& ofile ) {  
    osyncstream buf_file{ofile};  
    // ...complex code that writes to file and may throw  
} //ensures buf_file is flushed to outfile  
  
std::ofstream out_file("my_file");  
std::thread    out_thread1( &do_output, out_file )  
std::thread    out_thread2( &do_output, out_file )
```

# syncbuf and ostream synopsis

```
template <class charT,  
         class traits = char_traits<charT>,  
         class Allocator = allocator<charT>>  
class basic_syncbuf;  
  
using syncbuf = basic_syncbuf<char>;  
using wsyncbuf = basic_syncbuf<wchar_t>;  
  
template <class charT,  
         class traits = char_traits<charT>,  
         class Allocator = allocator<charT> >  
class basic_ostream;  
  
using ostream = basic_ostream<char>;  
using wostream = basic_ostream<wchar_t>;
```

## `syncbuf` and `osyncstream` references

- in working draft
- paper <http://wg21.link/p0053>
- [https://github.com/PeterSommerlad/SC22WG21\\_Pap](https://github.com/PeterSommerlad/SC22WG21_Pap)

# stringstream and basic\_stringbuf updates

- `basic_stringbuf` used in `stringstream`
- changes to allow for control of buffer allocator
- output interfaces to provide `string_view` from `stringstream`

# stringstream to string\_view example

```
#include <sstream>
//...
std::stringstream s;
s << "put some data into the stream";

std::string_view sv = s.view();
```



# stringstream and basic\_stringbuf resources

- <http://wg21.link/p0408>
- implementation  
[https://github.com/PeterSommerlad/SC22WG21\\_Pap](https://github.com/PeterSommerlad/SC22WG21_Pap)

# container and algorithm updates

- `span`
- `associative containers contains`
- `uniform container erase algorithms`
- `to_array` helper
- `erase` returns on `list`
- `shift` algorithms

# span

- span is a non-owning 'view' over contiguous sequence of object
- cheap to copy - implementation is a pointer and size
- constant time complexity for all member functions
- defined in header `<span>`
- unlike most 'view types' can mutate
- constexpr ready

# span construction

```
vector<int> vi = { 1, 2, 3, 4, 5 };  
span<int>    si ( vi );  
  
array<int, 5> ai = { 1, 2, 3, 4, 5 };  
span<int>    si2 ( ai );  
  
int cai[] = { 1, 2, 3, 4, 5 };  
span<int>    si3( cai );  
  
array<string, 2> sa = { "world" , "hello "};  
span<string> ss ( sa );
```

# span as a function parameter

```
void print_reverse(span<int> si) { //by value
    for ( auto i : std::ranges::reverse_view{si} ) {
        cout << i << " ";
    }
    cout << "\n";
}

int main () {

    vector<int> vi = { 1, 2, 3, 4, 5 };
    print_reverse( vi ); //5 4 3 2 1

    int cai[] = { 1, 2, 3, 4, 5 };
    print_reverse ( cai ); //5 4 3 2 1

    span<int> si ( vi );
    print_reverse( si.first(2) ); //2 1
    print_reverse( si.last(2) );  //5 4
}
```

## basic public accessors

- supports collection-like interfaces (`begin`, `end`, `operator[ ]`)
- `data` pointer to start of sequence
- `size` number of elements
- `size_bytes` memory size
- `empty` true if no elements

## static and dynamic extent

```
array<int, 5> ai = { 1, 2, 3, 4, 5 };  
span<int, 5> si5( ai );  
print_reverse ( si5 ); //5 4 3 2 1  
  
span<int, 5> si3( ai );
```

## span resources

- in the working paper
- <http://wg21.link/P0122>
- implementation <https://github.com/tcbrindle/span>
- should span be regular <http://wg21.link/p1085>



# contains method for associative containers

```
std::map<string, SomeType> m = initMap();  
if (m.find( "foo" ) != m.end()) {  
//or  
if (m.contains( "foo" )) { //uh huh -- easier to read
```

# contains method signatures

```
//adds member contains (2 overloads) to map, multimap,  
//set, and multiset  
  
//check with the key type  
bool contains(const key_type& x) const;  
  
//check with a type that is comparable to key_type  
template <class K> bool contains(const K& x) const;
```

# contains method signatures for unordered

```
// adds member contains to unordered_map, unordered_multimap,  
// unordered_set, and unordered_multiset  
  
//check with the key type  
bool contains(const key_type& x) const;
```

contains reference

- paper: <http://wg21.link/p0458>
- status: in working draft

# uniform container erasure

- motivation - removing elements unreasonably difficult
- when there's an idiom, that's a problem
- simplify user code
- Erase idiom

```
std::vector<SomeType> c = ...;  
c.erase(remove_if(c.begin(), c.end(), pred), c.end());
```

# erase and erase\_if signatures

```
template<class T, class Allocator, class U>  
void erase(list<T, Allocator>& c, const U& value);
```

Effects: Equivalent to:

```
erase_if(c, [&](auto& elem) return elem == value; );
```

```
template<class T, class Allocator, class Predicate>  
void erase_if(list<T, Allocator>& c, Predicate pred);
```

Effects: Equivalent to: `c.remove_if(pred);`

## erase and erase\_if scope

- both apply to `basic_string`, `deque`, `forward_list`, `list`, `vector`
- `erase_if` applies to `map`, `multimap`, `set`, `multiset`, `unordered_*`
- since associative containers already have `erase (keytype )` functions

## uniform container erase status and references

- merged in working draft
- starting paper <http://wg21.link/n4009>
- San Diego final <http://wg21.link/p1209>
- lifted from library fundamentals v2
- in vcc since 2015



`to_array` helper

## motivating example

```
// array<char, 4> from array{ "foo" } ? no  
// because: array{ "foo", "bar" } creates array<char const*, 2>  
  
auto a1 = to_array("foo");           // array<char, 4>  
auto a2 = array{ "foo" };             // array<char const*, 1>
```

- from the LTFsV2 with updates

to\_array references

- <http://wg21.link/P0325.html>
- status: LWG moved in Cologne

# return values on `remove`

- impacts `list` and `forwardlist`
- why throw away information?
- change `remove()`, `remove_if()` and `unique()`
- return `container::size_type` of number elements removed
- makes consistent with `map/set erase` function

return values status and reference

- in working draft
- <http://wg21.link/p0646>

# `shift_left` and `shift_right` algorithms

- move elements a fixed amount
- return an iterator to one past shift point
- `size` of sequence unchanged

# shift\_left example

```
vector<int>    vi = { 1, 2, 3, 4, 5 };

auto itr = shift_left(vi.begin(), vi.end(), 2);
//vi = 3 4 5 4 5
//      ^
//      itr
vi.erase( itr, vi.end());

for ( auto i : vi ) {
    cout << i << " ";
}
cout << "\n"; //3 4 5
```

# shift synopsis

```
// 25.6.14, shift
template<class ForwardIterator>
constexpr ForwardIterator
shift_left(ForwardIterator first, ForwardIterator last,
           typename iterator_traits<ForwardIterator>::difference_type n)

template<class ForwardIterator>
constexpr ForwardIterator
shift_right(ForwardIterator first, ForwardIterator last,
            typename iterator_traits<ForwardIterator>::difference_type n)

//also overloads for parallel execution
```



## shift reference

- merged in working draft
- <http://wg21.link/p0769>
- [https://github.com/danra/shift\\_proposal](https://github.com/danra/shift_proposal)
- will be ranges versions (not merged in c++20)

# chrono

- chrono fundamentals
- date support features
- i/o

# chrono design fundamentals

- extends existing capabilities to calendric functions
- two primary types in chrono
  - structure/field types like `year_month_day`
  - calculation types like `sys_days`, `time_point`, `seconds`
- field types are used more for i/o and conversions
- calculation types - small memory, fast computations
- supports `constexpr`

# simple calculations

```
year_month_day ymd(year(2019)/05/07);  
sys_days d( ymd );  
d += weeks(1);  
cout << ymd << "/n"; //2019-05-07  
cout << d << "/n";   //2019-05-14
```

## chrono error handling

- no exceptions
- the 'ok' method tells you if a value is good

```
year_month_day ymd(year(2019)/100/200);  
if ( !ymd.ok() ) {  
    cout << "bad ymd\n"; //executes  
}
```

# chrono i/o

- stream-based input/output
- `std::format` based output for strings
- uses `sprintf` style format strings

# chrono format output

```
//cant compile this...  
const std::chrono::year_month_day ymd(2019_y/05/07);  
std::string s = std::format("the date is {:%m-%d-%Y}.\n", ymd);  
cout << s << "\n"; //
```

## optional locale aware format output

```
//not compiled  
auto zt = std::chrono::zoned_time(...);  
std::cout << std::format(std::locale{"fi_FI"},  
                          "Localized time is {:%c}\n", zt);
```



# other c++20 chrono features

- leap seconds
- IANA timezone database support
- localized times
- many other types `year_month`, `gps_clock`, `weekday_indexed`

# chrono references and status

- reference implementation  
<https://github.com/HowardHinnant/date>
- chrono for c++20 <http://wg21.link/p0355> (in working paper)
- std::format and chrono <http://wg21.link/p1361> (LWG moved in Cologne)
- chrono format strings  
<https://en.cppreference.com/w/cpp/chrono/format>

# concurrency: threading and atomics

- `jthread` and `stop_token`
- `atomic<shared_ptr>`, `atomic<weak_ptr>`
- `atomic_ref<T>`

# jthread and stop\_token

- jthread (joining thread) is an update for `std::thread`
- automatically joins in the destructor
- stop\_token provides cooperative shutdown
- in header `<thread>`

# thread gone wrong?

```
#include <thread>

void do_nothing()
{
    this_thread::sleep_for( seconds( 1 ) );
}

int
main()
{
    cout << "start" << endl;
    {

        std::thread t(do_nothing);

    } // what bad thing happens here?
}
```

# jthread instead

```
#include <thread>

{

    std::jthread t(do_nothing);

} //I'm cool if you're cool


// fine, because...
~jthread()
{
    if( joinable() ) {
        request_stop();
        join();
    }
}
```

- jthread defaults to joining

# jthread cooperative shutdown

```
class jthread {  
public:  
    [[nodiscard]] stop_source get_stop_source() noexcept;  
    [[nodiscard]] stop_token get_stop_token() const noexcept;  
    bool request_stop() noexcept;  
    ....  
};
```

- 3 types stop\_source, stop\_token, stop\_callback

# jthread cooperative shutdown

```
// user code -- (not compiled)
{
    std::jthread t([]( std::stop_token token ){
        while ( !token.requested_stop() )
        {
            //do stuff
        }
    });
}
```



# new waiting interfaces with stop\_token

```
// 32.6.4.2 interruptible waits:
```

```
template<class Lock, class Predicate>
```

```
bool
```

```
wait_until(Lock& lock, Predicate pred, stop_token token);
```

```
template<class Lock, class Clock, class Duration, class Predicate>
```

```
bool
```

```
wait_until(Lock& lock,  
            const chrono::time_point<Clock, Duration>& abs_time,  
            Predicate pred,  
            stop_token token);
```

```
template<class Lock, class Rep, class Period, class Predicate>
```

```
bool
```

```
wait_for(Lock& lock,  
          const chrono::duration<Rep, Period>& rel_time,  
          Predicate pred,  
          stop_token token)
```

## jthread references and status

- motivation <http://wg21.link/p0660r0>
- wording <http://wg21.link/p0660>
- implementation <https://github.com/josuttis/jthread>

# atomic shared\_ptr and weak\_ptr

- atomic\_shared\_ptr interfaces are error prone
- removes atomic\_shared\_ptr<T> and atomic\_weak\_ptr<T> alias templates
- adds new atomic interface

```
// 23.11.3, atomic smart pointers
template <class T> struct atomic<shared_ptr<T>>;
template <class T> struct atomic<weak_ptr<T>>;
```

- reference <http://wg21.link/p0718>

# atomic\_ref

- High performance computing motivates
- like `atomic<T>` but for reference to non-atomic
- specializations for integral types (char, int, long, etc)
- specializations for floating point types
- `atomic_ref` <http://wg21.link/P0019>

## atomic\_ref example

```
vector<int> vi { 1, 2, 3, 4 }; //not atomics

for ( int i=0; i != vi.end(); i++) {
    value = atomic_ref<int>{ vi[i] };
    value += 1; //atomic fetch_add
}
```

# misc

- `shared_ptr`, `unique_ptr` default initialization
- `string starts_with ends_with`
- `source_location`
- `math`
- `[ [nodiscard] ]` on functions

# `shared_ptr`, `unique_ptr` default initialization

- All about performance
- arrays of built in types often initialized after creation
- hence initialization performed by `make_shared` and `make_unique` redundant

## default pointer init example

```
//shared_ptr to a default-initialized double[1024]  
//where each element has an indeterminate value  
shared_ptr<double[]> p = make_shared_default_init<double[]>(1024);
```



default initialization for `unique_ptr` and  
`shared_ptr` reference

- Smart pointer creation with default initialization  
<http://wg21.link/p1020>
- [http://boost.org/libs/smart\\_ptr](http://boost.org/libs/smart_ptr) [since 1.56]

# starts\_with and ends\_with for string and string\_view

```
//string_view
constexpr bool starts_with(basic_string_view x) const noexcept;
constexpr bool starts_with(charT x) const noexcept;
constexpr bool starts_with(const charT* x) const;
constexpr bool ends_with(basic_string_view x) const noexcept;
constexpr bool ends_with(charT x) const noexcept;
constexpr bool ends_with(const charT* x) const;
```

- prefix and suffix checks
- <http://wg21.link/p0457>

# source\_location

- modern replacement for FILE and LINE C macros
- more sophisticated abilities
- new header `<source_location>`
- originally planned for use in contracts and stacktrace

# source\_location synopsis

```
//[source_location.syn]
struct source_location {

    // source location construction
    static consteval source_location current() noexcept;
    constexpr source_location() noexcept;

    // source location field access
    constexpr uint_least32_t line() const noexcept;
    constexpr uint_least32_t column() const noexcept;
    constexpr const char* file_name() const noexcept;
    constexpr const char* function_name() const noexcept;
```

# source\_location usage

```
void f(source_location a = source_location::current()) {  
    // values in b refer to the line below  
    source_location b = source_location::current();  
    // ...  
}  
  
// f's first argument corresponds to this line of code  
f();  
  
// f's first argument gets the same values as c, above  
source_location c = source_location::current();  
f(c);
```

# source\_location usage

```
struct s {  
    source_location construct_loc = source_location::current();  
    int other_member;  
  
    // values of construct_loc refer to the location of the calling function  
    s(source_location loc = source_location::current())  
        : construct_loc(loc)  
    {}  
  
    // values of construct_loc refer to this location  
    s(int i) :  
        other_member( i )  
    {}  
  
    // values of construct_loc refer to this location  
    s(double)  
    {}  
};
```

## source\_location resources

- <http://wg21.link/p1208>
- source location implementations - not aware of one

# math constants

- new header `<math>`
- new namespace `std::math`
- provided for various floating point types
- nearest representable value



# math constants example

```
inline constexpr long double e1 = e_v<long double>;
inline constexpr long double log2e1 = log2e_v<long double>;
inline constexpr long double log10e1 = log10e_v<long double>;
inline constexpr long double pil = pi_v<long double>;
inline constexpr long double inv_pil = inv_pi_v<long double>;
inline constexpr long double inv_sqrtpil = inv_sqrtpi_v<long double>;
inline constexpr long double ln21 = ln2_v<long double>;
inline constexpr long double ln101 = ln10_v<long double>;
inline constexpr long double sqrt21 = sqrt2_v<long double>;
inline constexpr long double sqrt31 = sqrt3_v<long double>;
inline constexpr long double inv_sqrt31 = inv_sqrt3_v<long double>;
inline constexpr long double egamma1 = egamma_v<long double>;
inline constexpr long double phil = phi_v<long double>;
```

## math constants reference and status

- pre-reviewed for Cologne
- paper <http://wg21.link/p0631>
- boost <http://boost.org/libs/math>

# midpoint

- why is midpoint difficult?
- for integers, overflow is the primary problem
- for floating point it's rounding
- same issue for pointer type

```
template<typename A>  
constexpr A midpoint(A a, A b) noexcept;
```

# linear interpolation (lerp)

- handy for many domains
- why is linear interpolation difficult?
- really easy to get it wrong for non-mathematicians

```
//from P0811 -- wanted properties
exactness: lerp(a,b,0)==a && lerp(a,b,1)==b
monotonicity: cmp(lerp(a,b,t2),
                  lerp(a,b,t1)) * cmp(t2,t1) * cmp(b,a) >= 0,
              where cmp is an arithmetic three-way comparison function
determinacy: result of NaN only for lerp(a,a,INFINITY)
boundedness: t<0 || t>1 || isfinite(lerp(a,b,t))
consistency: lerp(a,a,t)==a
```

- interface

```
constexpr float lerp(float a, float b, float t);
constexpr double lerp(double a, double b, double t);
constexpr long double lerp(long double a, long double b, long double t);
```

## midpoint/lerp references and status

- midpoint and lerp are in p0811
- <http://wg21.link/p0811>
- in working draft

c++20 will addnodiscard to many  
functions

- including `empty`, `allocate`, `new`
- <http://wg21.link/p0600r1>

# bit manipulation functions

- rotation and count
- power functions
- `bit_cast`
- endian

# header `<bit>` overview

- new in c++20
- function templates to access, manipulate, and process individual bits and bit sequences.



# power functions

- Defined in header `<bit>` in namespace `std`
- `ispow2` checks if a number is an integral power of two
- `ceil2` finds the smallest integral power of two not less than the given value
- `floor2` finds the largest integral power of two not greater than the given value
- `log2p1` finds the smallest number of bits needed to represent the given value

# bit rotation and counting

```
//Header bit
// 23.20.2, rotating
template<class T>
    [[nodiscard]] constexpr T rotl(T x, int s) noexcept;
template<class T>
    [[nodiscard]] constexpr T rotr(T x, int s) noexcept;

// 23.20.3, counting
template<class T>
    constexpr int countl_zero(T x) noexcept;
template<class T>
    constexpr int countl_one(T x) noexcept;
template<class T>
    constexpr int countr_zero(T x) noexcept;
template<class T>
    constexpr int countr_one(T x) noexcept;
template<class T>
    constexpr int popcount(T x) noexcept;
}
```

# bit manipulation status and reference

- low level routines
- expect instruction level optimizations
- reference <http://wg21.link/p0553>
- earlier paper <http://wg21.link/n3864>

# bit\_cast

- low level facility to cast from type to type
- `reinterpret_cast` and `union` can be error prone
- size of the two types must match

```
// 25.5.3, bit_cast
//reinterpret the object representation of one type as that of another
template<typename To, typename From>
constexpr To bit_cast(const From& from) noexcept;
```

- <http://wg21.link/p0476>
- in working draft

# endian

```
//short description: detect the byte order at compile or runtime  
//header: <bit>
```

```
enum class endian {  
    little = /*implementation-defined*/,  
    big    = /*implementation-defined*/,  
    native = /*implementation-defined*/  
};
```

# compile time usage

```
class my_byte_dependent_type
{
public:
    static constexpr std::endian endian = std::endian::native;
```

# runtime usage

```
if (endian::native == endian::big)
    // do big endian encoding
else if (endian::native == endian::little)
    // handle little endian
else
    // handle mixed endian (mythical unicorn)
```

## endian status and reference

- merged in working draft
- Howard H. <http://wg21.link/p0463>
- boost endian <http://www.boost.org/libs/endian>
- other bit manipulation facilities
  - `std::bitset` (since 1998)
  - boost dynamic bitset  
[http://www.boost.org/libs/dynamic\\_bitset](http://www.boost.org/libs/dynamic_bitset)



# observations

- library updates for c++20 are big
- some key building blocks
- many small additions
- overall: nicer cleaner more capable code

# Thanks!!

*"Design is nothing if not decision making."*

*Henri Petroski 2007 - "Small Things  
Considered: Why There Is No Perfect Design"*