

The C++20 Synchronization Library

Bryce Adelstein Lelbach

```
unique_future<std::uint64_t>
fibonacci(boost_blockable_submitter auto&& s, std::uint64_t n) {
    if (n < 2) co_return n;

    auto n1 = async(s, fibonacci<decltype(s)>, s, n - 1);
    auto n2 = fibonacci(s, n - 2);

    co_return co_await n1 + co_await n2;
}
```



THE C++20 SYNCHRONIZATION LIBRARY

Bryce Adelstein Lelbach

CUDA C++ Core Libraries Lead



@blelbach

ISO C++ Library Evolution Incubator Chair, ISO C++ Tooling Study Group Chair

```
namespace stdr = std::ranges;  
namespace stdv = std::views;
```

Recipe For a Tasking Runtime

- ▶ Worker threads.
- ▶ Multi-consumer, multi-producer concurrent queue.
- ▶ Termination detection mechanism.
- ▶ Parallel algorithms.

Recipe For a Tasking Runtime

- ▶ **Worker threads.**
- ▶ Multi-consumer, multi-producer concurrent queue.
- ▶ Termination detection mechanism.
- ▶ Parallel algorithms.

```
struct thread_group {  
private:  
    std::vector<std::thread> members;  
  
public:  
    thread_group(std::uint64_t n, std::invocable auto&& f) {  
        for (auto i : std::iota(0, n)) members.emplace_back(f);  
    }  
};
```

```
struct thread_group {  
private:  
    std::vector<std::thread> members;  
  
public:  
    thread_group(std::uint64_t n, std::invocable auto&& f) {  
        for (auto i : std::iota(0, n)) members.emplace_back(f);  
    }  
};  
  
int main() {  
    std::atomic<std::uint64_t> count(0);  
  
    {  
        thread_group tg(6, [&] { ++count; });  
    }  
  
    std::cout << count << "\n";  
}
```

```
struct thread_group {  
private:  
    std::vector<std::thread> members;  
  
public:  
    thread_group(std::uint64_t n, std::invocable auto&& f) {  
        for (auto i : std::iota(0, n)) members.emplace_back(f);  
    }  
  
    ~thread_group() {  
        std::for_each(members, [] (std::thread& t) { t.join(); });  
    }  
};
```



```
int main() {  
    std::atomic<std::uint64_t> count(0);  
  
    {  
  
        thread_group tg(6,  
            [&] {  
                while (true)  
                    ++count;  
            }  
        );  
  
    }  
  
    std::cout << count << "\n";  
}
```

```
int main() {
    std::atomic<std::uint64_t> count(0);

    {
        std::atomic<bool> done(false);

        thread_group tg(6,
            [&] {
                while (!done.load(std::memory_order_relaxed))
                    ++count;
            }
        );

        done.store(true, std::memory_order_relaxed);
    }

    std::cout << count << "\n";
}
```

```
struct thread_group {  
private:  
    std::vector<std::jthread> members;  
  
public:  
    thread_group(std::uint64_t n, std::invocable<std::stop_token> auto&& f) {  
        for (auto i : std::iota(0, n)) members.emplace_back(f);  
    }  
};
```

`std::jthread`

- ▶ Just like `std::thread`, except:
- ▶ If the thread is joinable, it joins with it instead of calling `terminate`.

```
struct thread_group {  
private:  
    std::vector<std::jthread> members;  
  
public:  
    thread_group(std::uint64_t n, std::invocable<std::stop_token> auto&& f) {  
        for (auto i : std::iota(0, n)) members.emplace_back(f);  
    }  
  
    auto size() { return members.size(); }  
  
    void request_stop() {  
        std::for_each(members, []) { t.request_stop(); };  
    }  
};
```

std::jthread

- ▶ Just like `std::thread`, except:
- ▶ If the thread is joinable, it joins with it instead of calling `terminate`.
- ▶ It supports interruption.
 - ▶ `std::thread` invocable parameter: takes no arguments.
 - ▶ `std::jthread` invocable parameter: takes either no arguments, or a `std::stop_token`.

- ▶ Interruption API:

```
[[nodiscard]] stop_source std::jthread::get_stop_source() noexcept;  
[[nodiscard]] stop_token std::jthread::get_stop_token() const noexcept;  
bool std::jthread::request_stop() noexcept;
```

`std::stop_*`

- ▶ `std::stop_source` (analogous to a promise)
 - ▶ Producer of stop requests.
 - ▶ Owns the shared state (if any).
- ▶ `std::stop_token` (analogous to future)
 - ▶ Handle to a `std::stop_source`.
 - ▶ Consumer only; can query for stop requests, but can't make them.
- ▶ `std::stop_callback` (analogous to `future::then`)
 - ▶ Mechanism for registering invocables to be run upon receiving a stop request.

CV Interruption Support

```
struct condition_variable_any {  
    template <typename Lock, typename Predicate>  
        bool wait_until(Lock& lock, Predicate pred, stop_token token);  
    template <typename Lock, class Clock, typename Duration, typename Predicate>  
        bool wait_until(Lock& lock, const chrono::time_point<Clock, Duration>& abs_time  
                        Predicate pred, stop_token token);  
    template <typename Lock, typename Rep, typename Period, typename Predicate>  
        bool wait_for(Lock& lock, const chrono::duration<Rep, Period>& rel_time,  
                     Predicate pred, stop_token token);  
};
```



```
int main() {  
    std::atomic<std::uint64_t> count(0);  
  
    {  
        thread_group tg(6,  
            [&] (std::stop_token s) {  
                while (!s.stop_requested())  
                    ++count;  
            }  
        );  
    }  
  
    std::cout << count << "\n";  
}
```

Recipe For a Tasking Runtime

- ▶ Worker threads.
- ▶ **Multi-consumer, multi-producer concurrent queue.**
- ▶ Termination detection mechanism.
- ▶ Parallel algorithms.

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

    T pop();

public:
    constexpr concurrent_bounded_queue() = default;

    void enqueue(std::convertible_to<T> auto&& u);

    T dequeue();
    std::optional<T> try_dequeue();
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

    T pop();

public:
    constexpr concurrent_bounded_queue() = default;

    void enqueue(std::convertible_to<T> auto&& u);

    T dequeue();
    std::optional<T> try_dequeue();
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

    T pop();

public:
    constexpr concurrent_bounded_queue() = default;

    void enqueue(std::convertible_to<T> auto&& u);

    T dequeue();
    std::optional<T> try_dequeue();
};

```

```
template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

    T pop();

public:
    constexpr concurrent_bounded_queue() = default;

    void enqueue(std::convertible_to<T> auto&& u);

    T dequeue();
    std::optional<T> try_dequeue();
};
```

std::counting_semaphore

```
template <ptrdiff_t least_max_value = implementation-defined>
struct counting_semaphore {
    static constexpr ptrdiff_t max() noexcept;

    constexpr explicit counting_semaphore(ptrdiff_t desired);

    void release(ptrdiff_t update = 1);

    void acquire();
    bool try_acquire() noexcept;
    template <typename Rep, typename Period>
        bool try_acquire_for(const chrono::duration<Rep, Period>& rel_time);
    template <typename Clock, typename Duration>
        bool try_acquire_until(const chrono::time_point<Clock, Duration>& abs_time);
};
```

std::counting_semaphore

```
template <ptrdiff_t least_max_value = implementation-defined>
struct counting_semaphore {
    static constexpr ptrdiff_t max() noexcept;

    constexpr explicit counting_semaphore(ptrdiff_t desired);

    void release(ptrdiff_t update = 1);

    void acquire();
    bool try_acquire() noexcept;
    template <typename Rep, typename Period>
        bool try_acquire_for(const chrono::duration<Rep, Period>& rel_time);
    template <typename Clock, typename Duration>
        bool try_acquire_until(const chrono::time_point<Clock, Duration>& abs_time);
};
```


std::counting_semaphore

```
template <ptrdiff_t least_max_value = implementation-defined>
struct counting_semaphore {
    static constexpr ptrdiff_t max() noexcept;

    constexpr explicit counting_semaphore(ptrdiff_t desired);

    void release(ptrdiff_t update = 1);

    void acquire();
    bool try_acquire() noexcept;
    template <typename Rep, typename Period>
        bool try_acquire_for(const chrono::duration<Rep, Period>& rel_time);
    template <typename Clock, typename Duration>
        bool try_acquire_until(const chrono::time_point<Clock, Duration>& abs_time);
};
```

std::counting_semaphore

```
template <ptrdiff_t least_max_value = implementation-defined>
struct counting_semaphore {
    static constexpr ptrdiff_t max() noexcept;

    constexpr explicit counting_semaphore(ptrdiff_t desired);

    void release(ptrdiff_t update = 1);

    void acquire();
    bool try_acquire() noexcept;
    template <typename Rep, typename Period>
        bool try_acquire_for(const chrono::duration<Rep, Period>& rel_time);
    template <typename Clock, typename Duration>
        bool try_acquire_until(const chrono::time_point<Clock, Duration>& abs_time);
};
```

std::counting_semaphore

```
template <ptrdiff_t least_max_value = implementation-defined>
struct counting_semaphore {
    static constexpr ptrdiff_t max() noexcept;

    constexpr explicit counting_semaphore(ptrdiff_t desired);

    void release(ptrdiff_t update = 1);

    void acquire();
    bool try_acquire() noexcept;
    template <typename Rep, typename Period>
        bool try_acquire_for(const chrono::duration<Rep, Period>& rel_time);
    template <typename Clock, typename Duration>
        bool try_acquire_until(const chrono::time_point<Clock, Duration>& abs_time);
};
```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

    T pop();

public:
    constexpr concurrent_bounded_queue() = default;

    void enqueue(std::convertible_to<T> auto&& u);

    T dequeue();
    std::optional<T> try_dequeue();
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

    T pop() {
        std::optional<T> tmp;
        std::scoped_lock l(items_mtx);
        tmp = std::move(items.front());
        items.pop();
        return std::move(*tmp);
    }
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

    T pop() {
        std::optional<T> tmp;
        std::scoped_lock l(items_mtx);
        tmp = std::move(items.front());
        items.pop();
        return std::move(*tmp);
    }
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

    T pop() {
        std::optional<T> tmp;
        std::scoped_lock l(items_mtx);
        tmp = std::move(items.front());
        items.pop();
        return std::move(*tmp);
    }
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

public:
    void enqueue(std::convertible_to<T> auto&& u) {
        remaining_space.acquire();
        {
            std::scoped_lock l(items_mtx);
            items.emplace(std::forward<decltype(u)>(u));
        }
        items_produced.release();
    }
};

```



```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

public:
    void enqueue(std::convertible_to<T> auto&& u) {
        remaining_space.acquire();
        {
            std::scoped_lock l(items_mtx);
            items.emplace(std::forward<decltype(u)>(u));
        }
        items_produced.release();
    }
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

public:
    void enqueue(std::convertible_to<T> auto&& u) {
        remaining_space.acquire();
        {
            std::scoped_lock l(items_mtx);
            items.emplace(std::forward<decltype(u)>(u));
        }
        items_produced.release();
    }
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

public:
    void enqueue(std::convertible_to<T> auto&& u) {
        remaining_space.acquire();
        {
            std::scoped_lock l(items_mtx);
            items.emplace(std::forward<decltype(u)>(u));
        }
        items_produced.release();
    }
};

```

```
template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

public:
    T dequeue() {
        items_produced.acquire();
        T tmp = pop();
        remaining_space.release();
        return std::move(tmp);
    }
};
```

```
template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

public:
    T dequeue() {
        items_produced.acquire();
        T tmp = pop();
        remaining_space.release();
        return std::move(tmp);
    }
};
```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

public:
    T dequeue() {
        items_produced.acquire();
        T tmp = pop();
        remaining_space.release();
        return std::move(tmp);
    }
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

public:
    T dequeue() {
        items_produced.acquire();
        T tmp = pop();
        remaining_space.release();
        return std::move(tmp);
    }
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

public:
    std::optional<T> try_dequeue() {
        if (!items_produced.try_acquire()) return {};
        T tmp = pop();
        remaining_space.release();
        return std::move(tmp);
    }
};

```



```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

public:
    std::optional<T> try_dequeue() {
        if (!items_produced.try_acquire()) return {};
        T tmp = pop();
        remaining_space.release();
        return std::move(tmp);
    }
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    std::mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

    T pop();

public:
    constexpr concurrent_bounded_queue() = default;

    void enqueue(std::convertible_to<T> auto&& u);

    T dequeue();
    std::optional<T> try_dequeue();
};

```

```
struct spin_mutex {  
private:  
    std::atomic<bool> flag = ATOMIC_VAR_INIT(false);  
  
public:  
    void lock() {  
        while (flag.exchange(true, std::memory_order_acquire))  
            ;  
    }  
  
    void unlock() {  
        flag.store(false, std::memory_order_release);  
    }  
};
```

```
struct spin_mutex {  
private:  
    std::atomic_flag flag = ATOMIC_FLAG_INIT;  
  
public:  
    void lock() {  
        while (flag.test_and_set(std::memory_order_acquire))  
            ;  
    }  
  
    void unlock() {  
        flag.clear(std::memory_order_release);  
    }  
};
```

```

struct spin_mutex {
private:
    std::atomic_flag flag = ATOMIC_FLAG_INIT;

public:
    void lock() {
        for (std::uint64_t k = 0; !flag.test_and_set(std::memory_order_acquire); ++k) {
            if (k < 4) ;
            else if (k < 16) __asm__ __volatile__ ( "rep; nop" : : : "memory" );
            else if (k < 64) sched_yield();
            else {
                timespec rqtp = { 0, 0 };
                rqtp.tv_sec = 0; rqtp.tv_nsec = 1000;
                nanosleep(&rqtp, nullptr);
            }
        }
    }

    void unlock() {
        flag.clear(std::memory_order_release);
    }
};

```

```

struct spin_mutex {
private:
    std::atomic_flag flag = ATOMIC_FLAG_INIT;

public:
    void lock() {
        for (std::uint64_t k = 0; !flag.test_and_set(std::memory_order_acquire); ++k) {
            if (k < 4) ;
            else if (k < 16) __asm__ __volatile__ ( "rep; nop" : : : "memory" );
            else if (k < 64) sched_yield();
            else {
                timespec rqtp = { 0, 0 };
                rqtp.tv_sec = 0; rqtp.tv_nsec = 1000;
                nanosleep(&rqtp, nullptr);
            }
        }
    }

    void unlock() {
        flag.clear(std::memory_order_release);
    }
};

```

```

struct spin_mutex {
private:
    std::atomic_flag flag = ATOMIC_FLAG_INIT;

public:
    void lock() {
        for (std::uint64_t k = 0; !flag.test_and_set(std::memory_order_acquire); ++k) {
            if (k < 4) ;
            else if (k < 16) __asm__ __volatile__ ( "rep; nop" : : : "memory" );
            else if (k < 64) sched_yield();
            else {
                timespec rqtp = { 0, 0 };
                rqtp.tv_sec = 0; rqtp.tv_nsec = 1000;
                nanosleep(&rqtp, nullptr);
            }
        }
    }

    void unlock() {
        flag.clear(std::memory_order_release);
    }
};

```

```

struct spin_mutex {
private:
    std::atomic_flag flag = ATOMIC_FLAG_INIT;

public:
    void lock() {
        for (std::uint64_t k = 0; !flag.test_and_set(std::memory_order_acquire); ++k) {
            if (k < 4) ;
            else if (k < 16) __asm__ __volatile__( "rep; nop" : : : "memory" );
            else if (k < 64) sched_yield();
            else {
                timespec rqtp = { 0, 0 };
                rqtp.tv_sec = 0; rqtp.tv_nsec = 1000;
                nanosleep(&rqtp, nullptr);
            }
        }
    }

    void unlock() {
        flag.clear(std::memory_order_release);
    }
};

```



```

struct spin_mutex {
private:
    std::atomic_flag flag = ATOMIC_FLAG_INIT;

public:
    void lock() {
        for (std::uint64_t k = 0; !flag.test_and_set(std::memory_order_acquire); ++k) {
            if (k < 4) ;
            else if (k < 16) __asm__ __volatile__ ( "rep; nop" : : : "memory" );
            else if (k < 64) sched_yield();
            else {
                timespec rqtp = { 0, 0 };
                rqtp.tv_sec = 0; rqtp.tv_nsec = 1000;
                nanosleep(&rqtp, nullptr);
            }
        }
    }

    void unlock() {
        flag.clear(std::memory_order_release);
    }
};

```

```

struct spin_mutex {
private:
    std::atomic_flag flag = ATOMIC_FLAG_INIT;

public:
    void lock() {
        for (std::uint64_t k = 0; !flag.test_and_set(std::memory_order_acquire); ++k) {
            if (k < 4) ;
            else if (k < 16) __asm__ __volatile__ ( "rep; nop" : : : "memory" );
            else if (k < 64) sched_yield();
            else {
                timespec rqtp = { 0, 0 };
                rqtp.tv_sec = 0; rqtp.tv_nsec = 1000;
                nanosleep(&rqtp, nullptr);
            }
        }
    }

    void unlock() {
        flag.clear(std::memory_order_release);
    }
};

```

```
struct spin_mutex {  
private:  
    std::atomic_flag flag = ATOMIC_FLAG_INIT;  
  
public:  
    void lock() {  
        while (flag.test_and_set(std::memory_order_acquire))  
            flag.wait(true, std::memory_order_relaxed);  
    }  
  
    void unlock() {  
        flag.clear(std::memory_order_release);  
        flag.notify_one();  
    }  
};
```

```
struct spin_mutex {  
private:  
    std::atomic<bool> flag = ATOMIC_VAR_INIT(false);  
  
public:  
    void lock() {  
        while (flag.exchange(true, std::memory_order_acquire))  
            flag.wait(true, std::memory_order_relaxed);  
    }  
  
    void unlock() {  
        flag.store(false, std::memory_order_release);  
        flag.notify_one();  
    }  
};
```

`std::atomic{ _flag }` wait and notify

```
template <typename T>
struct atomic {
    void wait(T, memory_order = memory_order::seq_cst) const volatile noexcept;
    void wait(T, memory_order = memory_order::seq_cst) const noexcept;
    void notify_one() volatile noexcept;
    void notify_one() noexcept;
    void notify_all() volatile noexcept;
    void notify_all() noexcept;
};
```

`std::atomic{ _flag }` wait and notify

```
template <typename T>
struct atomic {
    void wait(T, memory_order = memory_order::seq_cst) const volatile noexcept;
    void wait(T, memory_order = memory_order::seq_cst) const noexcept;
    void notify_one() volatile noexcept;
    void notify_one() noexcept;
    void notify_all() volatile noexcept;
    void notify_all() noexcept;
};
```

std::atomic{ _flag } wait and notify

```
template <typename T>
struct atomic {
    void wait(T, memory_order = memory_order::seq_cst) const volatile noexcept;
    void wait(T, memory_order = memory_order::seq_cst) const noexcept;
    void notify_one() volatile noexcept;
    void notify_one() noexcept;
    void notify_all() volatile noexcept;
    void notify_all() noexcept;
};
```

`std::atomic{ _flag }` wait and notify

Some possible implementation strategies

- ▶ Futex. Supported for certain size objects on Linux and Windows.
- ▶ Condition Variables. Supported for certain size objects on Linux and Mac.
- ▶ Contention Table. Used to optimize futex notify or to hold CVs.
- ▶ Timed back-off. Supported on everything.
- ▶ Spinlock. Supported on everything. Note: performance is terrible.

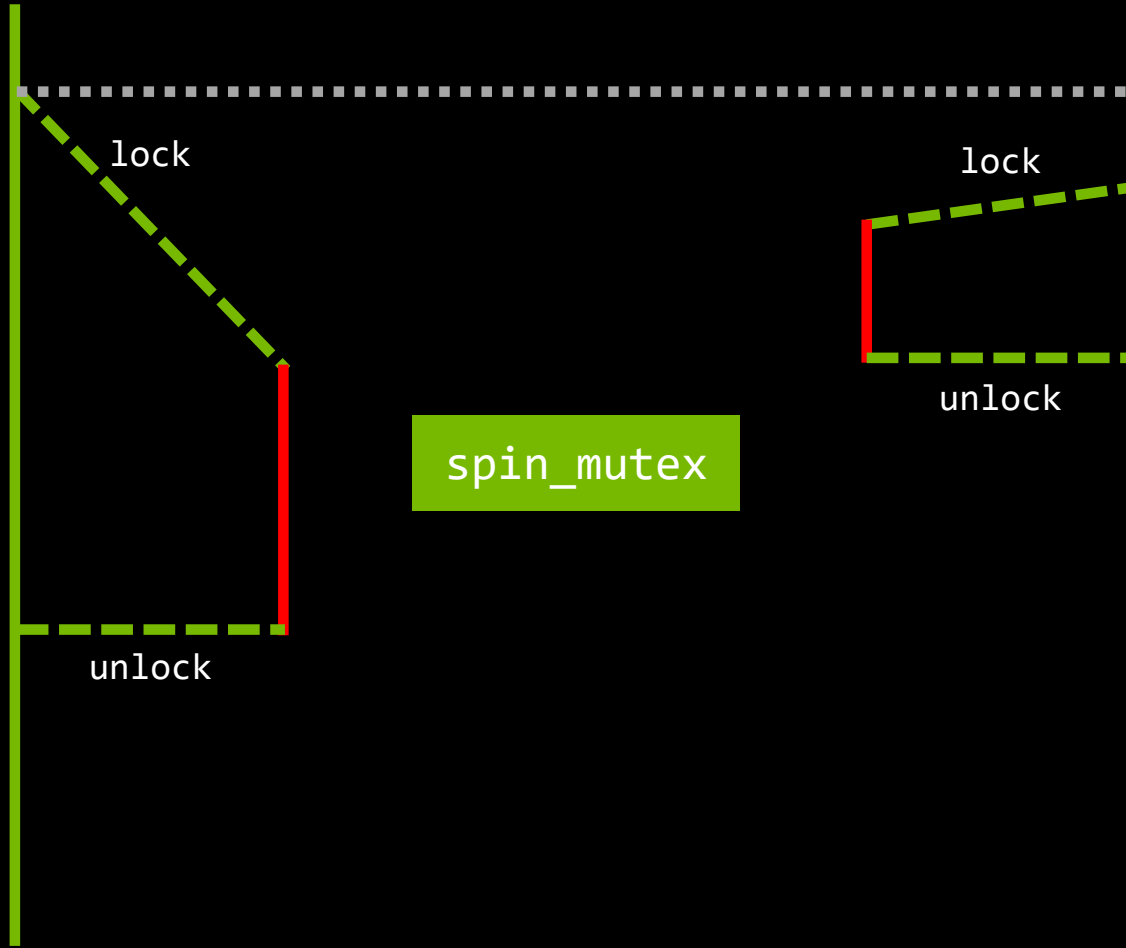
std::atomic_flag test

```
struct atomic_flag {  
    bool test(memory_order = memory_order::seq_cst) const volatile noexcept;  
    bool test(memory_order = memory_order::seq_cst) const noexcept;  
    bool test_and_set(memory_order = memory_order::seq_cst) volatile noexcept;  
    bool test_and_set(memory_order = memory_order::seq_cst) noexcept;  
    void clear(memory_order = memory_order::seq_cst) volatile noexcept;  
    void clear(memory_order = memory_order::seq_cst) noexcept;  
  
    void wait(bool, memory_order = memory_order::seq_cst) const volatile noexcept;  
    void wait(bool, memory_order = memory_order::seq_cst) const noexcept;  
    void notify_one() volatile noexcept;  
    void notify_one() noexcept;  
    void notify_all() volatile noexcept;  
    void notify_all() noexcept;  
};
```

```
struct spin_mutex {  
private:  
    std::atomic_flag flag = ATOMIC_FLAG_INIT;  
  
public:  
    void lock() {  
        while (flag.test_and_set(std::memory_order_acquire))  
            flag.wait(true, std::memory_order_relaxed);  
    }  
  
    void unlock() {  
        flag.clear(std::memory_order_release);  
        flag.notify_one();  
    }  
};
```


Thread A

Thread B



Thread A

Thread B



```
struct ticket_mutex {  
private:  
    std::atomic<int> in = ATOMIC_VAR_INIT(0);  
    std::atomic<int> out = ATOMIC_VAR_INIT(0);  
  
public:  
    void lock() {  
        auto const my = in.fetch_add(1, std::memory_order_acquire);  
        while (true) {  
            auto const now = out.load(std::memory_order_acquire);  
            if (now == my) return;  
            out.wait(now, std::memory_order_relaxed);  
        }  
    }  
  
    void unlock() {  
        out.fetch_add(1, std::memory_order_release);  
        out.notify_all();  
    }  
};
```

```
struct ticket_mutex {  
private:  
    std::atomic<int> in = ATOMIC_VAR_INIT(0);  
    std::atomic<int> out = ATOMIC_VAR_INIT(0);  
  
public:  
    void lock() {  
        auto const my = in.fetch_add(1, std::memory_order_acquire);  
        while (true) {  
            auto const now = out.load(std::memory_order_acquire);  
            if (now == my) return;  
            out.wait(now, std::memory_order_relaxed);  
        }  
    }  
  
    void unlock() {  
        out.fetch_add(1, std::memory_order_release);  
        out.notify_all();  
    }  
};
```

```

struct ticket_mutex {
private:
    std::atomic<int> in = ATOMIC_VAR_INIT(0);
    std::atomic<int> out = ATOMIC_VAR_INIT(0);

public:
    void lock() {
        auto const my = in.fetch_add(1, std::memory_order_acquire);
        while (true) {
            auto const now = out.load(std::memory_order_acquire);
            if (now == my) return;
            out.wait(now, std::memory_order_relaxed);
        }
    }

    void unlock() {
        out.fetch_add(1, std::memory_order_release);
        out.notify_all();
    }
};

```

```
struct ticket_mutex {  
private:  
    std::atomic<int> in = ATOMIC_VAR_INIT(0);  
    std::atomic<int> out = ATOMIC_VAR_INIT(0);  
  
public:  
    void lock() {  
        auto const my = in.fetch_add(1, std::memory_order_acquire);  
        while (true) {  
            auto const now = out.load(std::memory_order_acquire);  
            if (now == my) return;  
            out.wait(now, std::memory_order_relaxed);  
        }  
    }  
  
    void unlock() {  
        out.fetch_add(1, std::memory_order_release);  
        out.notify_all();  
    }  
};
```



```
struct ticket_mutex {
private:
    std::atomic<int> in = ATOMIC_VAR_INIT(0);
    std::atomic<int> out = ATOMIC_VAR_INIT(0);

public:
    void lock() {
        auto const my = in.fetch_add(1, std::memory_order_acquire);
        while (true) {
            auto const now = out.load(std::memory_order_acquire);
            if (now == my) return;
            out.wait(now, std::memory_order_relaxed);
        }
    }

    void unlock() {
        out.fetch_add(1, std::memory_order_release);
        out.notify_all();
    }
};
```

```

struct ticket_mutex {
private:
    std::atomic<int> in = ATOMIC_VAR_INIT(0);
    std::atomic<int> out = ATOMIC_VAR_INIT(0);

public:
    void lock() {
        auto const my = in.fetch_add(1, std::memory_order_acquire);
        while (true) {
            auto const now = out.load(std::memory_order_acquire);
            if (now == my) return;
            out.wait(now, std::memory_order_relaxed);
        }
    }

    void unlock() {
        out.fetch_add(1, std::memory_order_release);
        out.notify_all();
    }
};

```



```
struct ticket_mutex {  
private:  
    std::atomic<int> in = ATOMIC_VAR_INIT(0);  
    std::atomic<int> out = ATOMIC_VAR_INIT(0);  
  
public:  
    void lock() {  
        auto const my = in.fetch_add(1, std::memory_order_acquire);  
        while (true) {  
            auto const now = out.load(std::memory_order_acquire);  
            if (now == my) return;  
            out.wait(now, std::memory_order_relaxed);  
        }  
    }  
  
    void unlock() {  
        out.fetch_add(1, std::memory_order_release);  
        out.notify_all();  
    }  
};
```

```

struct ticket_mutex {
private:
    alignas(std::hardware_destructive_interference_size) std::atomic<int> in
        = ATOMIC_VAR_INIT(0);
    alignas(std::hardware_destructive_interference_size) std::atomic<int> out
        = ATOMIC_VAR_INIT(0);

public:
    void lock() {
        auto const my = in.fetch_add(1, std::memory_order_acquire);
        while (true) {
            auto const now = out.load(std::memory_order_acquire);
            if (now == my) return;
            out.wait(now, std::memory_order_relaxed);
        }
    }

    void unlock() {
        out.fetch_add(1, std::memory_order_release);
        out.notify_all();
    }
};

```

```

template <typename T, std::uint64_t QueueDepth>
struct concurrent_bounded_queue {
private:
    std::queue<T> items;
    ticket_mutex items_mtx;
    std::counting_semaphore<QueueDepth> items_produced{0};
    std::counting_semaphore<QueueDepth> remaining_space{QueueDepth};

    T pop();

public:
    constexpr concurrent_bounded_queue() = default;

    void enqueue(std::convertible_to<T> auto&& u);

    T dequeue();
    std::optional<T> try_dequeue();
};

```

Recipe For a Tasking Runtime

- ▶ Worker threads.
- ▶ Multi-consumer, multi-producer concurrent queue.
- ▶ **Termination detection mechanism.**
- ▶ Parallel algorithms.

```
template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

    void process_tasks(std::stop_token s);

public:
    bounded_depth_task_manager(std::uint64_t n);

    void submit(std::invocable auto&& f);
};
```



```
template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

    void process_tasks(std::stop_token s);

public:
    bounded_depth_task_manager(std::uint64_t n);

    void submit(std::invocable auto&& f);
};
```

```
template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

    void process_tasks(std::stop_token s);

public:
    bounded_depth_task_manager(std::uint64_t n);

    void submit(std::invocable auto&& f);
};
```

```
template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

    void process_tasks(std::stop_token s);

public:
    bounded_depth_task_manager(std::uint64_t n);

    void submit(std::invocable auto&& f);
};
```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue>();
    }

public:
    bounded_depth_task_manager(std::uint64_t n)
        : threads(n, [&] (std::stop_token s) { process_tasks(s); })
    {}
};

```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue>();
    }

public:
    bounded_depth_task_manager(std::uint64_t n)
        : threads(n, [&] (std::stop_token s) { process_tasks(s); })
    {}
};

```

```
template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

public:
    void submit(std::invocable auto&& f) {
        tasks.enqueue(std::forward<decltype(f)>(f));
    }
};
```

```
int main() {  
    std::atomic<std::uint64_t> count(0);  
  
    {  
        bounded_depth_task_manager<64> tm(6);  
  
        for (auto i : stdv::iota(0, 256))  
            tm.submit([&] { ++count; });  
    }  
  
    std::cout << count << "\n";  
}
```

```
int main() {  
    std::atomic<std::uint64_t> count(0);  
  
    {  
        bounded_depth_task_manager<64> tm(6);  
  
        for (auto i : std::iota(0, 256))  
            tm.submit([&] { ++count; });  
    }  
  
    std::cout << count << "\n";  
}
```



```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue>();
    }

public:
    bounded_depth_task_manager(std::uint64_t n)
        : threads(n, [&] (std::stop_token s) { process_tasks(s); })
    {}
};

```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue>();
        while (true) {
            if (auto f = tasks.try_dequeue()) std::move(*f)();
            else break;
        }
    }

public:
    bounded_depth_task_manager(std::uint64_t n)
        : threads(n, [&] (std::stop_token s) { process_tasks(s); })
    {}
};

```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue>();
        while (true) {
            if (auto f = tasks.try_dequeue()) std::move(*f)();
            else break;
        }
    }
public:
    ~bounded_depth_task_manager() {
        std::latch l(threads.size() + 1);
        for (auto i : std::iota(0, threads.size()))
            submit([&] { l.arrive_and_wait(); });
        threads.request_stop();
        l.count_down();
    }
};

```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue>();
        while (true) {
            if (auto f = tasks.try_dequeue()) std::move(*f)();
            else break;
        }
    }
public:
    ~bounded_depth_task_manager() {
        std::latch l(threads.size() + 1);
        for (auto i : std::iota(0, threads.size()))
            submit([&] { l.arrive_and_wait(); });
        threads.request_stop();
        l.count_down();
    }
};


```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue>();
        while (true) {
            if (auto f = tasks.try_dequeue()) std::move(*f)();
            else break;
        }
    }
public:
    ~bounded_depth_task_manager() {
        std::latch l(threads.size() + 1);
        for (auto i : std::iota(0, threads.size()))
            submit([&] { l.arrive_and_wait(); });
        threads.request_stop();
        l.count_down();
    }
};


```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue(); 
        while (true) {
            if (auto f = tasks.try_dequeue()) std::move(*f)();
            else break;
        }
    }
public:
    ~bounded_depth_task_manager() {
        std::latch l(threads.size() + 1);
        for (auto i : stdv::iota(0, threads.size()))
            submit([& { l.arrive_and_wait(); }]);
        threads.request_stop();
        l.count_down();
    }
};

```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue(); 
        while (true) {
            if (auto f = tasks.try_dequeue()) std::move(*f)();
            else break;
        }
    }
public:
    ~bounded_depth_task_manager() {
        std::latch l(threads.size() + 1);
        for (auto i : std::iota(0, threads.size()))
            submit([&] { l.arrive_and_wait(); });
        threads.request_stop();
        l.count_down();
    }
};

```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue>();
        while (true) {
            if (auto f = tasks.try_dequeue()) std::move(*f)();
            else break;
        }
    }
public:
    ~bounded_depth_task_manager() {
        std::latch l(threads.size() + 1);
        for (auto i : std::iota(0, threads.size()))
            submit([&] { l.arrive_and_wait(); });
        threads.request_stop();
        l.count_down();
    }
};

```



```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    void process_tasks(std::stop_token s) {
        while (!s.stop_requested())
            tasks.dequeue>();
        while (true) {
            if (auto f = tasks.try_dequeue()) std::move(*f)();
            else break;
        }
    }
public:
    ~bounded_depth_task_manager() {
        std::latch l(threads.size() + 1);
        for (auto i : std::iota(0, threads.size()))
            submit([&] { l.arrive_and_wait(); });
        threads.request_stop();
        l.count_down();
    }
};

```

std::latch

```
struct latch {  
    constexpr explicit latch(ptrdiff_t expected);  
  
    latch(const latch&) = delete;  
    latch& operator=(const latch&) = delete;  
  
    void count_down(ptrdiff_t update = 1);  
    bool try_wait() const noexcept;  
    void wait() const;  
    void arrive_and_wait(ptrdiff_t update = 1);  
};
```

std::latch

```
struct latch {  
    constexpr explicit latch(ptrdiff_t expected);  
  
    latch(const latch&) = delete;  
    latch& operator=(const latch&) = delete;  
  
    void count_down(ptrdiff_t update = 1);  
    bool try_wait() const noexcept;  
    void wait() const;  
    void arrive_and_wait(ptrdiff_t update = 1);  
};
```

std::latch

```
struct latch {  
    constexpr explicit latch(ptrdiff_t expected);  
  
    latch(const latch&) = delete;  
    latch& operator=(const latch&) = delete;  
  
    void count_down(ptrdiff_t update = 1);  
    bool try_wait() const noexcept;  
    void wait() const;  
    void arrive_and_wait(ptrdiff_t update = 1);  
};
```

std::latch

```
struct latch {  
    constexpr explicit latch(ptrdiff_t expected);  
  
    latch(const latch&) = delete;  
    latch& operator=(const latch&) = delete;  
  
    void count_down(ptrdiff_t update = 1);  
    bool try_wait() const noexcept;  
    void wait() const;  
    void arrive_and_wait(ptrdiff_t update = 1);  
};
```

std::latch

```
struct latch {  
    constexpr explicit latch(ptrdiff_t expected);  
  
    latch(const latch&) = delete;  
    latch& operator=(const latch&) = delete;  
  
    void count_down(ptrdiff_t update = 1);  
    bool try_wait() const noexcept;  
    void wait() const;  
    void arrive_and_wait(ptrdiff_t update = 1);  
};
```

```
void submit_tree(auto& tm, std::atomic<std::uint64_t>& count, std::uint64_t level) {
    ++count;
    if (0 != level) {
        tm.submit([&tm, &count, level] { submit_tree(tm, count, level - 1); });
        tm.submit([&tm, &count, level] { submit_tree(tm, count, level - 1); });
    }
}

int main() {
    std::atomic<std::uint64_t> count(0);

    {
        bounded_depth_task_manager<64> tm(6);

        submit_tree(tm, count, 8);
    }

    std::cout << count << "\n";
}
```

```

void submit_tree(auto& tm, std::atomic<std::uint64_t>& count, std::uint64_t level) {
    ++count;
    if (0 != level) {
        tm.submit([&tm, &count, level] { submit_tree(tm, count, level - 1); });
        tm.submit([&tm, &count, level] { submit_tree(tm, count, level - 1); });
    }
}

int main() {
    std::atomic<std::uint64_t> count(0);

    {
        bounded_depth_task_manager<64> tm(6);

        submit_tree(tm, count, 8);
    }

    std::cout << count << "\n";
}

```



```

void submit_tree(auto& tm, std::atomic<std::uint64_t>& count, std::uint64_t level) {
    ++count;
    if (0 != level) {
        tm.submit([&tm, &count, level] { submit_tree(tm, count, level - 1); });
        tm.submit([&tm, &count, level] { submit_tree(tm, count, level - 1); });
    }
}

int main() {
    std::atomic<std::uint64_t> count(0);

    {
        bounded_depth_task_manager<64> tm(6);

        submit_tree(tm, count, 8);
    }

    std::cout << count << "\n";
}

```

```

void submit_tree(auto& tm, std::atomic<std::uint64_t>& count, std::uint64_t level) {
    ++count;
    if (0 != level) {
        tm.submit([&tm, &count, level] { submit_tree(tm, count, level - 1); });
        tm.submit([&tm, &count, level] { submit_tree(tm, count, level - 1); });
    }
}

int main() {
    std::atomic<std::uint64_t> count(0);

    {
        bounded_depth_task_manager<64> tm(6);

        submit_tree(tm, count, 8);
    }

    std::cout << count << "\n";
}

```

```

void submit_tree(auto& tm, std::atomic<std::uint64_t>& count, std::uint64_t level) {
    ++count;
    if (0 != level) {
        tm.submit([&tm, &count, level] { submit_tree(tm, count, level - 1); });
        tm.submit([&tm, &count, level] { submit_tree(tm, count, level - 1); });
    }
}

int main() {
    std::atomic<std::uint64_t> count(0);

    {
        bounded_depth_task_manager<64> tm(6);

        submit_tree(tm, count, 8);
    }

    std::cout << count << "\n";
}

```

```
template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

public:
    void submit(std::invocable auto&& f) {
        tasks.enqueue(std::forward<decltype(f)>(f));
    }
};
```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

public:
    void submit(std::invocable auto&& f) {
        while (!tasks.try_enqueue(std::forward<decltype(f)>(f)))
            make_progress();
    }

    void make_progress() {
        if (auto f = tasks.try_dequeue()) std::move(*f)();
    }
};

```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

public:
    void submit(std::invocable auto&& f) {
        while (!tasks.try_enqueue(std::forward<decltype(f)>(f)))
            make_progress();
    }

    void make_progress() {
        if (auto f = tasks.try_dequeue()) std::move(*f)();
    }
};

```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

public:
    void submit(std::invocable auto&& f) {
        while (!tasks.try_enqueue(std::forward<decltype(f)>(f)))
            make_progress();
    }

    void make_progress() {
        if (auto f = tasks.try_dequeue()) std::move(*f)();
    }
};

```

```

template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

public:
    void submit(std::invocable auto&& f) {
        while (!tasks.try_enqueue(std::forward<decltype(f)>(f)))
            make_progress();
    }

    void make_progress() {
        if (auto f = tasks.try_dequeue()) std::move(*f)();
    }
};

```



```
template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

    void process_tasks(std::stop_token s);

public:
    bounded_depth_task_manager(std::uint64_t n);

    ~bounded_depth_task_manager();

    void submit(std::invocable auto&& f);

    void make_progress();
};
```

```
template <std::uint64_t QueueDepth>
struct bounded_depth_task_manager {
private:
    concurrent_bounded_queue<any_invocable<void()>, QueueDepth> tasks;
    thread_group threads;

    void process_tasks(std::stop_token s);

public:
    bounded_depth_task_manager(std::uint64_t n);

    ~bounded_depth_task_manager();

    void submit(std::invocable auto&& f);

    void make_progress();
};
```

Recipe For a Tasking Runtime

- ▶ Worker threads.
- ▶ Multi-consumer, multi-producer concurrent queue.
- ▶ Termination detection mechanism.
- ▶ **Parallel algorithms.**

```
template <std::range I, std::random_access_iterator O,  
        typename T, std::invocable</* ... */> BO>  
    requires /* ... */  
void histogram(I&& input, O output, T inc, OP op) {  
    std::for_each(input, [&] (auto&& t) { output[op(t)] += inc; });  
}
```

```
template <std::range I, std::random_access_iterator O,  
        typename T, std::invocable</* ... */> BO>  
    requires /* ... */  
void histogram(I&& input, O output, T inc, OP op) {  
    std::for_each(input, [&] (auto&& t) { output[op(t)] += inc; });  
}
```

```
template <std::range I, std::random_access_iterator O,  
         typename T, std::invocable</* ... */> BO>  
    requires /* ... */  
void histogram(I&& input, O output, T inc, OP op) {  
    std::for_each(input, [&] (auto&& t) { output[op(t)] += inc; });  
}
```

```
template <std::range I, std::random_access_iterator O,  
         typename T, std::invocable</* ... */> BO>  
    requires /* ... */  
void histogram(I&& input, O output, T inc, OP op) {  
    std::for_each(input, [&] (auto&& t) { output[op(t)] += inc; });  
}
```

```
template <execution_policy EP,  
         std::random_access_range I, std::random_access_iterator O,  
         typename T, std::invocable</* ... */> BO>  
    requires /* ... */  
void histogram(EP&& exec, I&& input, O output, T inc, OP op);
```



```
template <execution_policy EP,  
         std::random_access_range I, std::random_access_iterator O,  
         typename T, std::invocable</* ... */> BO>  
    requires /* ... */  
void histogram(EP&& exec, I&& input, O output, T inc, OP op);
```

```
template <execution_policy EP,  
         std::random_access_range I, std::random_access_iterator O,  
         typename T, std::invocable</* ... */> BO>  
    requires /* ... */  
void histogram(EP&& exec, I&& input, O output, T inc, OP op);
```

```
void histogram(EP&& exec, I&& input, O output, T inc, OP op) {  
    std::uint64_t const elements = stdr::distance(input);  
    std::uint64_t const chunks   = (1 + ((elements - 1) / exec.chunk_size()));  
  
    // ...  
}
```

```
void histogram(EP&& exec, I&& input, O output, T inc, OP op) {  
    std::uint64_t const elements = stdr::distance(input);  
    std::uint64_t const chunks   = (1 + ((elements - 1) / exec.chunk_size()));  
  
    // ...  
}
```

```
void histogram(EP&& exec, I&& input, O output, T inc, OP op) {  
    std::uint64_t const elements = stdr::distance(input);  
    std::uint64_t const chunks    = (1 + ((elements - 1) / exec.chunk_size()));  
  
    std::latch l(chunks);  
  
    // ...  
}
```

```
void histogram(EP&& exec, I&& input, O output, T inc, OP op) {  
    // ...  
  
    for (auto chunk : stdv::iota(0, chunks))  
        exec.submit(  
            [&, =chunk] {  
                // ...  
            }  
        );  
}
```

```
void histogram(EP&& exec, I&& input, O output, T inc, OP op) {  
    // ...  
  
    for (auto chunk : stdv::iota(0, chunks))  
        exec.submit(  
            // ...  
        );  
}
```

```

void histogram(EP&& exec, I&& input, O output, T inc, OP op) {
    // ...

    for (auto chunk : stdv::iota(0, chunks))
        exec.submit(
            [=, &exec, &input, &l] {
                auto const my_begin = chunk * exec.chunk_size();
                auto const my_end   = std::min(elements, (chunk + 1) * exec.chunk_size());

                // ...
            }
        );
}

```



```

void histogram(EP&& exec, I&& input, O output, T inc, OP op) {
    // ...

    for (auto chunk : stdv::iota(0, chunks))
        exec.submit(
            [=, &exec, &input, &l] {
                auto const my_begin = chunk * exec.chunk_size();
                auto const my_end   = std::min(elements, (chunk + 1) * exec.chunk_size());

                // ...
            }
        );
}

```

```

void histogram(EP&& exec, I&& input, O output, T inc, OP op) {
    // ...

    for (auto chunk : stdv::iota(0, chunks))
        exec.submit(
            [=, &exec, &input, &l] {
                auto const my_begin = chunk * exec.chunk_size();
                auto const my_end   = std::min(elements, (chunk + 1) * exec.chunk_size());

                std::for_each(stdr::begin(input) + my_begin,
                             stdr::begin(input) + my_end,
                             [&] (auto&& t) {
                                 output[op(t)] += inc;
                             });
            });
    // ...
}

```

```

void histogram(EP&& exec, I&& input, O output, T inc, OP op) {
    // ...

    for (auto chunk : stdv::iota(0, chunks))
        exec.submit(
            [=, &exec, &input, &l] {
                auto const my_begin = chunk * exec.chunk_size();
                auto const my_end   = std::min(elements, (chunk + 1) * exec.chunk_size());

                std::for_each(stdr::begin(input) + my_begin,
                             stdr::begin(input) + my_end,
                             [&] (auto&& t) {
                                 output[op(t)] += inc;
                             });
            });

    // ...
}
};
}

```

```

void histogram(EP&& exec, I&& input, O output, T inc, OP op) {
    // ...

    for (auto chunk : stdv::iota(0, chunks))
        exec.submit(
            [=, &exec, &input, &l] {
                auto const my_begin = chunk * exec.chunk_size();
                auto const my_end   = std::min(elements, (chunk + 1) * exec.chunk_size());

                std::for_each(stdr::begin(input) + my_begin,
                             stdr::begin(input) + my_end,
                             [&] (auto&& t) {
                                 std::atomic_ref r(output[op(t)]);
                                 r.fetch_add(inc, std::memory_order_relaxed);
                             });
            });

    // ...
};
}

```

`std::atomic_ref<T>`

`std::atomic<T>` holds a T.

```
template <struct T>
struct atomic {
private:
    T data; // exposition only
public:
    // ...
};
```

`std::atomic<T>` does not hold a T.

```
template <struct T>
struct atomic_ref {
private:
    T* ptr; // exposition only
public:
    explicit atomic_ref(T&);

    // Otherwise, same API as std::atomic.
};
```

```

void histogram(EP&& exec, I&& input, O output, T inc, OP op) {
    // ...

    for (auto chunk : stdv::iota(0, chunks))
        exec.submit(
            [=, &exec, &input, &l] {
                auto const my_begin = chunk * exec.chunk_size();
                auto const my_end   = std::min(elements, (chunk + 1) * exec.chunk_size());

                std::for_each(stdr::begin(input) + my_begin,
                             stdr::begin(input) + my_end,
                             [&] (auto&& t) {
                                 std::atomic_ref r(output[op(t)]);
                                 r.fetch_add(inc, std::memory_order_relaxed);
                             });
            });

    // ...
};
}

```

`std::atomic<floating-point>`

```
template<> struct atomic<floating-point> {  
    floating-point fetch_add(floating-point,  
                             memory_order = memory_order_seq_cst) volatile noexcept;  
    floating-point fetch_add(floating-point,  
                             memory_order = memory_order_seq_cst) noexcept;  
    floating-point fetch_sub(floating-point,  
                             memory_order = memory_order_seq_cst) volatile noexcept;  
    floating-point fetch_sub(floating-point,  
                             memory_order = memory_order_seq_cst) noexcept;  
};
```

```

void histogram(EP&& exec, I&& input, O output, T inc, OP op) {
    // ...

    for (auto chunk : stdv::iota(0, chunks))
        exec.submit(
            [=, &exec, &input, &l] {
                auto const my_begin = chunk * exec.chunk_size();
                auto const my_end   = std::min(elements, (chunk + 1) * exec.chunk_size());

                std::for_each(stdr::begin(input) + my_begin,
                             stdr::begin(input) + my_end,
                             [&] (auto&& t) {
                                 std::atomic_ref r(output[op(t)]);
                                 r.fetch_add(inc, std::memory_order_relaxed);
                             });

                l.count_down();
                while (!l.try_wait()) exec.make_progress();
            });
}

```



```

void histogram(EP&& exec, I&& input, O output, T inc, OP op) {
    // ...

    for (auto chunk : stdv::iota(0, chunks))
        exec.submit(
            [=, &exec, &input, &l] {
                auto const my_begin = chunk * exec.chunk_size();
                auto const my_end   = std::min(elements, (chunk + 1) * exec.chunk_size());

                std::for_each(stdr::begin(input) + my_begin,
                             stdr::begin(input) + my_end,
                             [&] (auto&& t) {
                                 std::atomic_ref r(output[op(t)]);
                                 r.fetch_add(inc, std::memory_order_relaxed);
                             });

                l.count_down();
                while (!l.try_wait()) exec.make_progress();
            }
        );
}

```

```
void histogram(EP&& exec, I&& input, O output, T inc, OP op) {  
    // ...  
  
    std::latch l(chunks);  
  
    for (std::uint64_t chunk = 0; chunk < chunks; ++chunk)  
        exec.submit(  
            // ...  
        );  
  
    while (!l.try_wait())  
        exec.make_progress();  
}
```

```
template <std::range I, std::weakly_incrementable O,  
          typename T, std::invocable</* ... */> BO>  
    requires /* ... */  
    O exclusive_scan(I&& input, O output, T init, OP op);
```

Exclusive Scan

a	b	c	d	e	f	g	h	i
---	---	---	---	---	---	---	---	---

Exclusive Scan

a	b	c	d	e	f	g	h	i
X	Xa	Xab	Xabc	Xabcd	Xabcde	Xabcdef	Xabcdefg	Xabcdefgh

Exclusive Scan

a	b	c	d	e	f	g	h	i
X	Xa	Xab	Xabc	Xabcd	Xabcde	Xabcdef	Xabcdefg	Xabcdefgh

Exclusive Scan

a	b	c	d	e	f	g	h	i
---	---	---	---	---	---	---	---	---

	a	ab
--	---	----

`std::exclusive_scan`

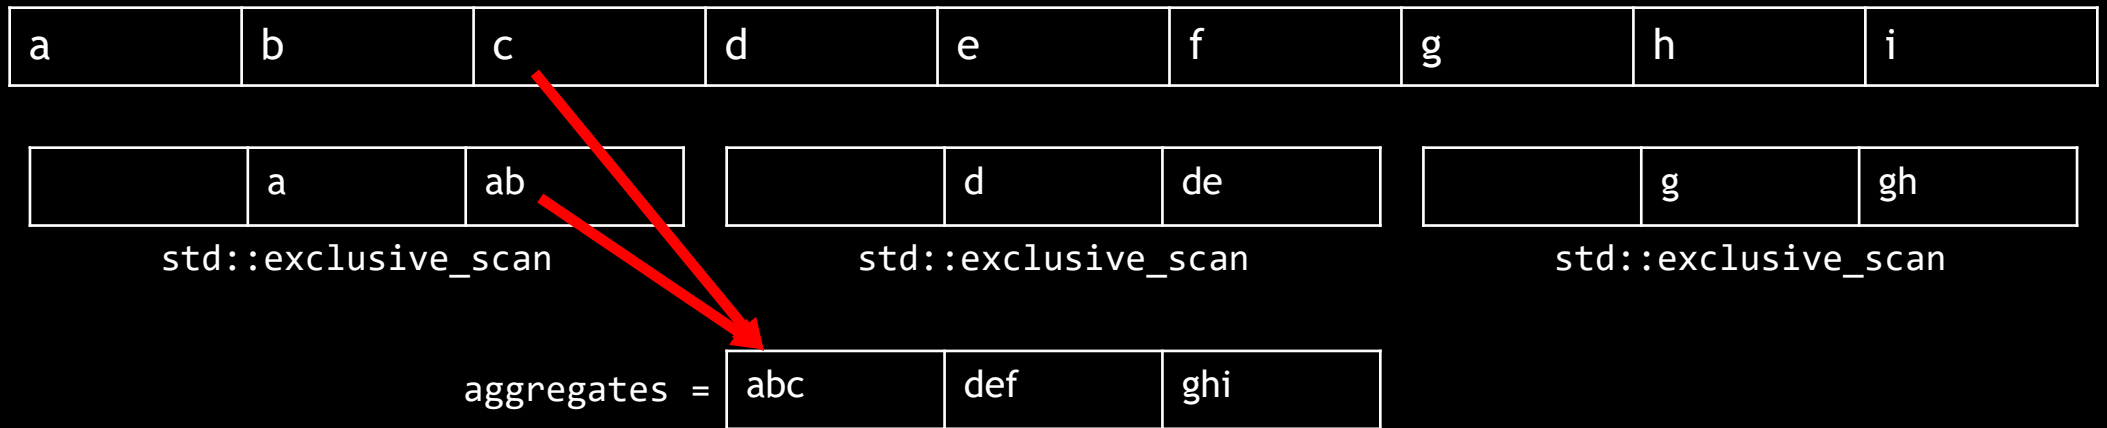
	d	de
--	---	----

`std::exclusive_scan`

	g	gh
--	---	----

`std::exclusive_scan`

Exclusive Scan



Exclusive Scan

a	b	c	d	e	f	g	h	i
---	---	---	---	---	---	---	---	---

	a	ab
--	---	----

`std::exclusive_scan`

	d	de
--	---	----

`std::exclusive_scan`

	g	gh
--	---	----

`std::exclusive_scan`

aggregates =	abc	def	ghi
--------------	-----	-----	-----

`std::exclusive_scan (init = X)`

Exclusive Scan

a	b	c	d	e	f	g	h	i
---	---	---	---	---	---	---	---	---

	a	ab
--	---	----

`std::exclusive_scan`

	d	de
--	---	----

`std::exclusive_scan`

	g	gh
--	---	----

`std::exclusive_scan`

aggregates =	X	Xabc	Xabcdef
--------------	---	------	---------

`std::exclusive_scan (init = X)`

Exclusive Scan

a	b	c	d	e	f	g	h	i
---	---	---	---	---	---	---	---	---

	a	ab
--	---	----

`std::exclusive_scan`

	d	de
--	---	----

`std::exclusive_scan`

	g	gh
--	---	----

`std::exclusive_scan`

`aggregates =`

X	Xabc	Xabcdef
---	------	---------

`std::exclusive_scan (init = X)`

X	Xa	Xab
---	----	-----

Increment by `aggregates[0]`

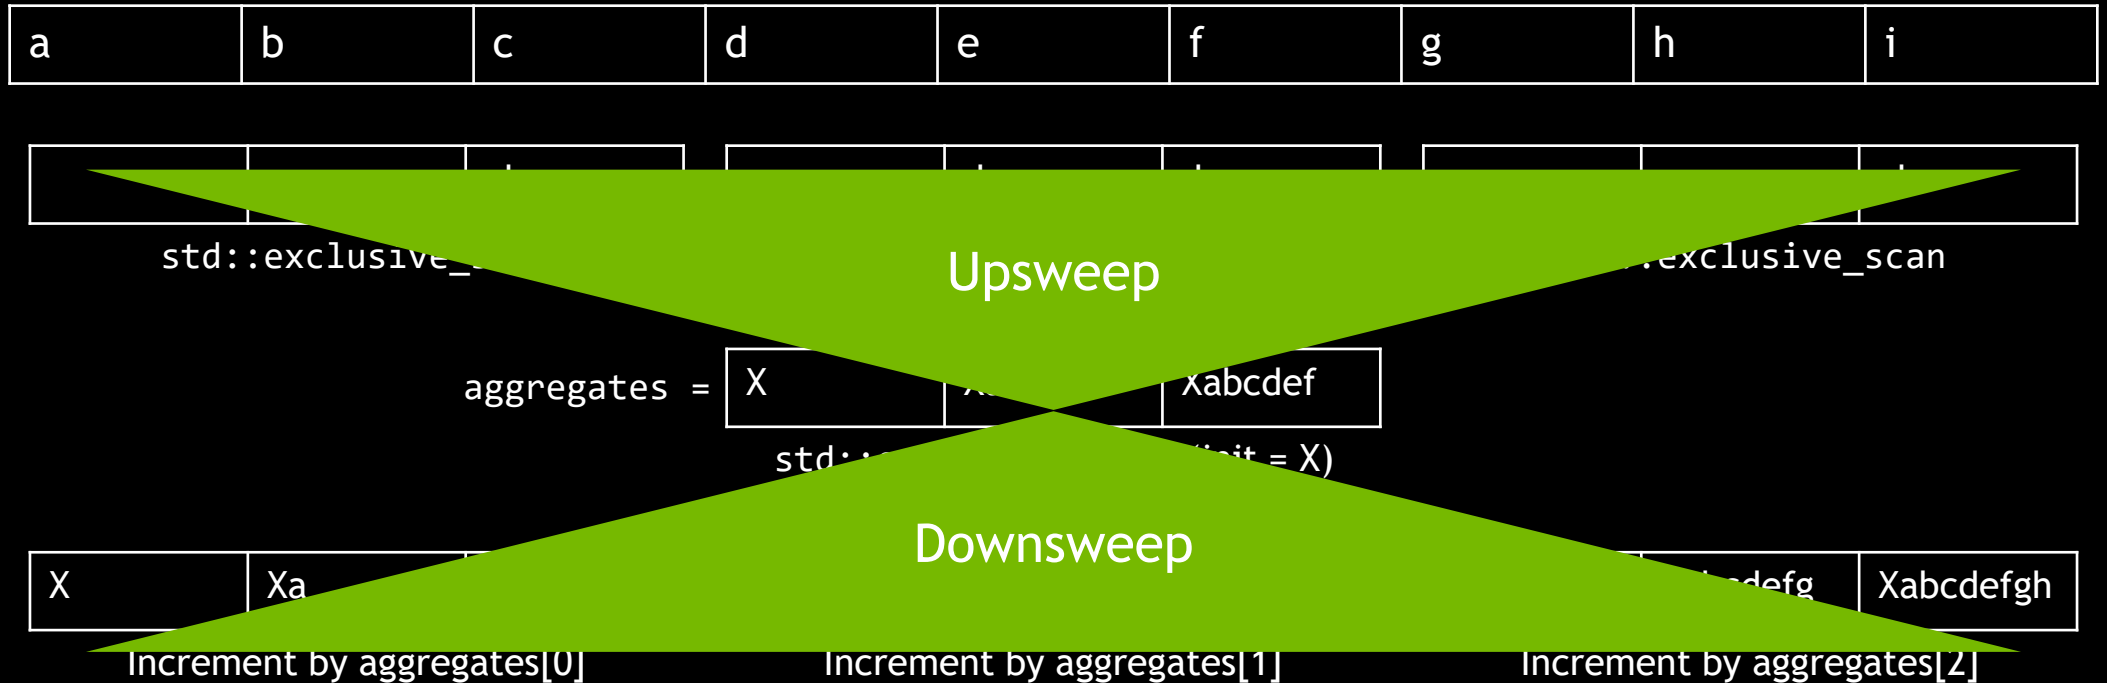
Xabc	Xabcd	Xabcde
------	-------	--------

Increment by `aggregates[1]`

Xabcdef	Xabcdefg	Xabcdefgh
---------	----------	-----------

Increment by `aggregates[2]`

Exclusive Scan



```
template <std::random_access_range I, std::random_access_iterator O,  
          typename T, std::invocable</* ... */> BO>  
    requires /* ... */  
struct exclusive_scanner {  
    // ...  
};
```

```
template <std::random_access_range I, std::random_access_iterator O,  
         typename T, std::invocable</* ... */> BO>  
    requires /* ... */  
    struct exclusive_scanner {  
private:  
    I input;  
    O output;  
    T init;  
    BO op;  
  
    // ...  
};
```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
    struct exclusive_scanner {
private:
    I input;
    O output;
    T init;
    BO op;

    std::uint64_t const concurrency;
    std::uint64_t const elements;
    std::uint64_t const chunk_size;

    // ...
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
struct exclusive_scanner {
private:
    I input;
    O output;
    T init;
    BO op;

    std::uint64_t const concurrency;
    std::uint64_t const elements;
    std::uint64_t const chunk_size;

    std::vector<T> aggregates;
    std::barrier<std::function<void()>> upsweep_barrier;
    std::latch downsweep_latch;

    // ...
};

```



```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
struct exclusive_scanner {
private:
    I input;
    O output;
    T init;
    BO op;

    std::uint64_t const concurrency;
    std::uint64_t const elements;
    std::uint64_t const chunk_size;

    std::vector<T> aggregates;
    std::barrier<std::function<void()>> upsweep_barrier;
    std::latch downswep_latch;

    // ...
};

```

std::barrier

```
template <typename CompletionFunction = see below>
struct barrier {
    using arrival_token = see below;

    constexpr explicit barrier(ptrdiff_t expected,
                               CompletionFunction f = CompletionFunction());

    [[nodiscard]] arrival_token arrive(ptrdiff_t update = 1);
    void wait(arrival_token&& arrival) const;

    void arrive_and_wait();
    void arrive_and_drop();
};
```

std::barrier

```
template <typename CompletionFunction = see below>
struct barrier {
    using arrival_token = see below;

    constexpr explicit barrier(ptrdiff_t expected,
                               CompletionFunction f = CompletionFunction());

    [[nodiscard]] arrival_token arrive(ptrdiff_t update = 1);
    void wait(arrival_token&& arrival) const;

    void arrive_and_wait();
    void arrive_and_drop();
};
```

std::barrier

```
template <typename CompletionFunction = see below>
struct barrier {
    using arrival_token = see below;

    constexpr explicit barrier(ptrdiff_t expected,
                               CompletionFunction f = CompletionFunction());

    [[nodiscard]] arrival_token arrive(ptrdiff_t update = 1);
    void wait(arrival_token&& arrival) const;

    void arrive_and_wait();
    void arrive_and_drop();
};
```

std::barrier

```
template <typename CompletionFunction = see below>
struct barrier {
    using arrival_token = see below;

    constexpr explicit barrier(ptrdiff_t expected,
                               CompletionFunction f = CompletionFunction());

    [[nodiscard]] arrival_token arrive(ptrdiff_t update = 1);
    void wait(arrival_token&& arrival) const;

    void arrive_and_wait();
    void arrive_and_drop();
};
```

std::barrier

```
template <typename CompletionFunction = see below>
struct barrier {
    using arrival_token = see below;

    constexpr explicit barrier(ptrdiff_t expected,
                               CompletionFunction f = CompletionFunction());

    [[nodiscard]] arrival_token arrive(ptrdiff_t update = 1);
    void wait(arrival_token&& arrival) const;

    void arrive_and_wait();
    void arrive_and_drop();
};
```

std::latch

```
struct latch {  
    constexpr explicit latch(ptrdiff_t expected);  
  
    latch(const latch&) = delete;  
    latch& operator=(const latch&) = delete;  
  
    void count_down(ptrdiff_t update = 1);  
    bool try_wait() const noexcept;  
    void wait() const;  
    void arrive_and_wait(ptrdiff_t update = 1);  
};
```

`std::latch` vs `std::barrier`

`std::latch`

- ▶ Supports asynchronous arrival.
- ▶ Single phase.
- ▶ No thread identity:
 - ▶ Threads may arrive multiple times.
 - ▶ Any thread may wait on a latch.
- ▶ No completion function.

`std::barrier`

- ▶ Supports asynchronous arrival.
- ▶ Multi phase.
- ▶ Thread identity:
 - ▶ A thread may arrive only once per phase.
 - ▶ Only a thread who has arrived may wait.
- ▶ Supports completion functions.


```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
struct exclusive_scanner {
private:
    I input;
    O output;
    T init;
    BO op;

    std::uint64_t const concurrency;
    std::uint64_t const elements;
    std::uint64_t const chunk_size;

    std::vector<T> aggregates;
    std::barrier<std::function<void()>> upsweep_barrier;
    std::latch downswEEP_latch;

    // ...
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
struct exclusive_scanner {
private:
    I input;
    O output;
    T init;
    BO op;

    std::uint64_t const concurrency;
    std::uint64_t const elements;
    std::uint64_t const chunk_size;

    std::vector<T> aggregates;
    std::barrier<std::function<void()>> upsweep_barrier;
    std::latch downsweep_latch;

    // ...
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
struct exclusive_scanner {
private:
    I input;
    O output;
    T init;
    BO op;

    std::uint64_t const concurrency;
    std::uint64_t const elements;
    std::uint64_t const chunk_size;

    std::vector<T> aggregates;
    std::barrier<std::function<void()>> upsweep_barrier;
    std::latch downswep_latch;

    void scan_aggregates();
public:
    exclusive_scanner(std::uint64_t concurrency_, I input_, O output_, T init_, BO op_);

    void launch(execution_policy auto&& exec);

    void process_chunk(std::uint64_t chunk, std::random_access_range auto&& my_input, O my_output);
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
    struct exclusive_scanner {
    public:
        exclusive_scanner(std::uint64_t concurrency_, I input_, O output_, T init_, BO op_);
        : input(std::move(input_))
        , output(std::move(output_))
        , init(std::move(init_))
        , op(std::move(op_))
        , concurrency(concurrency_)
        , elements(std::distance(input_))
        , chunk_size((elements + concurrency - 1) / concurrency)
        , aggregates(concurrency)
        , upsweep_barrier(concurrency, [&] { scan_aggregates(); })
        , downsweep_latch(concurrency)
        {}
    };

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
    struct exclusive_scanner {
public:
    exclusive_scanner(std::uint64_t concurrency_, I input_, O output_, T init_, BO op_);
        : input(std::move(input_))
        , output(std::move(output_))
        , init(std::move(init_))
        , op(std::move(op_))
        , concurrency(concurrency_)
        , elements(std::distance(input_))
        , chunk_size((elements + concurrency - 1) / concurrency)
        , aggregates(concurrency)
        , upsweep_barrier(concurrency, [&] { scan_aggregates(); })
        , downsweep_latch(concurrency)
    {}
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> B0>
    requires /* ... */
struct exclusive_scanner {
private:
    std::vector<T> aggregates;

    void scan_aggregates() {
        std::exclusive_scan(aggregates, aggregates.begin(), std::move(init), op);
    }
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
    struct exclusive_scanner {
public:
    void launch(execution_policy auto&& exec) {
        for (auto chunk : stdv::iota(0, concurrency))
            exec.submit(
                [&, chunk] {
                    auto const my_begin = chunk * chunk_size;
                    auto const my_end   = std::min(elements, (chunk + 1) * chunk_size);
                    process_chunk(chunk,
                                std::span(stdr::begin(input) + my_begin,
                                           stdr::begin(input) + my_end)
                                output + my_begin);
                }
            );
    }

    while (!downsweep_latch.try_wait())
        exec.make_progress();
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
    struct exclusive_scanner {
    public:
        void launch(execution_policy auto&& exec) {
            for (auto chunk : stdv::iota(0, concurrency))
                exec.submit(
                    [&, chunk] {
                        auto const my_begin = chunk * chunk_size;
                        auto const my_end   = std::min(elements, (chunk + 1) * chunk_size);
                        process_chunk(chunk,
                                    std::span(stdr::begin(input) + my_begin,
                                                stdr::begin(input) + my_end)
                                    output + my_begin);
                    }
                );
        }

        while (!downsweep_latch.try_wait())
            exec.make_progress();
    };
};

```



```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
    struct exclusive_scanner {
public:
    void launch(execution_policy auto&& exec) {
        for (auto chunk : stdv::iota(0, concurrency))
            exec.submit(
                [&, chunk] {
                    auto const my_begin = chunk * chunk_size;
                    auto const my_end   = std::min(elements, (chunk + 1) * chunk_size);
                    process_chunk(chunk,
                                std::span(stdr::begin(input) + my_begin,
                                           stdr::begin(input) + my_end)
                                output + my_begin);
                }
            );
    }

    while (!downsweep_latch.try_wait())
        exec.make_progress();
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
struct exclusive_scanner {
public:
    void process_chunk(std::uint64_t chunk, std::random_access_range auto&& my_input, O my_output) {
        // Upsweep.
        auto const& last_element = *(last - 1);
        aggregates[chunk] = op(*--std::exclusive_scan(my_input, my_output, T{}, op), last_element);

        // Barrier completion function does the aggregate sum.
        upsweep_barrier.arrive_and_wait();

        // Downsweep.
        std::for_each(my_output, my_output + std::distance(my_input),
                      [&, chunk] (auto& t) { t = op(std::move(t), aggregates[chunk]); });

        downsweep_latch.count_down();
    }
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
    struct exclusive_scanner {
    public:
        void process_chunk(std::uint64_t chunk, std::random_access_range auto&& my_input, O my_output) {
            // Upsweep.
            auto const& last_element = *(last - 1);
            aggregates[chunk] = op(*--std::exclusive_scan(my_input, my_output, T{}, op), last_element);

            // Barrier completion function does the aggregate sum.
            upsweep_barrier.arrive_and_wait();

            // Downsweep.
            std::for_each(my_output, my_output + std::distance(my_input),
                          [&, chunk] (auto& t) { t = op(std::move(t), aggregates[chunk]); });

            downsweep_latch.count_down();
        }
    };

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
struct exclusive_scanner {
public:
    void process_chunk(std::uint64_t chunk, std::random_access_range auto&& my_input, O my_output) {
        // Upsweep.
        auto const& last_element = *(last - 1);
        aggregates[chunk] = op(*--std::exclusive_scan(my_input, my_output, T{}, op), last_element);

        // Barrier completion function does the aggregate sum.
        upsweep_barrier.arrive_and_wait();

        // Downsweep.
        std::for_each(my_output, my_output + std::distance(my_input),
                      [&, chunk] (auto& t) { t = op(std::move(t), aggregates[chunk]); });

        downsweep_latch.count_down();
    }
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
struct exclusive_scanner {
public:
    void process_chunk(std::uint64_t chunk, std::random_access_range auto&& my_input, O my_output) {
        // Upsweep.
        auto const& last_element = *(last - 1);
        aggregates[chunk] = op(*--std::exclusive_scan(my_input, my_output, T{}, op), last_element);

        // Barrier completion function does the aggregate sum.
        upsweep_barrier.arrive_and_wait();

        // Downsweep.
        std::for_each(my_output, my_output + std::distance(my_input),
                      [&, chunk] (auto& t) { t = op(std::move(t), aggregates[chunk]); });

        downsweep_latch.count_down();
    }
};

```

```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
    struct exclusive_scanner {
    public:
        void process_chunk(std::uint64_t chunk, std::random_access_range auto&& my_input, O my_output) {
            // Upsweep.
            auto const& last_element = *(last - 1);
            aggregates[chunk] = op(*--std::exclusive_scan(my_input, my_output, T{}, op), last_element);

            // Barrier completion function does the aggregate sum.
            upsweep_barrier.arrive_and_wait();

            // Downsweep.
            std::for_each(my_output, my_output + std::distance(my_input),
                          [&, chunk] (auto& t) { t = op(std::move(t), aggregates[chunk]); });

            downsweep_latch.count_down();
        }
    };

```


`std::barrier`

Synchronization

N x arrives

1 x completion

N x waits complete



```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
    struct exclusive_scanner {
public:
    void process_chunk(std::uint64_t chunk, std::random_access_range auto&& my_input, O my_output) {
        // Upsweep.
        auto const& last_element = *(last - 1);
        aggregates[chunk] = op(*--std::exclusive_scan(my_input, my_output, T{}, op), last_element);

        // Barrier completion function does the aggregate sum.
        upsweep_barrier.arrive_and_wait();

        // Downsweep.
        std::for_each(my_output, my_output + std::distance(my_input),
                      [&, chunk] (auto& t) { t = op(std::move(t), aggregates[chunk]); });

        downsweep_latch.count_down();
    }
};

```



```

template <std::random_access_range I, std::random_access_iterator O,
          typename T, std::invocable</* ... */> BO>
    requires /* ... */
    struct exclusive_scanner {
public:
    void process_chunk(std::uint64_t chunk, std::random_access_range auto&& my_input, O my_output) {
        // Upsweep.
        auto const& last_element = *(last - 1);
        aggregates[chunk] = op(*--std::exclusive_scan(my_input, my_output, T{}, op), last_element);

        // Barrier completion function does the aggregate sum.
        upsweep_barrier.arrive_and_wait();

        // Downsweep.
        std::for_each(my_output, my_output + std::distance(my_input),
                      [&, chunk] (auto& t) { t = op(std::move(t), aggregates[chunk]); });

        downsweep_latch.count_down();
    }
};

```

C++20 Synchronization Library

- ▶ `std::atomic<T>` et al
 - ▶ wait/notify interface
 - ▶ `std::atomic_ref<T>`
 - ▶ test interface for `std::atomic_flag`
 - ▶ Floating-point specializations
- ▶ `std::latch` & `std::barrier`
- ▶ `std::counting_semaphore`
- ▶ `std::jthread`
 - ▶ Joining destructor
 - ▶ `std::stop_*` interruption mechanism



@blelbach

