

From Functions to Lambdas

From Functions to Lambdas: How Do C++ Callables *Really* Work?

WALTER E. BROWN, PH.D.

< webrown.cpp @ gmail.com >



Edition: 2019-09-20. Copyright © 2019 by Walter E. Brown. All rights reserved.

A little about me

- B.A. (math's); M.S., Ph.D. (computer science).
- Professional programmer for over 50 years, programming in C++ since 1982.
- Experienced in industry, academia, consulting, and research:
 - Founded a Computer Science Dept.; served as Professor and Dept. Head; taught and mentored at all levels.
 - Managed and mentored the programming staff for a reseller.
 - Lectured internationally as a software consultant and commercial trainer.
 - Retired from the Scientific Computing Division at Fermilab, specializing in C++ programming and in-house consulting.
- Not dead — still doing training & consulting. (Email me!)



Copyright © 2019 by Walter E. Brown. All rights reserved.

2

Emeritus participant in C++ standardization

- Written >160 papers for WG21, proposing such now-standard C++ library features as `gcd/lcm`, `cbegin/cend`, `common_type`, and `void_t`, as well as all of headers `<random>` and `<ratio>`.
- Influenced such core language features as *alias templates*, *contextual conversions*, and *variable templates*; recently worked on *requires-expressions*, `operator<=>`, and more!
- Conceived and served as Project Editor for *Int'l Standard on Mathematical Special Functions in C++* (ISO/IEC 29124), now incorporated into C++17's `<cmath>`.
- Be forewarned: Based on my training and experience, I hold some rather strong opinions about computer software and programming methodology — these opinions are not shared by all programmers, but they should be! ☺



Copyright © 2019 by Walter E. Brown. All rights reserved.

3

In today's talk

- The Big Picture
- Callable Types and Their Call Syntax
- Callbacks
- Function Objects and Their Types
- Higher-Order Functions
- Function Objects in Disguise
- An Advanced Example (as time permits)

Copyright © 2019 by Walter E. Brown. All rights reserved.

4

The Big Picture

Copyright © 2019 by Walter E. Brown. All rights reserved.

Our theme

- `()`:
 - Usually termed the call/invoke operator.
 - Distinct from parentheses used to group (an operator w/ its operands) or to punctuate (a comma-separated list).
- `g(...)` is the commonest syntax to apply the operator:
 - Intending to yield control (of the CPU) to entity `g`, and ...
 - Expecting to resume control when `g` returns (finishes).
 - Akin to a news anchor who calls upon a field reporter, later resuming when that reporter says "Back to you."


Copyright © 2019 by Walter E. Brown. All rights reserved.

6

From Functions to Lambdas

Comparing basic flow of control (many variations; not in this talk)

- Sequential



Copyright © 2013 by Walter E. Brown. All rights reserved.

7

Example

- ```
void second() {
 puts("Second\n");
 return;
}
```
- ```
int main() {  
    puts( "First\n" );  
    second();  
    puts( "Third\n" );  
    return 0;  
}
```
- Fctn main calls upon fctn second via a call expression:
 - So main is the caller/client ...
 - And second is the callee/server.
- When executed, emits:

```
First  
Second  
Third
```

Copyright © 2013 by Walter E. Brown. All rights reserved.

8

Such control flow has many variations

- E.g., if you have multiple CPUs:
 - Could give each a separate task to do, while you keep going with your own, and all rendezvous at the end.
 - Such independent executions are a form of parallelism.
- E.g., if you have multiple threads on a single CPU:
 - Akin to parallelism, but now tasks can synchronize with each other to share common resources.
 - Such cooperative executions are a form of concurrency.
- Exploring these (and more!) would need many hours:
 - So today we'll focus on the classical essentials only, ...
 - And leave such special topics to other talks/speakers.

Copyright © 2013 by Walter E. Brown. All rights reserved.

9

Making a call is like hiring someone

- E.g., imagine hiring an attorney:
 - The hiree (att'y) may specify certain documents needed from the hirer (you) to provide the requested service.
- Similarly, when applying a call operator:
 - Hirer/hiree ~ caller/callee ~ client/server relationship.
 - The callee may specify certain values needed from the caller to provide the requested service.
- The needed items are termed parameters, while the corresponding supplied items are the arguments:
 - Establishing the correspondence, usually early in a call, is termed parameter passage.

Copyright © 2013 by Walter E. Brown. All rights reserved.

10

Parameter syntax

- In C++, the type of each callee param encodes both:
 - The kind of value required from the caller, and ...
 - How that corresponding arg will be used.

Syntax	Nomenclature	Callee's promise
<code>T</code>	variable	will be modified
<code>T const &</code>	reference to const	will not be modified
<code>T & T &&</code>	reference to reference	will not be modified

Copyright © 2013 by Walter E. Brown. All rights reserved.

11

How is a call carried out? (since C++17; see [expr.call])

- A call expression `L (R)` evaluates **L before R**:
 - The left operand **L** is a *postfix-expression* (often just a name) that denotes a callable entity (often a function).
 - The right operand **R** is a sequence of *initializer-clauses* that will be evaluated in some non-overlapping order.
- To invoke/call **L with arguments R** means to apply the `()` operator to those evaluated operands:
 - Each call yields control to **L** via a **prolog**, and ...
 - Each call resumes CPU control via an **epilog**.
 - The details of this calling convention vary per platform.

Copyright © 2013 by Walter E. Brown. All rights reserved.

12

From Functions to Lambdas

Prolog: a call is initiated in the context of the caller/client

- Sufficient storage (a stack frame) is reserved (allocated) for all of the callee's automatic lifetime needs:
 - CPU registers may be allocated, too (but their contents may need first to be saved/spilled for restoration upon return).
 - This storage is used for parameters (explicitly declared + implicit/hidden), as well as for other fctn-local variables.
- Examples of implicit parameters:
 - A return address for resuming control.
 - `this`, a mbr fctn's pointer to its invoking object (*`this`).
 - A ptr/ref to storage for the call's result (unless void).

Copyright © 2013 by Walter E. Brown. All rights reserved.

13

More prolog: can't yet yield control

- Parameter passage now takes place by initializing each param's storage with the corresponding arg:
 - Each param `p` (of type `P`) is init'd from corresp arg `a` just as in a declaration `P p = a;` // recall: `P` may be a ref type.
 - (Recall that if `a`'s type is not `P`, a temporary may be created to hold `a`'s promoted/decayed/converted value.)
 - Each such init (`arg eval + any side effects`) is indeterminately sequenced with respect to any other param's init.
- The implicit parameters' storage is likewise init'd:
 - With arg's arranged/provided by the compiler.
 - But (of course) not yet the call's result.

Copyright © 2013 by Walter E. Brown. All rights reserved.

14

What can go wrong?

- An exception might arise during an init:
 - The prolog discontinues further parameter passage, ...
 - Destroys each successfully init'd param, ...
 - Then reclaims all the stack frame storage.
- Now process the exception:
 - Search for a matching exception handler, starting in the caller (since we never yielded control to the callee), ...
 - Init the handler's param with (a copy of) the exception object, ...
 - Finally, yield control to the handler.

Copyright © 2013 by Walter E. Brown. All rights reserved.

15

Still not ready for the callee to have control (C++23?)

- Preconditions `[[pre: ...]]` are part of the contract between the caller and callee functions:
 - Each precondition must be satisfied (true) "immediately before starting [execution] of the [callee] function ..."
 - Otherwise, the program cannot trust its state, so the contract violation handler function will be called.
- A violation handler typically ends program execution:
 - Why? Because a callee is not obliged to behave sanely when a caller fails to satisfy callee preconditions.
 - In the standard library, "Violation of any preconditions ... results in undefined behavior" (see `[res.on.required]`).

Copyright © 2013 by Walter E. Brown. All rights reserved.

16

"There and Back Again" (apologies to J. R. R. Tolkien)

- Only now can control be yielded (flow) to the callee:
 - The callee begins execution with its 1st statement, ...
 - And continues with its own internal control flow ...
 - Possibly making calls of its own along the way ...
 - Until it executes a return statement to start the epilog.
- "The result of a fctn call is the result of the operand [if any] of the evaluated return stmt..." (see `[expr.call]/9`):
 - That result then initializes the implicit result parameter.
 - Certain optimizations (e.g., RVO, NRVO, ...) can reduce (sometimes significantly) the costs of that initialization.

Copyright © 2013 by Walter E. Brown. All rights reserved.

17

More epilog to conclude a call normally

- Before leaving the callee, d'tors are invoked for non-trivial locals (i.e., for var's and temp's, not yet param's):
 - Happens in the reverse order of construction.
- Then the caller regains control (per the return address hidden parameter):
 - Now back in the caller's context (where they were constructed), param's (and their temp's) are destroyed.
 - Usually happens promptly for param's, but delayed until the end of the call's enclosing *full-expression* for temp's.
 - Finally, registers are restored, the stack frame storage is reclaimed, and the call is considered complete.

Copyright © 2013 by Walter E. Brown. All rights reserved.

18

From Functions to Lambdas

Postconditions (C++23?)

- Postconditions are also part of the contract between the caller and callee functions:
 - “A postcondition is checked by evaluating its predicate immediately before returning control to the caller ...
 - [By which time the] lifetime of local variables and temporaries [but not of parameters] has ended.”
- If false, the program cannot trust its state, so the contract violation handler function will be called:
 - The caller is not obliged to behave sanely if the callee fails to satisfy its postconditions.
 - If the violation handler doesn't abort the program, whatever happens afterwards is undefined behavior.

Copyright © 2019 by Walter E. Brown. All rights reserved.

19

std:: fctns can conclude a call abnormally (decl'd in various hdrs)

- `[[noreturn]] void exit(int exit_code);`
 - Causes normal program termination w/ partial cleanup.
 - Similar to returning from main.
- `[[noreturn]] void _Exit(int exit_code) noexcept;`
 - Causes normal program termination w/out cleanup.
- `[[noreturn]] void abort() noexcept;`
 - Causes abnormal program termination w/out cleanup.
- `[[noreturn]] void quick_exit(int exit_code) noexcept;`
 - Causes quick program termination w/ partial cleanup.
- `[[noreturn]] void terminate() noexcept;`
 - Calls the current terminate_handler (abort by default).
 - Is implicitly called when exception handling fails.

Copyright © 2019 by Walter E. Brown. All rights reserved.

20

Callable Types and Their Call Syntax

Copyright © 2019 by Walter E. Brown. All rights reserved.

We can invoke/call an entity of type F (or F& or F&&) ...

- ① If F is a *fcn* type `R (...)` (also `noexcept(...)` since C++17).
- ② If F is a *ptr-to-fcn* type `R (*) (...)`.
- ③ If F is a *ptr-to-mbr-fcn* type `R (C :: *) (...)`.
- ④ If F is a *class* type defining a non-static member that is a fcn call operator `R F :: operator () (...)`.
- ⑤ If F is a *class* type ...
 - With an implicit conversion fcn `F :: operator G () ...`
 - Where G is a { *ptr, ref, or ref-to-ptr* } -to-fcn type.

Copyright © 2019 by Walter E. Brown. All rights reserved.

22

Speaking of function types ...

- Can a fcn type be cv- and/or ref-qualified?
 - E.g., using `bizarre_t = int () const &&;` // *valid type?*
 - Syntactically yes, but (despite appearances) such a fcn type is neither a cv-qualified type nor a ref-qualified type!
- So `std::is_const_v< F const >` is false if F is a fcn type!
 - (Or if F is a reference type.)
- Such abominable types are rarely useful:
 - Except perhaps in torture-testing a compiler ...
 - And in implementing the `std :: is_function` type trait.

Copyright © 2019 by Walter E. Brown. All rights reserved.

23

It is said ...

“In the beginning was the word.

“But by the time
the second word was added to it,
there was trouble.

“For with it came syntax....”

— John Simon,
1981



Copyright © 2019 by Walter E. Brown. All rights reserved.

24

From Functions to Lambdas

Call syntax depends mostly on the callee's type

- If an entity *g* has (ref to) *fcn*, *ptr-to-fcn*, or *class* type, use traditional call operator syntax *g*(*a*₁, *a*₂, ..., *a*_N).
- But if *C::g* is a non-static mbr fcn, the hidden *this* param must be init'd to point to the invoking object:
 - When *decltype(a₁)* is (or inherits from) type *C*, use call syntax *a₁.g* (*a*₂, ..., *a*_N).
 - Otherwise (e.g., *decltype(a₁)* is a ptr/smart-ptr type), use call syntax *((*a₁).g)* (*a*₂, ..., *a*_N).
- Similarly, if *g* has *ptr-to-mbr-fcn* type:
 - Use call syntax *(a₁.*g)* (*a*₂, ..., *a*_N) for case ①, and ...
 - Use call syntax *((*a₁).*g)* (*a*₂, ..., *a*_N) for case ②.

Copyright © 2013 by Walter E. Brown. All rights reserved.

25

Copyability of callables

If <i>decltype(g)</i> is a ...	Then <i>g</i> is a ...	Is <i>g</i> copyable?
		decay
<i>ptr-to-fcn</i> type	variable	yes, naturally so
<i>ptr-to-mbr-fcn</i> type	variable	yes, naturally so
<i>a</i> (<i>...</i>)	function object	depends on its type
reference type	reference	yes, but can't mutate non-reseatable another variable

But why do we care about copyability?

Copyright © 2013 by Walter E. Brown. All rights reserved.

26

Callbacks

Copyright © 2013 by Walter E. Brown. All rights reserved.

Example

- Let's use this function:
`bool is_even(int k) { return k % 2 == 0; }`
- ... in the following context:
`int a[N] = { ... };
std::partition(a + 0, a + N, is_even); // ⇒ &is_even`
- ... in order to rearrange the array's contents, placing all the even values ahead of all the odd values:
 - For each array item *a*[*k*], the partition algorithm calls *is_even*(*a*[*k*]) to decide where the item belongs.
 - Executable code (e.g., a function) that is passed as an arg to other code (e.g., to an algorithm) is termed a callback.

Copyright © 2013 by Walter E. Brown. All rights reserved.

28

How do callbacks work in C++?

- While functions can't be copied, they decay into *ptr-to-fcn* values that can be copied (and called, too).
- E.g., here's an int-specific partition algorithm:
 - `void partition(int * b, int * e // e is past-the-last
 , bool belongs_in_front(int) // bool(*)(int)
) {
 for(; ;) { // do find_if_not, then do rfind_if
 while(b != e and belongs_in_front(* b)) ++b;
 while(b != e and not belongs_in_front(* --e));
 if (b == e) return;
 std::swap_iter(b++, e);
 }
}`

Copyright © 2013 by Walter E. Brown. All rights reserved.

29

Equivalently, using named helper algorithms ...

- Given helpers *find_if_not* and *rfind_if*, our int-specific partition algorithm gains a much simpler structure:
 - `void partition(int * b, int * e
 , bool belongs_in_front(int)
) {
 while(b = find_if_not(b, e, belongs_in_front)
 , e = rfind_if(b, e, belongs_in_front)
 , b != e)
 swap_iter(b++, e);
}`
- Note that *partition* forwarded its callback to those helper algorithms for their use.

Copyright © 2013 by Walter E. Brown. All rights reserved.

30

From Functions to Lambdas

Possible generic helpers (C++20-style)

- Algorithm `find_if` not:
 - `template< InputIterator Iter, Predicate<...> Pred >`
`Iter find_if_not(Iter b, Iter e, Pred p) { // e is past-the-last`
`for(; b != e and p(*b); ++b) ;`
`return b;`
`}`
- Algorithm `rfind_if`:
 - `template< BidirectionalIterator Iter, Predicate<...> Pred >`
`Iter rfind_if(Iter b, Iter e, Pred p) { // e is past-the-last`
`while(b != e and not p(*--e)) ;`
`return b;`
`}`

Copyright © 2019 by Walter E. Brown. All rights reserved.

31

So what characterizes a callback?

- A callback is an entity such that:
 - Its type `F` satisfies `invocable<F, ...>` (C++20).
 - It is supplied to an algorithm, usually as an argument of known or deduced type.
 - The algorithm then invokes the callback as needed, ...
 - Thereby tailoring the algorithm's behavior with respect to the client's data.
- The callback thus serves as the go-between:
 - It has/provides the client-specific details ...
 - That the general algorithm lacks.

Copyright © 2019 by Walter E. Brown. All rights reserved.

32

C++ callbacks have a long history

- E.g., in the C library headers `<stdlib.h>/<cstdlib>`:
 - `void* qsort(void* ptr, size_t count, size_t size`
`, int comp(void const*, void const*));`
 - `void* bsearch(void const* key`
`, void const* ptr, size_t count, size_t size`
`, int comp(void const*, void const*));`
- E.g., in the C++ library header `<exception>`:
 - using `terminate_handler = void (*) ();`
 - `terminate_handler`
`set_terminate(terminate_handler h) noexcept;`

Copyright © 2019 by Walter E. Brown. All rights reserved.

33

INVOKE and std::invoke (C++11/C++17, resp.)

- For specification purposes, the std lib denotes the call-syntax selection algorithm as *INVOKE*(g, a_1, \dots, a_N):
 - Provides a uniform interface to all callable/invokable.
- INVOKE* exists only notionally, but `std::invoke` is real:
 - `template< class G, class... A >`
`constexpr invoke_result_t< G, A... > invoke(G&& g, A&&... a)`
`noexcept(is_nothrow_invocable_v< G, A... >);`
- Each also correctly handles ref types for g and for a_i :
 - To decide the correct call syntax, ref types `F&/F&&` and `std::reference_wrapper<F>` types are treated as type `F`.
 - And if a_i has a `std::reference_wrapper<...>` type, it is unwrapped (via `.get()`) in the selected call syntax.

Copyright © 2019 by Walter E. Brown. All rights reserved.

34

The std::invocable concepts (C++20)

- `template< class G, class... A >`
`concept invocable = requires(G&& g, A&&... a) {`
`invoke(forward<G>(g), forward<A>(a)...);`
`// not required to be equality-preserving`
`};`
- `template< class G, class... A >`
`concept regular_invocable = invocable< G, A... >;`
`// is required to be equality-preserving`
- An expression is said to be equality-preserving iff:
 - Multiple evaluations with identical inputs/operands ...
 - Always give the same outputs/results/side effects.
 - E.g., a URBG satisfies `invocable`, not `regular_invocable`.

Copyright © 2019 by Walter E. Brown. All rights reserved.

35

Function Objects and Their Types

Copyright © 2019 by Walter E. Brown. All rights reserved.

From Functions to Lambdas

Characteristics of C++ function objects

- Any function object (e.g., a ptr-to-fctn) is used both:
 - As a value (construct/copy/destroy it), and ...
 - As a callable entity (invoke it to get a result + side effects).
- Thus, a function object of class type typically has:
 - ① A default c'tor if it's stateless (no data members), else a c'tor that initializes its state (data members), ...
 - ② A copy c'tor (and maybe also a move c'tor), ...
 - ③ A d'tor, and ...
 - ④ At least one member fctn (could be a fctn template) named operator (). (A suitable conv op will do, too.)

Copyright © 2013 by Walter E. Brown. All rights reserved.

37

Example

(C++20)

- A (stateful) function object type:
 - class is_less {
int n;
public:
is_less(int n = 0) : n{ n } { } // serves as default c'tor
// implicitly copyable and destructible
bool operator() (int val) const noexcept
{ return val < n; }
};
- A function object, instantiated for use as a callback:
 - int a[N] = { ... };
std::partition(a + 0, a + N, is_less{ 42 });

Copyright © 2013 by Walter E. Brown. All rights reserved.

38

Function object types in the standard library

- Header <functional> provides these fctn object types:
 - plus, minus, multiplies, divides, modulus, negate;
 - equal_to, not_equal_to;
 - greater, less, greater_equal, less_equal;
 - logical_and, logical_or, logical_not;
 - bit_and, bit_or, bit_xor, bit_not;
- These are often instantiated for use with <algorithm>:
 - std::greater<> gt;
 - std::sort(from, upto, gt); // sort will call fctn object gt
 - std::sort(from, upto, std::greater<>()); // prefer { }
 - // sort will call the (anonymous, temporary) fctn object

Copyright © 2013 by Walter E. Brown. All rights reserved.

39

How to convert a fctn obj type to a ptr-to-fctn type

- struct Doubler {
using fctn_t = int (int); // convenient alias, not essential
constexpr int operator () (int k) const noexcept
{ return 2 * k; }
static constexpr int decayed(int k) noexcept
{ return Doubler{}(k); }
constexpr operator fctn_t * () const noexcept
{ return &decayed; }
};
- int (*fp)(int) = Doubler{}; // uses conversion operator
... fp(3) ... // yields 6

Copyright © 2013 by Walter E. Brown. All rights reserved.

40

std::function as a universal callable handle

- template< class > class function; // not defined in general
- template< class R (P...) >
class function< R (P...) > { // template arg is a fctn type
R operator() (P...) const; // behaves as INVOKE does
};
- Think of it as a generalized ptr-to-fctn type:
 - Because it can be initialized from any callable that ...
 - Obtains a result of (or convertible to) type R, and ...
 - Takes arg's that can init param's of types P... .

Copyright © 2013 by Walter E. Brown. All rights reserved.

41

Example: using std::function as a callback parameter

- Revising the earlier int-specific partition algorithm:
 - void partition(int * b, int * e // e is past-the-last
// ~~bool belongs_in_front(int)~~ // bool (*)(int)
, std::function<bool(int)> belongs_in_front
) {
// algorithm's body is unchanged
}
- This still accepts a callback arg of ptr-to-fctn type:
 - But an arg of any fctn obj type is equally accepted, ...
 - As long as that arg is int-taking and bool-returning.

Copyright © 2013 by Walter E. Brown. All rights reserved.

42

From Functions to Lambdas

Another example: a container of callables

- `void fctn() { cout << "I'm an ordinary function.\n"; }`
- `struct obj {
 void operator() () const
 { cout << "I'm a stateless function object.\n"; }
};`
- `struct state {
 int n = 0; //stateful
 void operator() () const
 { cout << "Times I've been called: " << ++n << ".\n"; }
};`
- *// a container of callables, each taking/returning nothing:*
`vector< function<void()> > v{ &fctn, obj{ }, stateful{ } };`
`for(auto&& g : v) g(); //invoke each callable in v`

Copyright © 2019 by Walter E. Brown. All rights reserved.

43

Proposed for C++next

- `std::function_ref`:
 - Designed as a (non-owning) reference to a Callable.
 - Interface very much like that of `std::function`.
- Like any reference (or handle or view) type:
 - Has the classical coordination of lifetime issue arises.
 - Must keep the Callable alive while the `function_ref` lives.
 - Temporary Callables are especially vulnerable.

Copyright © 2019 by Walter E. Brown. All rights reserved.

44

Callbacks Generalized: Higher-Order Functions

Copyright © 2019 by Walter E. Brown. All rights reserved.

Characteristics

- A higher-order function is a ~~function~~ callable that:
 - Has at least one callback parameter, and/or ...
 - Has a result of a callable type.
- Mathematics is full of higher-order functions:
 - Prominent examples include differentiation and integration operations, each of which is a function ...
 - That produces a function as its result, having taken a function as its parameter.
- In computer programming languages, higher-order functions have been around since LISP (1958):
 - Modern LISP is based on lambda calculus (Church, 1936).

Copyright © 2019 by Walter E. Brown. All rights reserved.

46

Example: composition of callables

- `using fctn_t = std::function<int(int)>;`
- `struct Composer {
 fctn_t f1, f2;
 Composer(fctn_t f1, fctn_t f2) f1(f1), f2(f2) { }
 int operator() (int k) { return f1(f2(k)); }
};`
- `int double(int k) { return 2 * k; }`
`int triple(int k) { return 3 * k; }`
- `static_assert(Composer{ double, triple } (5)
 == 30); // == 2 * (3 * 5)`

Copyright © 2019 by Walter E. Brown. All rights reserved.

47

Function Objects in Disguise: Lambda Expressions and Their Closures

Copyright © 2019 by Walter E. Brown. All rights reserved.

From Functions to Lambdas

C++ lambda expression syntax

- The fundamental syntax is `[...] (...) { ... } ;`
 - Brackets provide for captures. (Stay tuned!)
 - Parentheses provide for function parameters.
 - Braces provide for a function body.
- C++20's full syntax also permits such extra items as:
 - A return type and/or an exception specification, ...
 - Constrained function parameters, ...
 - Template parameters (optionally constrained), ...
 - And more.

Copyright © 2019 by Walter E. Brown. All rights reserved.

49

What can we do with a C++ lambda?

- We can initialize a variable with one:
 - *E.g.*, `auto greater = [] (int a, int b) { return a > b; };`
`bool b = greater(1, 2); //later`
- We can pass one as a callback:
 - *E.g.*, `std::sort(from, upto`
`, [] (int a, int b) { return a > b; });`
- We can return one from a function (*e.g.*, factory-style):
 - *E.g.*, `auto make()`
`{ return [] (int a, int b) { return a > b; }; }`
`auto greater = make();`
- We can even invoke one immediately:
 - *E.g.*, `bool b = [] (int a, int b) { return a > b; } (1, 2);`

Copyright © 2019 by Walter E. Brown. All rights reserved.

50

A lambda is a kind of expression, nothing more

- Like any expression, a *lambda-expression* is evaluated:
 - Such evaluation produces a temporary function object termed a closure (closure object).
 - *I.e.*, it has class type, with a public operator `()` member and most/all of the usual special member functions.
 - Thus, a closure has both value and callable semantics.
- It's always the closure that serves:
 - As the initializer in a declaration, or ...
 - As the callback (argument) in a function call, or ...
 - As the callee in a call expression, or ...
 - As the result in a return statement.

Copyright © 2019 by Walter E. Brown. All rights reserved.

51

A closure's type

- What's the type of this lambda's closure?
 - `auto greet = [] () { std::cout << "Hello!"; };`
 - `struct { // 1st approximation; more detail forthcoming`
`inline void operator() () const { std::cout << "Hello!"; }`
`} greet;`
- Each closure has a uniquely named type:
 - That type's true name is known to the compiler only.
 - The type is nonetheless deducible (as usual) via `auto` or via a template type parameter.
 - Can also use `decltype(greet)`, if there's a need.

Copyright © 2019 by Walter E. Brown. All rights reserved.

52

Calling a closure

- A variable of a closure type is called exactly as we would call any other function object:
 - `... greet () ...`
- An anonymous closure is called only immediately:
 - `... [] () { std::cout << "Hello!"; } () ...`
- Either can be passed as a callback argument:
 - `template< class G >`
`void emit (G g) { g (); }`
 - `... emit(greet) ...`
 - `... emit([] () { std::cout << "Hello!"; }) ...`

Copyright © 2019 by Walter E. Brown. All rights reserved.

53

Return type?

- The lambda's return type can often be deduced:
 - Exactly as is done for a function with `auto` return type.
- No deduction if there's a late-specified return type:
 - Placed after the parameters') and before the body's {.
 - Exactly as is done for a function with `auto` return type.
- Example:
 - `auto greet = [] ()`
`-> void`
`{ std::cout << "Hello!"; };`
 - More items syntactically permitted there.

Copyright © 2019 by Walter E. Brown. All rights reserved.

54

From Functions to Lambdas

The body of a lambda

- C++11 limited the body to a single stmt, a return.
 - Since, most limitations have been lifted.
 - But there remains an inherent scope/lifetime issue.
- Although a closure starts out as a temporary:
 - Like any copyable value, it can be saved for later use, ...
 - Possibly from an entirely different scope, ...
 - Thus outliving any nonlocal variables used in its body!
- When a closure outlives the lifetime of variables it uses when called, Bad Things usually ensue.

Copyright © 2013 by Walter E. Brown. All rights reserved.

55

Problematic example

(doesn't compile)

- Note the lifetimes of parameter pi and closure L :
 - `auto get_work(double pi) { // automatic lifetime ...`
`auto L = [] () { return pi; }; // ... leads to ...`
`return L;`
`}`
 - `auto work = get_work(3.14);`
`auto answer2 = work(); // ... undefined behavior!`
- C++ forbids a lambda to use any automatic-lifetime variable from an enclosing scope:
 - Unless it is first captured, by value or by reference, ...
 - Much like a function argument is passed.

Copyright © 2013 by Walter E. Brown. All rights reserved.

56

Captures

- The contents of a *lambda-introducer* (the introductory brackets) specify the lambda's captures:
 - For each captured variable, the compiler injects a corresponding data member into the closure type.
 - If the brackets are empty, the lambda has no captures and its closure is described as stateless.
 - Otherwise, the closure is described as stateful.
- Captures can be specified individually or in bulk:
 - If individually, there will be 1 data member per capture.
 - If in bulk, the compiler will determine how many data members are needed.

Copyright © 2013 by Walter E. Brown. All rights reserved.

57

Capture syntax for stateful closures

(partial list; most common only)

Syntax (between [])	Will inject (into the closure) ...	Since
<code>&</code>		
<code>Identifier</code>	A copy of the spec'd automatic var	C++11
<code>&&</code>		
<code>&*</code>		
<code>*&</code>		
<code>&*</code>		
<code>*&</code>		

Copyright © 2013 by Walter E. Brown. All rights reserved.

58

Example of captures

(C++20)

- `auto echo(auto x) {`
`auto C1 = [=] () // by-value capture in bulk`
`{ return x; };`
`auto C2 = [y = C1()] () // single init-capture`
`{ return y; };`
`return C2(); // just returns x`
`}`
- `int main() {`
`return echo(6); // tedious form of exit status 6`
`}`

Copyright © 2013 by Walter E. Brown. All rights reserved.

59

Closure types' special member functions have evolved

- Recall that C++'s special member functions are those class c'tors, d'tor, and assignment op's that a compiler can sometimes generate for us.

Spec mbr fctn	C++11	C++14/17, C++20 stateful	C++20 stateless
destructor	= default	= default	= default
default c'tor	= delete	none	= default
copy c'tor, move c'tor	"implicitly declared"	= default	= default
copy assign	= delete	= delete	= default
move assign	not declared	not declared	= default

Copyright © 2013 by Walter E. Brown. All rights reserved.

60

From Functions to Lambdas

An Advanced Example

Copyright © 2019 by Walter E. Brown. All rights reserved.

Scenario

- Suppose I will want to make a call $w(a_1, \dots, a_n)$:
 - Of the n arg's, I already know a_1, \dots, a_k now, but ...
 - Only later will I know the remaining a_{k+1}, \dots, a_n .
- I want to combine (bind) w with the k known arg's:
 - Producing an equivalent new callable (say, b) that has the known a_1, \dots, a_k hard-wired/built-in.
 - I.e., $\text{auto } b = \text{combiner}(w, a_1, \dots, a_k)$;
 - Later, once the remaining a_{k+1}, \dots, a_n are known, I can call $b(a_{k+1}, \dots, a_n)$ and have the effect of calling $w(a_1, \dots, a_n)$.
 - C++20 has such a combiner, `std::bind_front`.

Copyright © 2019 by Walter E. Brown. All rights reserved.

62

A form of currying (C++20)

```
template< class callable_t, class... bound_t >
auto my_bind_front( callable_t && c, bound_t &&... b ) {
    return [ c = c
           , ...b = b
           ]
    ( auto && ... free ) //mutable? noexcept?
    -> decltype(auto)
    { return c( b...
               , free ...
               );
    };
}
```

Copyright © 2019 by Walter E. Brown. All rights reserved.

63

A form of currying ② (fwd \Rightarrow std::forward) (C++20)

```
template< class callable_t, class... bound_t >
auto my_bind_front( callable_t && c, bound_t &&... b ) {
    return [ c = fwd<callable_t>( c )
           , ...b = fwd<bound_t>( b )
           ]
    ( auto && ... free ) //mutable? noexcept?
    -> decltype(auto)
    { return c( fwd<bound_t>( b ) ...
               , fwd<decltype(free)>( free ) ...
               );
    };
}
```

Copyright © 2019 by Walter E. Brown. All rights reserved.

64

A form of currying ③ (invoke \Rightarrow std::invoke) (C++20)

```
template< class callable_t, class... bound_t >
auto my_bind_front( callable_t && c, bound_t &&... b ) {
    return [ c = fwd<callable_t>( c )
           , ...b = fwd<bound_t>( b )
           ]
    ( auto && ... free ) //mutable? noexcept?
    -> decltype(auto)
    { return invoke( fwd<callable_t>( c )
                   , fwd<bound_t>( b ) ...
                   , fwd<decltype(free)>( free ) ...
                   );
    };
}
```

Copyright © 2019 by Walter E. Brown. All rights reserved.

65

From Functions to Lambdas: How Do C++ Callables *Really* Work?

FIN

WALTER E. BROWN, PH.D.

< webrown.cpp @ gmail.com >



Copyright © 2019 by Walter E. Brown. All rights reserved.