

High Performance Graphics and Text Rendering on the GPU

Barbara Geller & Ansel Sermersheim
CppCon - September 2019

Introduction

- Prologue
- Brief Introduction to Graphics
- Why Graphics Matters to <almost> Everyone
- Using Graphics in a C++ Application
- Graphics Demo
- Graphics Terminology
- Program/Data Flow
- Rendering Text as Graphics
- Roadmap

Who is CopperSpice

- Maintainers and Co-Founders
 - CopperSpice
 - set of cross platform C++ libraries (*linux, os x, windows*)
 - CsString
 - support for UTF-8 and UTF-16, extensible to other encodings
 - CsSignal
 - thread aware signal / slot library
 - libGuarded
 - library for managing access to data shared between threads
 - DoxyPress
 - documentation generator for C++ and various other languages to HTML, LaTeX, DocBook, and others

Who is CopperSpice

- Credentials
 - all libraries and applications are open source
 - developed using cutting edge C++ technology
 - source code hosted on github
 - prebuilt binaries available on our download site
 - documentation generated by DoxyPress
 - youtube channel for C++ with over 40 videos
 - speakers at multiple CppCon and CppNow conferences
 - speakers at emBO++ 2019
 - numerous presentations in the US and Europe

- Additional Information
 - graphics timeline from 1989 through 2018
 - what drives the graphics industry
 - changes in GPU design
 - overview of the various API specifications
 - Evolution of Graphics Technology
 - <https://www.youtube.com/watch?v=u5SNd9sKn94>
 - GPU, Pipeline, and the Vector Graphics API
 - <https://www.youtube.com/watch?v=CrKxMrLczis>
 - Rendering 3D Graphics
 - <https://www.youtube.com/watch?v=MXz2t0gvRxl>

- Brief Introduction to Graphics
 - graphics are responsible for displaying images and text in a way which is effective and meaningful to the consumer
 - classified into distinct categories
 - raster graphics
 - bitmap image rendered by addressing discrete pixels
 - typically processed on the CPU
 - resolution dependent, does not scale without loss of quality
 - vector graphics
 - an image defined by vertex coordinates
 - drawing commands are based on complicated mathematics
 - should be processed on the GPU
 - scaled to a larger size the image quality is not compromised

- Image Quality
 - VGA
 - 307,000 pixels
 - 640 x 480
 - 1080p, Full HD
 - 2.07 million pixels
 - 1920 x 1080
 - 4k, UHD
 - 8.2 million pixels
 - 3840 x 2160
 - 8k, 8k UHD
 - 33.1 million pixels
 - 7680 x 4320

- Processors

- CPU

- average around 4-8 cores
 - high end around 10-20 cores
 - typically 1 or 2 threads run on each core

- GPU

- average around 1000 cores
 - high end around 3000 cores
 - many are hyperthreaded which means there are typically 5-60 threads running on each core

- Using Graphics in a C++ Application
 - gaming industry has the biggest influence over changes in graphics
 - changes in the GPU design require new programming tools
 - graphics are more than just explosions, robots, and lightsabers
 - are graphics important to the user of my desktop program
 - graphics for the GUI developer
 - high dpi support for free
 - responsive user interface
 - smooth scrolling of text
 - clean colors and edges
 - scalable charts, graphs, animations

- Demo Dependencies
 - C++17
 - CMake version 3.8 or newer
 - SDL2 library for window management
 - handles user input like keyboard and mouse events
 - open source answer to DirectX
 - cross platform
 - lots of games leverage this library
 - similar to GLUT, GLEW, or GLFW
 - <insert demo here>

- Structure of our Demo Program
 - demo.cpp
 - setup, main event loop (60 lines)
 - demo.h
 - two structures, function declarations (60 lines)
 - data.cpp
 - transform calculations (200 lines)
 - resources
 - mesh files, font image, shaders
 - links with SDL and one other library

- What is behind the Demo
 - CsPaint library
 - open source BSD license
 - encapsulates the Vulkan API exposing a higher level API
 - works on any GPU which supports Vulkan 1.1
 - uses vulkan.hpp which is the C++ interface
 - currently 38 source files
 - GLM
 - used for matrix math / linear algebra
 - provides buffer containers like the class `glm::vec3`
 - bundled with CsPaint
 - MIT license

- **CsPaint**
 - platform independent
 - windows and linux have native drivers
 - android native drivers available since version 7
 - MoltenVK is a Vulkan wrapper on OS X and iOS
 - using the CsPaint library allows you to render graphics without having to focus on the tedious and repetitive sections
 - instead you are free to concentrate on writing your shaders and modeling 3D images or simply displaying text

- Why Vulkan
 - Khronos Group took OpenGL and brought it closer to the metal
 - OpenGL is unable to send multiple commands
 - GPU sits idle too much of the time
 - CPU spends too much time waiting for the GPU
 - burden is now on the developer
 - no built in memory management, must be implemented and controlled by the developer
 - fixed function pipeline does not exist so developers must supply their own shader implementations
 - more configuration is required with fewer defaults supplied
 - same API for Linux, Windows, and Android
 - well supported translation layer for OS X and iOS

- Demo Source Code (set up, demo.cpp)

```
auto window                = init_window();
auto [context, surface] = init_vulkan(window);

auto device                = surface->graphicsDevice();
auto vertexShader          = device->createShader(vertexShaderData);
auto fragmentShader        = device->createShader(fragmentShaderData);
auto textFragmentShader    = device->createShader(textFragmentShaderData);

auto graphicsQueue         = device->graphicsQueue();
auto transferQueue         = device->transferQueue();
auto renderPass            = device->createRenderPass();
auto commandpool           = device->graphicsCommandPool();
auto commandbuffer         = commandpool->createCommandBuffer();

init_render(device, surface, commandpool, vertexShader, fragmentShader,
            textFragmentShader, renderPass);
```

- Demo Source Code (event loop, demo.cpp)

```
while (run) {
    SDL_Event evt;

    while (SDL_PollEvent(&evt)) {
        if (evt.type == SDL_QUIT) {
            run = 0;

        } else if (evt.type == SDL_KEYDOWN) {
            // process key events like the arrow keys to rotate the copper pot
        }
    }

    draw_frame(device, surface, renderPass,
               transform_matrix(glm::vec3(x_rotation, y_rotation, 0.0)), zoom_factor);
}
```


- Demo Data (fragment shader, shader.frag)

```
layout(location = 0) in  vec3 inNormal;
layout(location = 1) in  vec3 inColor;
layout(location = 2) in  vec3 inEyePos;
layout(location = 3) in  vec3 inLightVec;
layout(location = 0) out vec4 outColor;

void main() {
    vec3 Eye          = normalize(-inEyePos);
    vec3 Reflected    = normalize(reflect(-inLightVec, inNormal));

    vec4 IAmbient      = vec4(0.4, 0.4, 0.4, 1.0);
    vec4 IDiffuse       = vec4(0.5, 0.5, 0.5, 0.5) * max(dot(inNormal, inLightVec),0.0);
    vec4 ISpecular      = vec4(0.9, 0.9, 0.9, 1.0) * pow(max(dot(Reflected, Eye), 0.0),
                                                         1.5) * 0.50;

    outColor = vec4((IAmbient + IDiffuse) * vec4(inColor, 1.0) + ISpecular);
}
```

- Graphics Terminology
 - instance, context, surface
 - memory heap
 - vertex buffer, uniform buffer
 - shaders (vertex, tessellation, geometry, fragment)
 - pipeline
 - command buffer
 - textures
 - frame buffer, depth buffer, render pass
 - swapchain
 - queues (graphics, present, transfer, compute)
 - fences, semaphores
 - face culling, winding
 - push constants, scissor, extent, viewport

- Graphics Definitions

- instance

- connection between your application and the Vulkan library
 - typically your application will only create one
 - every other Vulkan call relies on passing this instance or some data associated with this instance

- context

- Vulkan instance is similar to an OpenGL context
 - not a term used in the Vulkan API

- surface

- a term for the window region where images are rendered
 - native window support is not handled directly in the API, implemented in platform extensions

- Graphics Definitions

- memory heap
 - provides storage for buffers which will be accessed by the GPU
 - every buffer must be created and managed by the user
 - memory is allocated from a memory heap
 - the buffer is then bound to the allocated memory
- any memory which is visible to the GPU will be in the list of all memory heaps
- there are usually multiple heaps available
- certain heaps can only be used for specific types of buffers

- Buffers
 - vertex buffer
 - coordinates of all vertices which describes the geometry for the image being rendered
 - triangles are the most widely used shape
 - there can be as few as three vertices, typically thousands for a given image
 - uniform buffer
 - contains data which is applied to every vertex
 - for example, transformations can be used to shift the geometry and render the image somewhere other than the center

- What is a Shader
 - a program which is written in a specialized language
 - shaders are intended to be run on the GPU
 - OpenGL shaders are written in GLSL (OpenGL Shading Language)
 - Direct3D shaders are written in HLSL
 - Vulkan shaders are typically written in GLSL or HLSL and must be compiled to the SPIR-V binary format
 - there are various categories of shaders
 - each is responsible for a different aspect of image rendering
 - some shaders are required others are optional
 - ability to provide your own shader programs is considered one of the largest and most valuable graphics API advancements

- Shader Categories
 - vertex shader
 - transforms a 3D position into 2D coordinates
 - executes once per vertex
 - tessellation shader
 - decompose shapes into smaller components, optional
 - geometry shader
 - alters the vertex shader output, optional
 - fragment shader
 - uses lighting and textures to calculate colors
 - executes at least once per pixel or fragment (partial pixel)

- Graphics Definitions

- pipeline

- before a **draw command** is added to a command buffer you must create a pipeline object which sets a whole lot of options
 - shader handles, descriptor sets, push constants, depth buffer
 - winding direction, culling options, viewport
 - stencil, scissor, blending

- command buffer

- accumulates draw commands which are executed on the GPU at a later time
 - only way to synchronize these commands is by using a barrier
 - set semaphore, query semaphore
 - wait for pipeline stage (such as wait for a vertex shader to finish)

- Graphics Definitions

- frame buffer

- no default buffer exists since displaying an image is optional
 - typically contains one image
 - not required if no images are displayed

- depth buffer

- a 3D mesh has perspective and to draw it realistically some triangles must appear in front of other triangles
 - a depth buffer is used to determine what part of the mesh is closer to the camera
 - misconfiguring this buffer will result in far away objects being rendered on top of closer objects

- Graphics Definitions

- render pass
 - added for improved performance on mobile
 - OpenGL and DirectX 12 do not have a render pass equivalent
 - for tile based GPU support
 - common in mobile, uncommon on desktop computers
 - a pipeline needs information about color channels, the depth buffer, possible image sampling, etc
 - all of this information is split into one or more render passes
 - begin command buffer, begin render pass
 - add commands for this one pass
 - end render pass, end command buffer

- Graphics Definitions

- swapchain

- a set of buffers where the GPU can render
 - in Vulkan there is no default frame buffer
 - creating a swapchain is the only way to create an empty frame buffer
 - each frame buffer will hold a single image which is waiting to be presented on the surface
 - a single swap chain can contain many frame buffers
 - platform specific, supplied through “extensions” which are requested when you create the Vulkan instance

- Queues

- similar to a CPU thread, fixed number of queues on the GPU
- each queue is intended to execute a certain kind of operation
- a mechanism to send commands to the GPU for processing

- graphics queue
 - used to generate an image
- present queue
 - used to copy images to the screen
- computation queue
 - used to run general purpose computations
- transfer queue
 - used to copy data between CPU memory and GPU memory

- Synchronization
 - `draw_frame()` is a function you write which puts a sequence of commands in a queue
 - this function is responsible for the following operations
 - retrieve an inactive frame buffer from the swapchain
 - build a command buffer
 - submit the command buffer to the graphics queue
 - submit the frame buffer to the presentation queue
 - some of these operations are executed **asynchronously**
 - unfortunately some tasks must wait until prior ones have completed
 - additional constraints are required for operations which have dependencies and these must be provided by the programmer

- Graphics Definitions
 - there are two ways to synchronize these operations
 - fence
 - used to synchronize work between the CPU and the GPU
 - status of a fence can be accessed from your program
 - semaphore
 - used to synchronize work occurring on the GPU
 - semaphore dependencies are set up in your program however they are observed only on the GPU
 - not visible to the CPU nor can you query their status

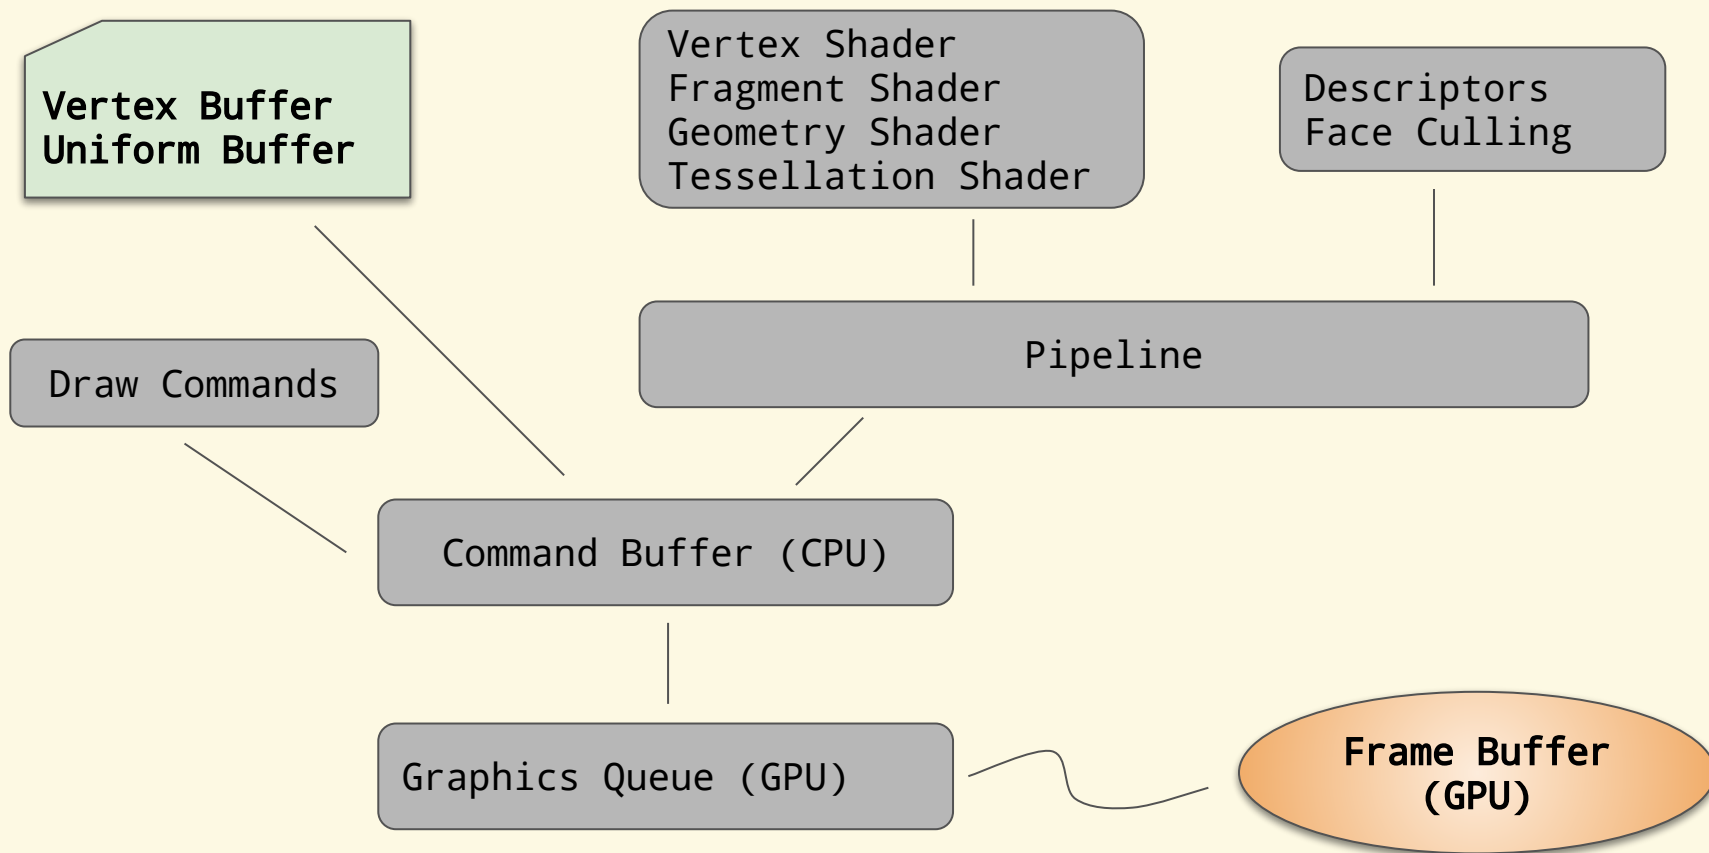
- Graphics Definitions

- winding direction
 - must be specified as part of the pipeline
 - conventionally clockwise winding is used
 - the direction is used to determine whether a given triangle is facing the camera (front facing) or facing away (back facing)
- back face culling
 - is the process of discarding triangles which are back facing since these are not normally visible to the viewer
 - there is no default so it must be set on or off in the pipeline
 - triangles are "culled" early in the rendering pipeline
 - increases efficiency by reducing the number of fragments which are processed

- Graphics Definitions

- push constants
 - a way to provide a small amount of uniform data to shaders
- scissor
 - scissor test restricts drawing to the specified rectangle
 - test is always enabled and can not be turned off
- extent
 - structure containing width and height
 - defines things like the size of the screen or the render area
- viewport
 - specifies a region of the frame buffer which should be refreshed, potentially complicated to determine

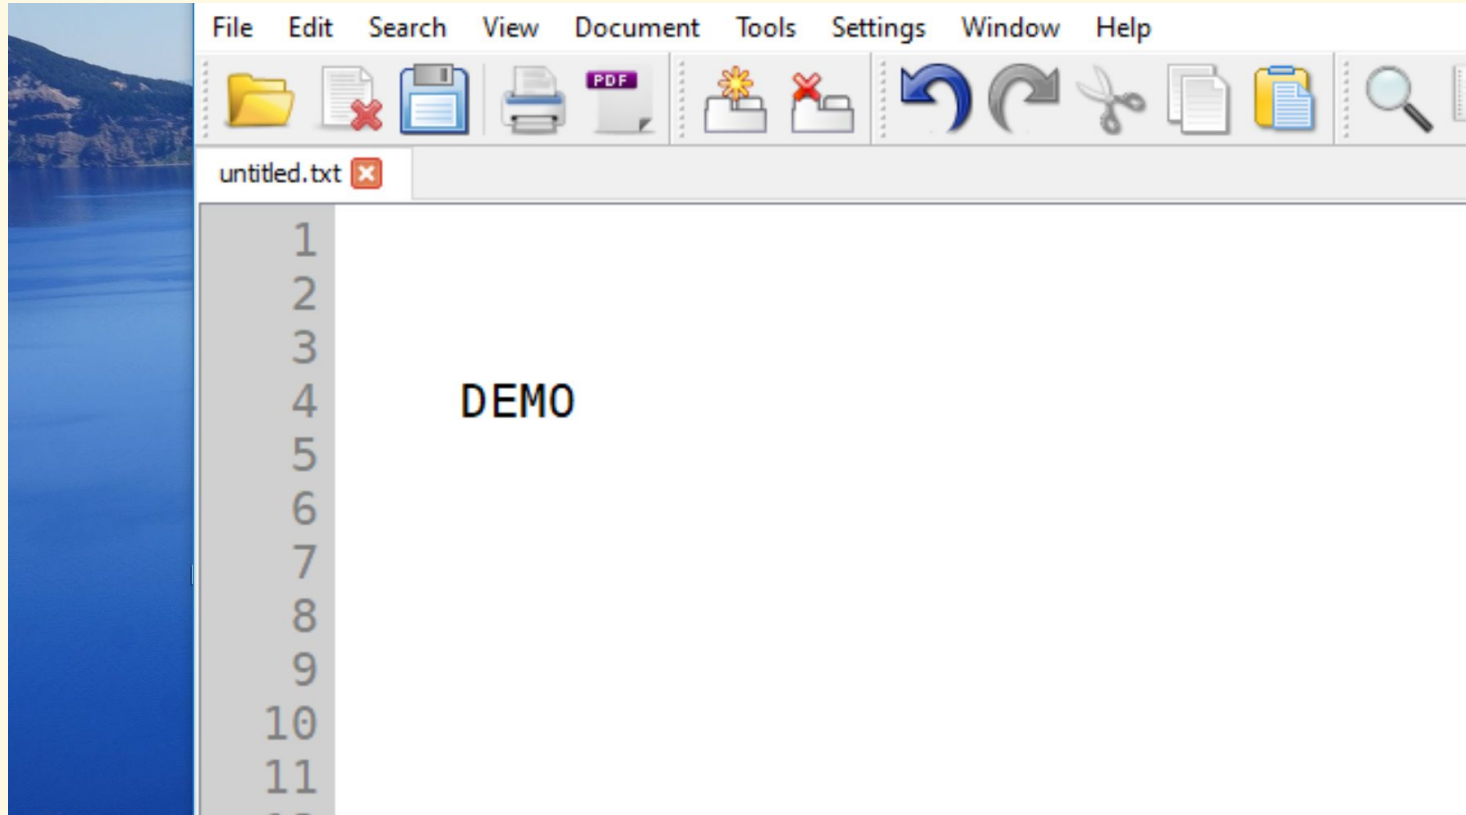
Program / Data Flow



When the frame buffer is ready, a previously configured semaphore is triggered which presents the output

- Putting the Graphics back in GUI
 - intent of a GUI is to present a graphical interface
 - graphics have traditionally been rendered on the CPU
 - most GUI libraries still use software rendering
 - why not use the GPU for all types of graphics
 - text is graphics too
 - CsPaint can render text as graphics on the GPU
 - more efficient
 - infinitely scalable
 - cleaner edges
 - textured, reflective, shadowing . . .

CPU Text vs GPU Rendering



This text is a 16 point monospace font.
Generated screenshot was taken from a Windows GUI application.

CPU Text vs GPU Rendering



Increased the previous image size by 1400 percent.



DEMO

In the “Graphics Demo” moved the camera closer to the text.
Generated a screenshot with no alterations.

- Rendering Text on the GPU
 - walk the string and create a “rectangle” for each letter
 - uses a process similar to applying a standard texture to a model
 - texture for text does not contain a color or pattern
 - contains a distorted font image generated from any normal font file
 - currently supports monospace fonts for Latin-1
 - fragment shader
 - computes a high resolution output using a multi channel (color) signed distance field algorithm
 - all the heavy lifting is part of CsPaint

- Vulkan API

- LunarG is the organization which provides the main Vulkan SDK
- key element is the **loader** which is responsible for the bridge between your program and the vendor supplied graphics drivers
 - supports various layers
 - supports multiple GPUs
- tools for compiling shaders to **SPIR-V** binary output
- other development **tools**
 - <https://www.lunarg.com/vulkan-sdk>
- informative overview of LunarG
 - <https://www.youtube.com/watch?v=wWYRFwIHdJc>

- Vulkan Tutorials
 - <https://vulkan-tutorial.com/>
 - <https://vulkan.lunarg.com/doc/view/1.1.101.0/windows/tutorial/html/index.html>
 - <https://gpuopen.com/understanding-vulkan-objects/>
 - <https://developer.nvidia.com/transitioning-opengl-vulkan>

- Modeling
 - numerous open source and proprietary programs
 - selected Blender
 - open source 3D graphics modeling application
 - professional quality
 - version 2.8 released July 2019
 - several good youtube channels which use Blender
 - blender guru has a 9 part tutorial series
 - <https://www.blenderguru.com/>

Where is CsPaint Headed

- Roadmap

- enhance the graphics demo
- allow some of the defaulted parameters to be configured
- add a slightly higher level API to CsPaint
- provide sample fragment shaders
- enhance 2D graphics support
- create distance field textures on the fly for a given font
- add text shaping using the Harfbuzz version 2 library
- potentially link with CsString for Unicode support
- enhance API documentation
- incorporate user contributions

Integrating CopperSpice with CsPaint

- Roadmap

- main application window will be a Vulkan surface
- internal paint system redesigned to call CsPaint
- all controls rendered using Vulkan

- majority of controls will render on the GPU without modification
 - push button, radio button, check box, combo box
table view, tree view, calendar, line edit, text edit

- on Windows and Unix CopperSpice will use CsPaint directly
- on OS X MoltenVK wrapper will be used to translate Vulkan

- some level of fallback may be maintained for systems which do not have current Vulkan 1.1 support

Presentations

- ❑ Why CopperSpice
- ❑ Why DoxyPress
- ❑ Compile Time Counter
- ❑ Modern C++ Data Types (references)
- ❑ Modern C++ Data Types (value categories)
- ❑ Modern C++ Data Types (move semantics)
- ❑ CsString library (unicode)
- ❑ CsString library (library design)
- ❑ Multithreading in C++
- ❑ Multithreading using libGuarded
- ❑ Signals and Slots
- ❑ Build Systems
- ❑ Templates in the Real World
- ❑ Copyright Copyleft
- ❑ What's in a Container
- ❑ Modern C++ Threads
- ❑ C++ Undefined Behavior
- ❑ Regular Expressions
- ❑ Using DoxyPress
- ❑ Type Traits
- ❑ C++ Tapas (typedef, forward declarations)
- ❑ Lambdas in C++
- ❑ C++ Tapas (typename, virtual, pure virtual)
- ❑ Overload Resolution
- ❑ Futures & Promises
- ❑ Special Member Functions
- ❑ C++ in Review
- ❑ Thread Safety
- ❑ Constexpr Static Const
- ❑ When Your Codebase is Old Enough to Vote
- ❑ Sequencing
- ❑ Linkage
- ❑ Inheritance
- ❑ Evolution of Graphics Technology
- ❑ GPU, Pipeline, and the Vector Graphics API
- ❑ Rendering 3D Graphics
- ❑ Declarations and Type Conversions
- ❑ C++ ISO Standard
- ❑ Inline Namespaces
- ❑ Lambdas in Action
- ❑ Any Optional

Please subscribe to our YouTube Channel
<https://www.youtube.com/copperspice>

Libraries

- CopperSpice
 - libraries for developing GUI applications
- CsPaint Library
 - standalone C++ library for rendering graphics on the GPU
- CsSignal Library
 - standalone thread aware signal/slot library
- CsString Library
 - standalone unicode aware string library
- libGuarded
 - standalone multithreading library for shared data

Applications

- **KitchenSink**
 - contains 30 demos and links with almost every CopperSpice library
- **Diamond**
 - programmers editor which uses the CopperSpice libraries
- **DoxyPress & DoxyPressApp**
 - application for generating source code and API documentation

Where to find CopperSpice

- www.copperspice.com
- ansel@copperspice.com
- barbara@copperspice.com
- source, binaries, documentation files
 - download.copperspice.com
- source code repository
 - github.com/copperspice
- discussion
 - forum.copperspice.com