

The Networking TS in Practice: Patterns for Real World Problems

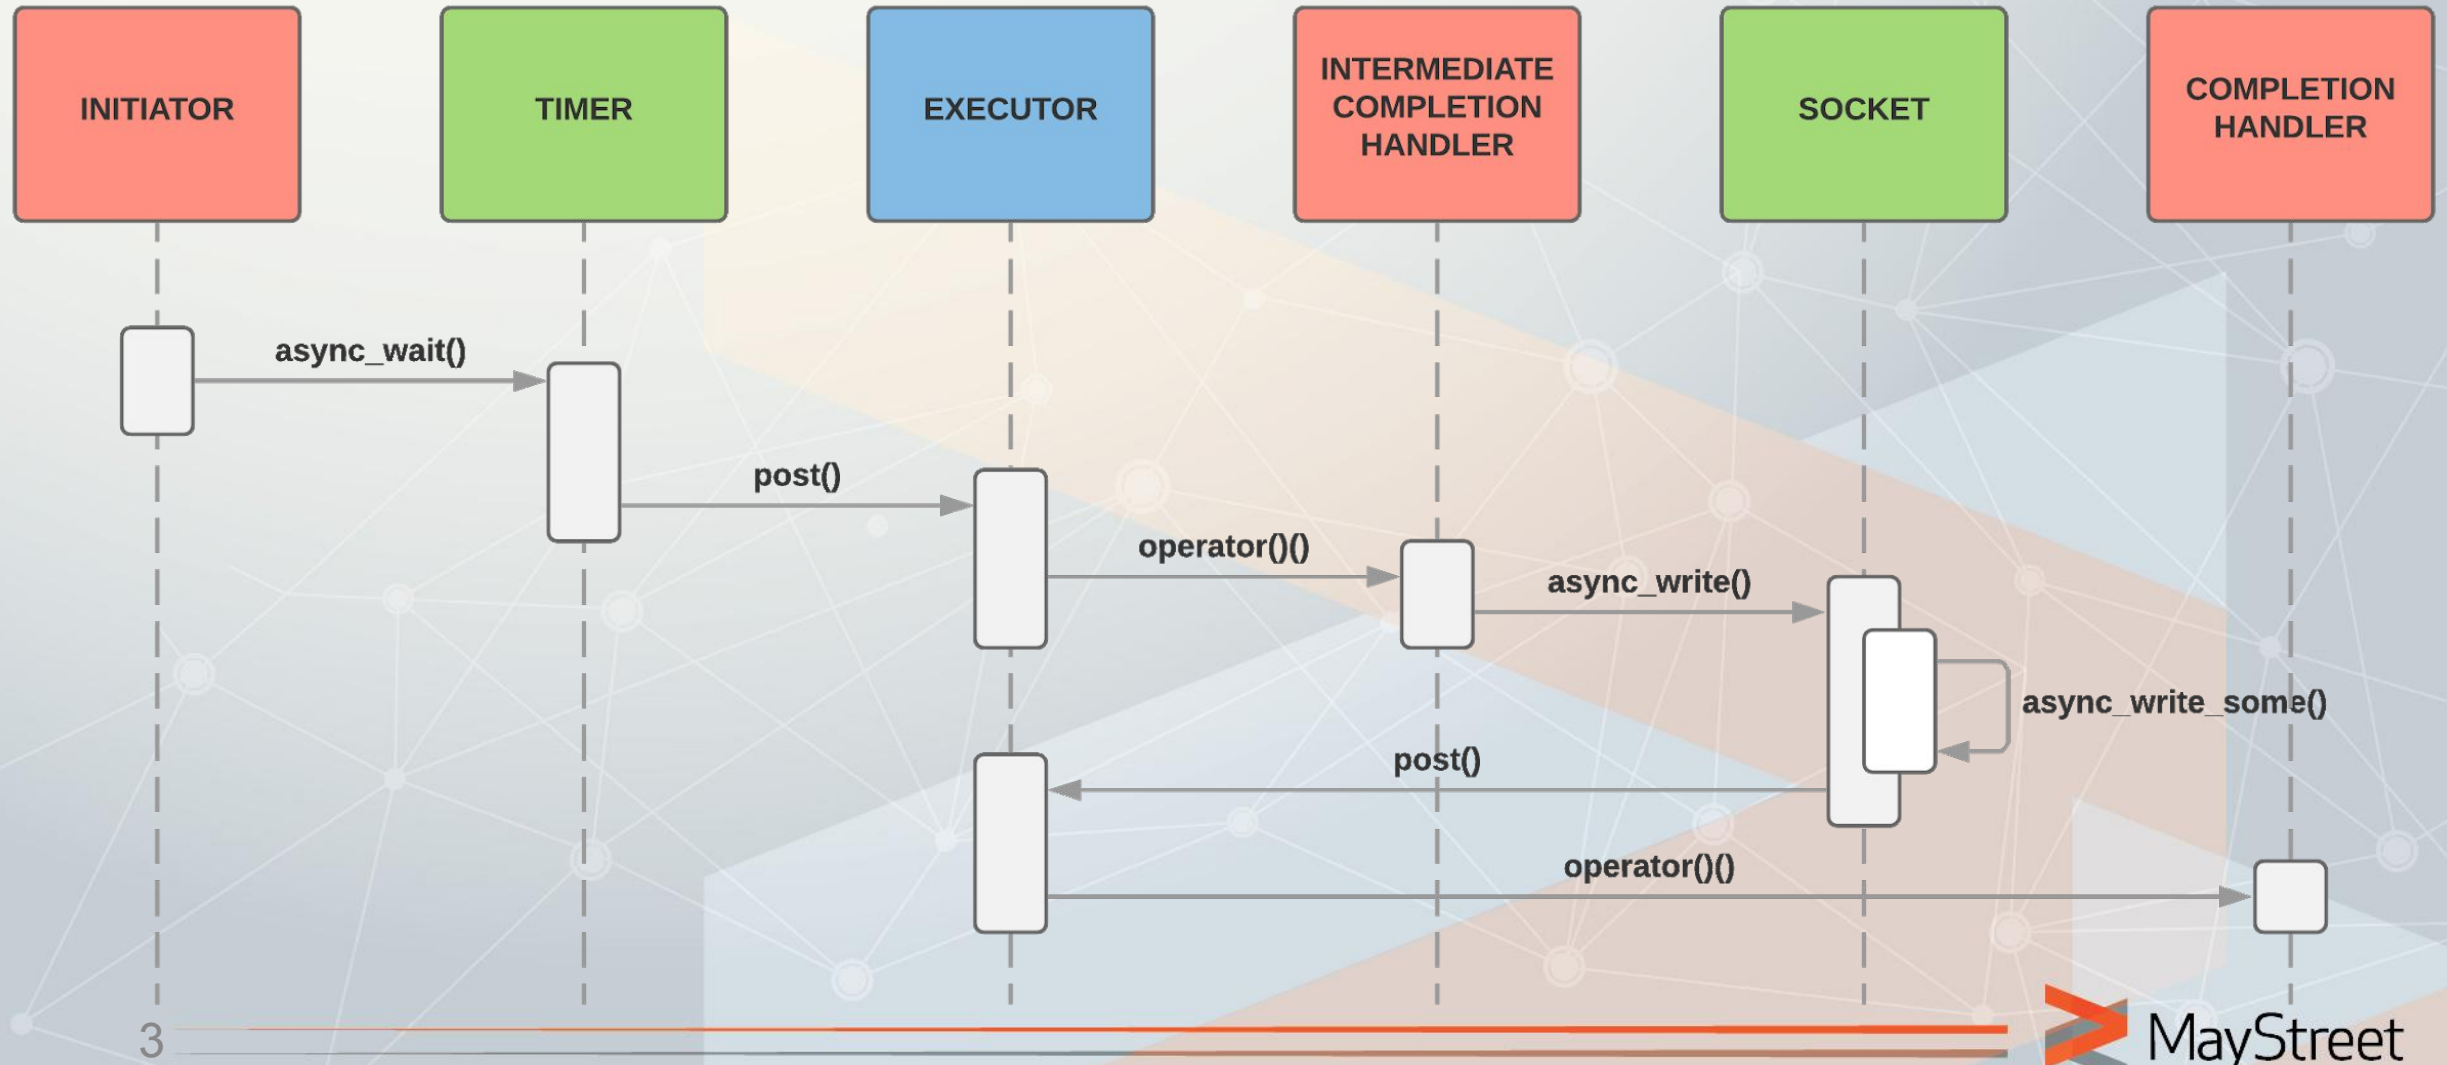
Robert Leahy

Senior Software Engineer
MayStreet Inc.
rleahy@rleahy.ca

Networking TS, Boost.Asio, & ISO C++

- Had hoped for Networking TS in C++20
- Now hoping for C++23
- Boost.Asio interface compatible
 - `boost::asio` namespace
 - Used to prepare these slides
- GCC 9 experimental implementation
 - `std::experimental::net` namespace
 - Incomplete

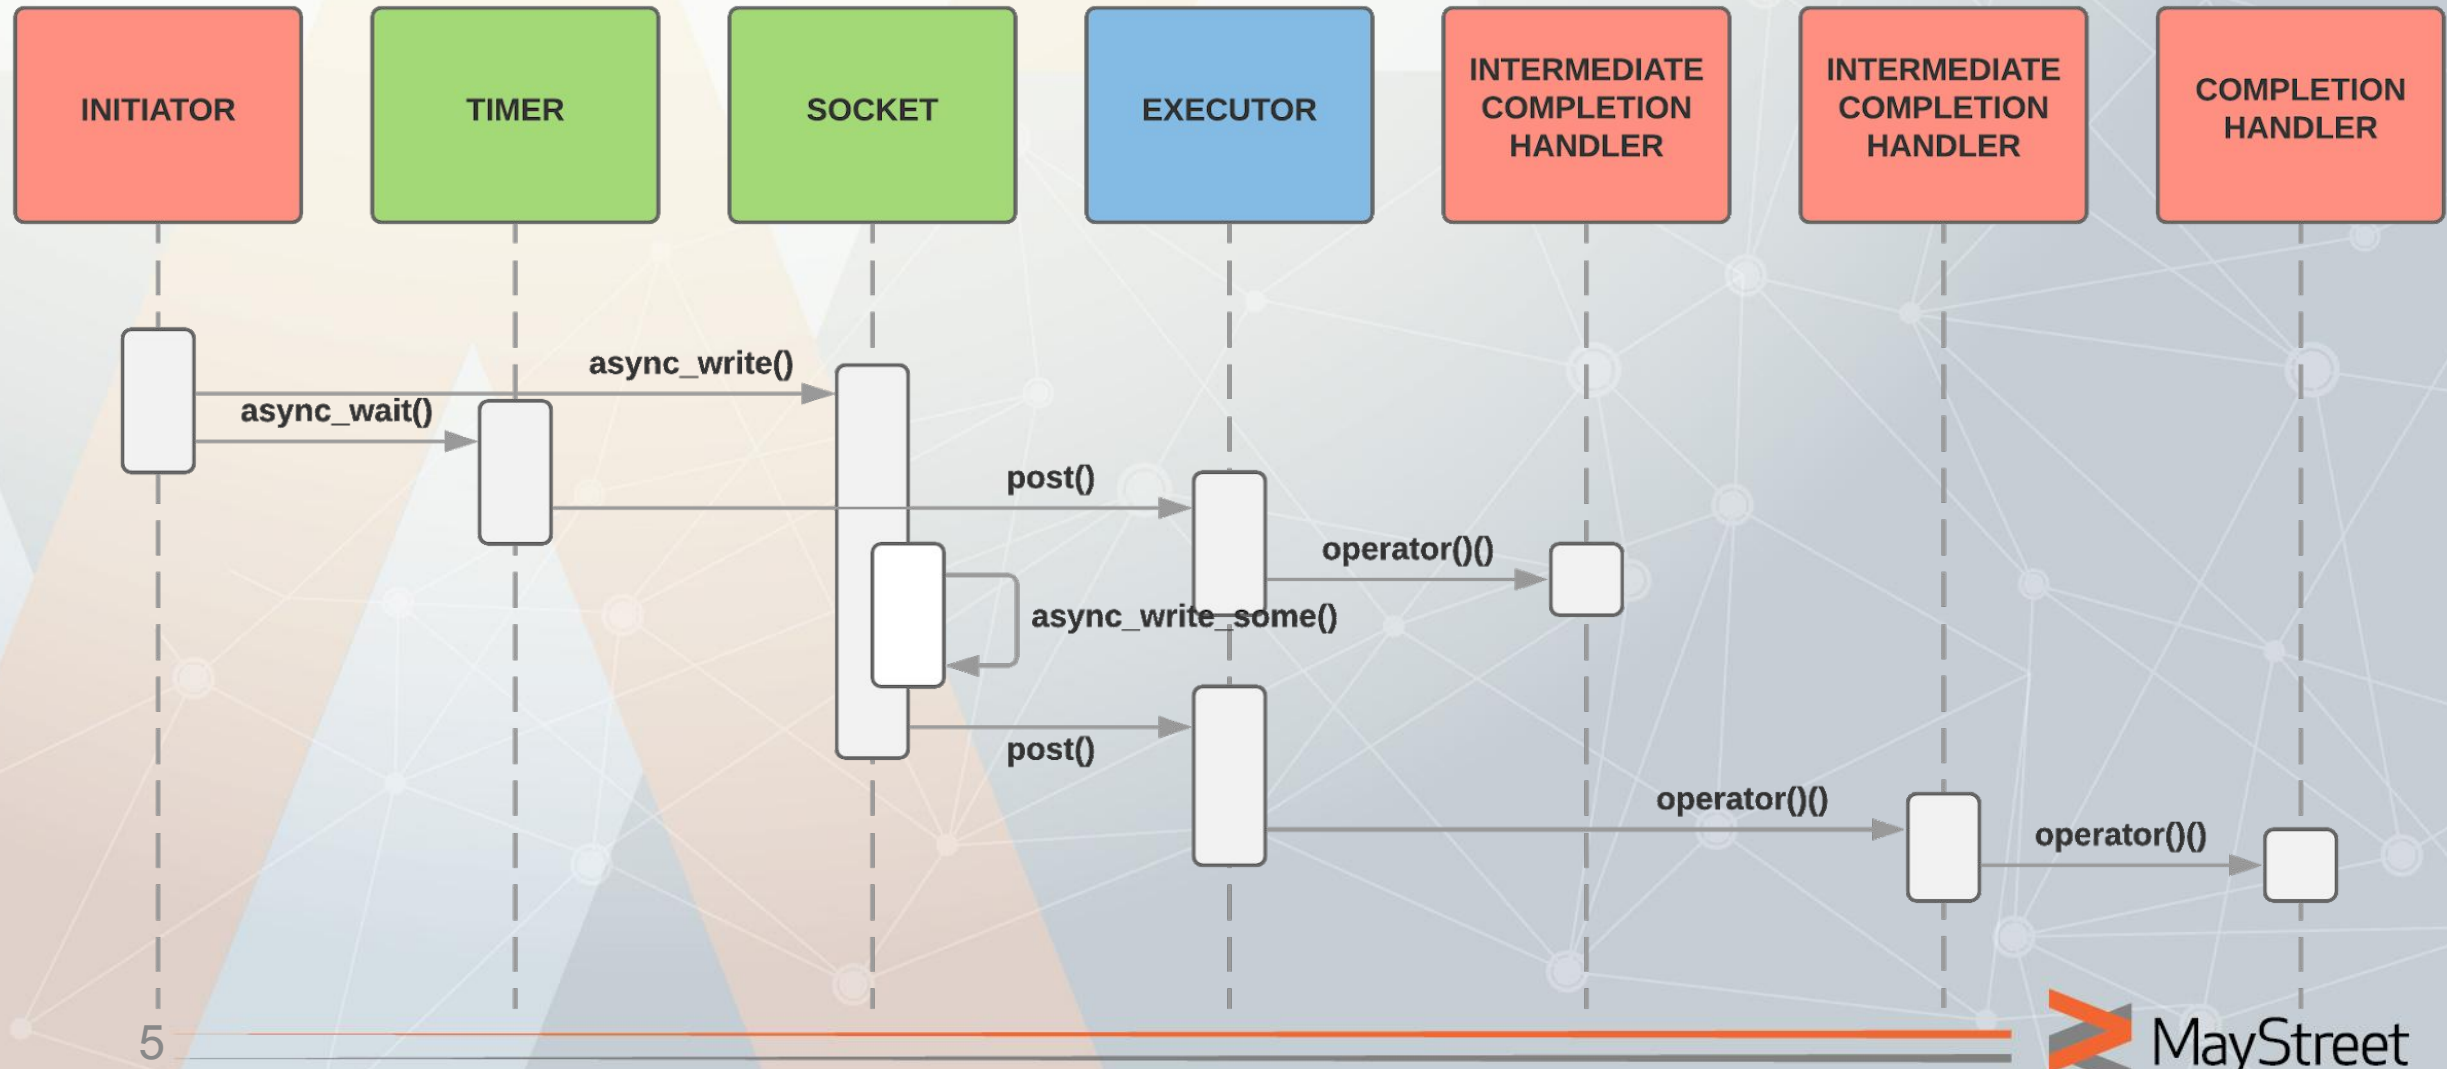
async_wait_then_write



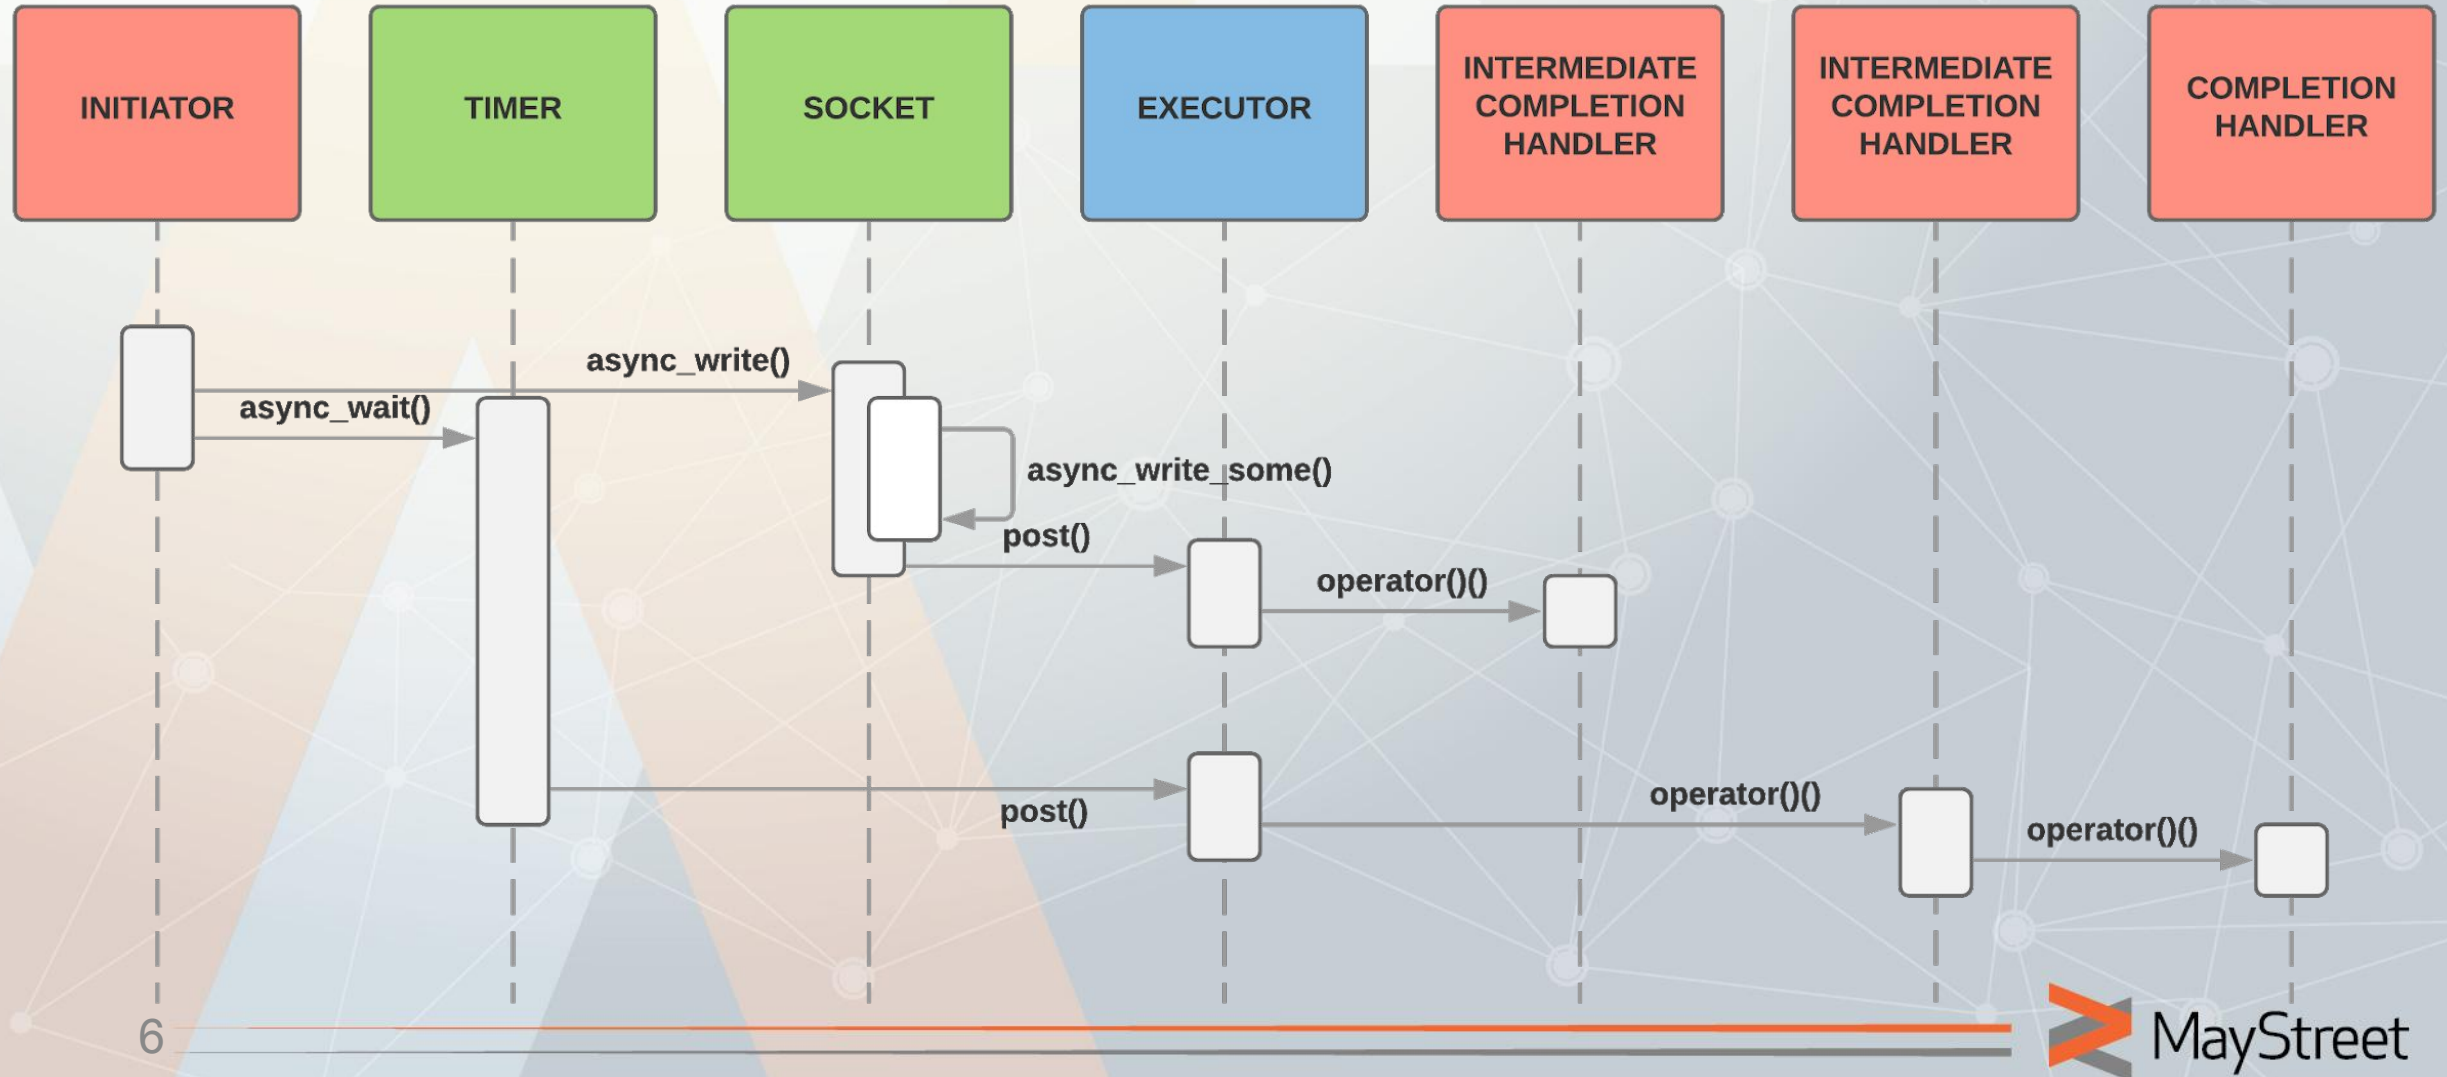
Series Composition

- Each successive operation waits for completion of previous operation
- Pattern can be extended to compose any number of operations **in series**
- Series composition only covers some use cases

async_write_with_timeout



async_write_with_timeout (Alternate)



Shared State

- Intermediate completion handlers must have shared state to:
 - Track completion of intermediate completion handlers
 - Collect results
 - Store final completion handler

```
1 template<class AsyncWriteStream, class ConstBuffers, class Timer, class Token>
2 auto async_write_with_timeout(AsyncWriteStream& stream, ConstBuffers cb, Timer& timer,
3     typename Timer::duration timeout, Token&& token)
4 {
5     using completion_type = /* ... */;
6     using handler_type = /* ... */;
7     struct state : private boost::noncopyable { /* ... */ };
18    class write_op { /* ... */ };
51    class timeout_op { /* ... */ };
80    completion_type completion(token);
81    auto a = std::net::get_associated_allocator(completion.completion_handler);
82    auto ptr = std::allocate_shared<state>(a, stream, timer,
83        std::move(completion.completion_handler));
84    std::net::async_write(ptr->stream, cb, write_op(ptr));
85    timer.expires_after(timeout);
86    timer.async_wait(timeout_op(std::move(ptr)));
87    return completion.result.get();
88 }
```



```
7 struct state : private boost::noncopyable {
8     state(AsyncWriteStream& s, Timer& t, handler_type h) noexcept(/* ... */)
9         : stream(s), timer(t), handler(std::move(h))
10    {}
11    AsyncWriteStream& stream;
12    Timer& timer;
13    handler_type handler;
14    std::error_code ec;
15    std::size_t bytes_transferred = 0;
16    std::size_t outstanding = 2;
17 };
```

```
18  class write_op {
19      std::shared_ptr<state> state_;
20  public:
21      using allocator_type = std::net::associated_allocator_t<handler_type>;
22      auto get_allocator() const noexcept {
23          return std::net::get_associated_allocator(state_>handler);
24      }
25      using executor_type = std::net::associated_executor_t<handler_type,
26          decltype(stream.get_executor())>;
27      auto get_executor() const noexcept {
28          return std::net::get_associated_executor(state_>handler,
29              state_>stream.get_executor());
30      }
31      explicit write_op(std::shared_ptr<state> s) noexcept : state_(std::move(s)) {}
32      write_op(write_op&&) = default;
33      write_op& operator=(write_op&&) = delete;
34      write_op(const write_op&) = delete;
35      write_op& operator=(const write_op&) = delete;
36      // Next slide...
48  };
```

```
18  class write_op {
19      // Previous slide...
36  void operator()(std::error_code ec, std::size_t bytes_transferred) {
37      if (--state_->outstanding) {
38          state_->ec = ec;
39          state_->bytes_transferred = bytes_transferred;
40          state_.reset();
41      } else {
42          auto h = std::move(state_->handler);
43          ec = state_->ec;
44          state_.reset();
45          h(ec, bytes_transferred);
46      }
47  }
48  };
```

```
49  class timeout_op {
50      std::shared_ptr<state> state_;
51  public:
52      using allocator_type = std::net::associated_allocator_t<handler_type>;
53      auto get_allocator() const noexcept {
54          return std::net::get_associated_allocator(state_>handler);
55      }
56      using executor_type = std::net::associated_executor_t<handler_type,
57          decltype(timer.get_executor())>;
58      auto get_executor() const noexcept {
59          return std::net::get_associated_executor(state_>handler,
60              state_>timer.get_executor());
61      }
62      explicit timeout_op(std::shared_ptr<state> s) noexcept : state_(std::move(s)) {}
63      timeout_op(timeout_op&&) = default;
64      timeout_op& operator=(timeout_op&&) = delete;
65      timeout_op(const timeout_op&) = delete;
66      timeout_op& operator=(const timeout_op&) = delete;
67      // Next slide...
79  };
```

```
49  class timeout_op {
50      // Previous slide...
67      void operator()(std::error_code) {
68          if (--state_->outstanding) {
69              state_->ec = make_error_code(std::errc::timed_out);
70              state_.reset();
71          } else {
72              auto h = std::move(state_->handler);
73              auto ec = state_->ec;
74              auto bytes_transferred = state_->bytes_transferred;
75              state_.reset();
76              h(ec, bytes_transferred);
77          }
78      }
79  };
```


Cancellation

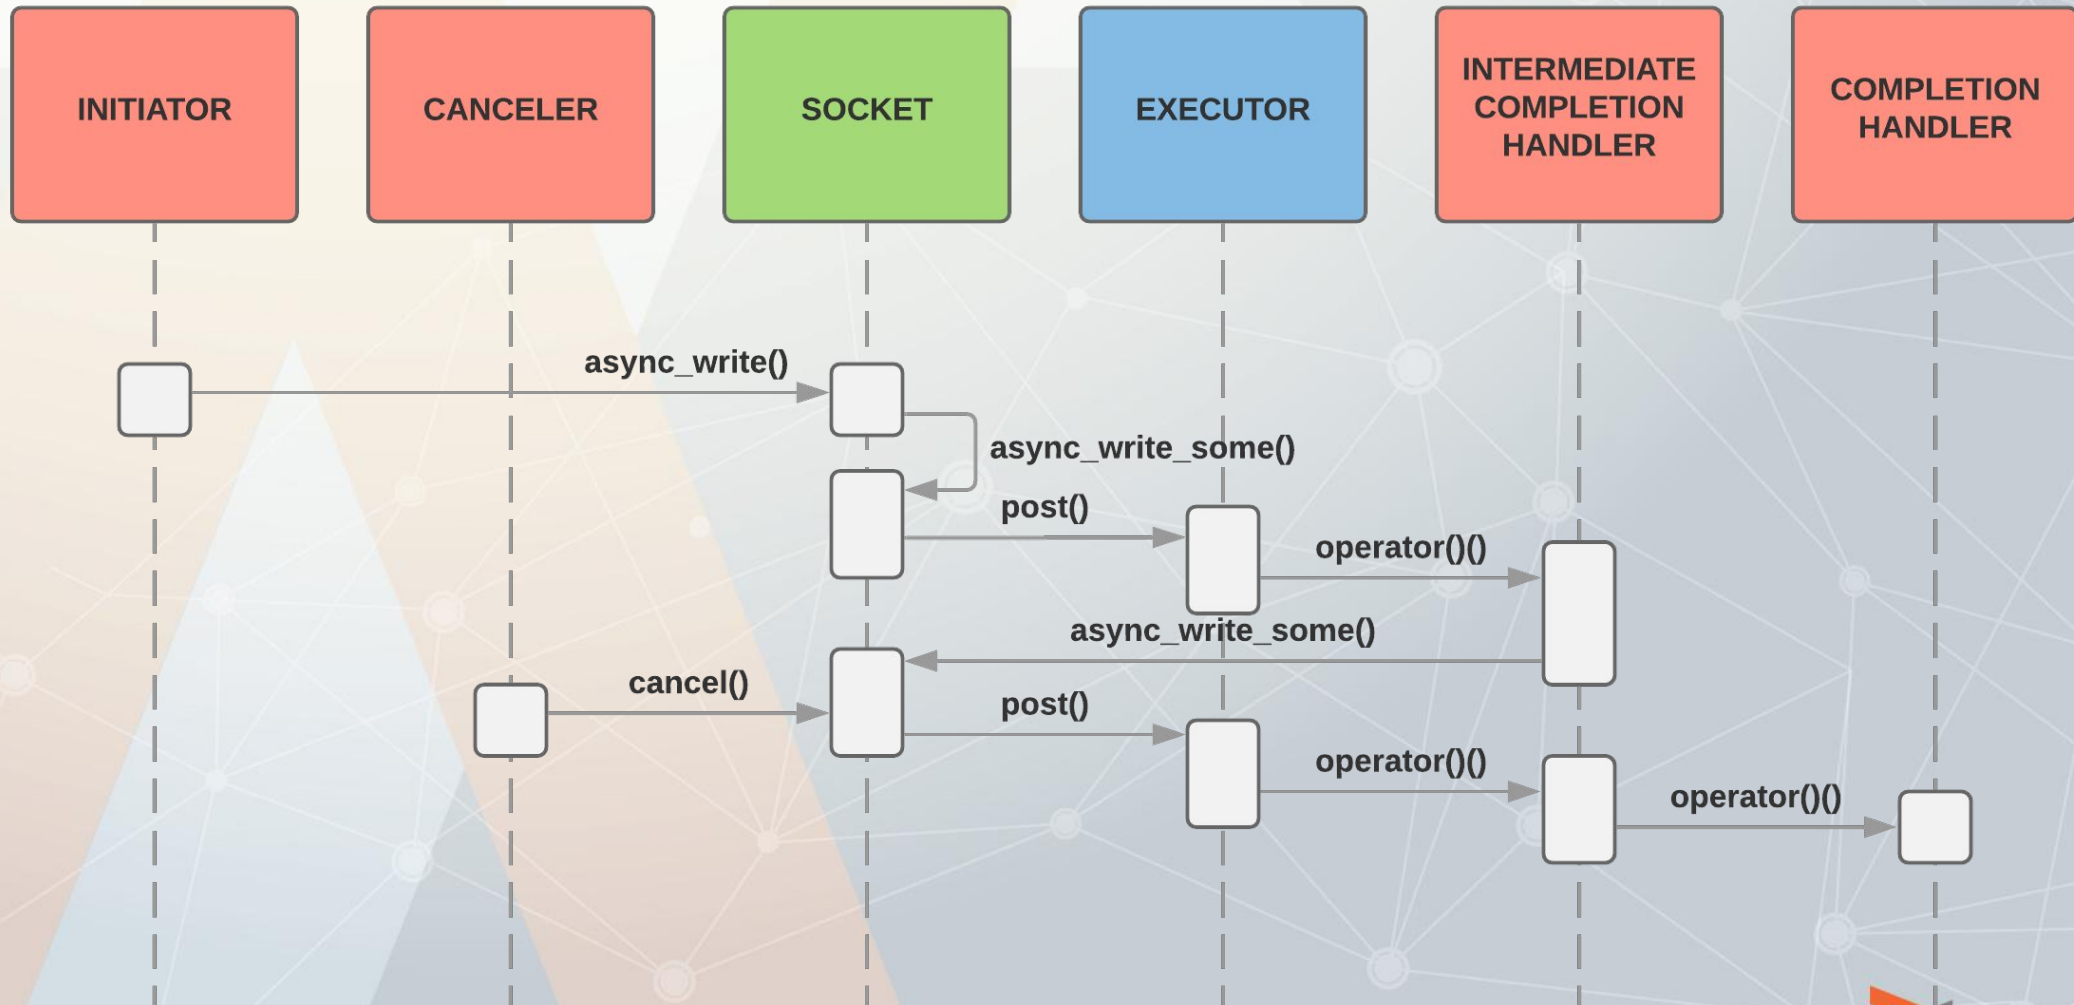
`size_t cancel();`

Effects: Causes any outstanding asynchronous wait operations to complete. Completion handlers for canceled operations are passed an error code `ec` such that `ec == errc::operation_cancelled` yields true.

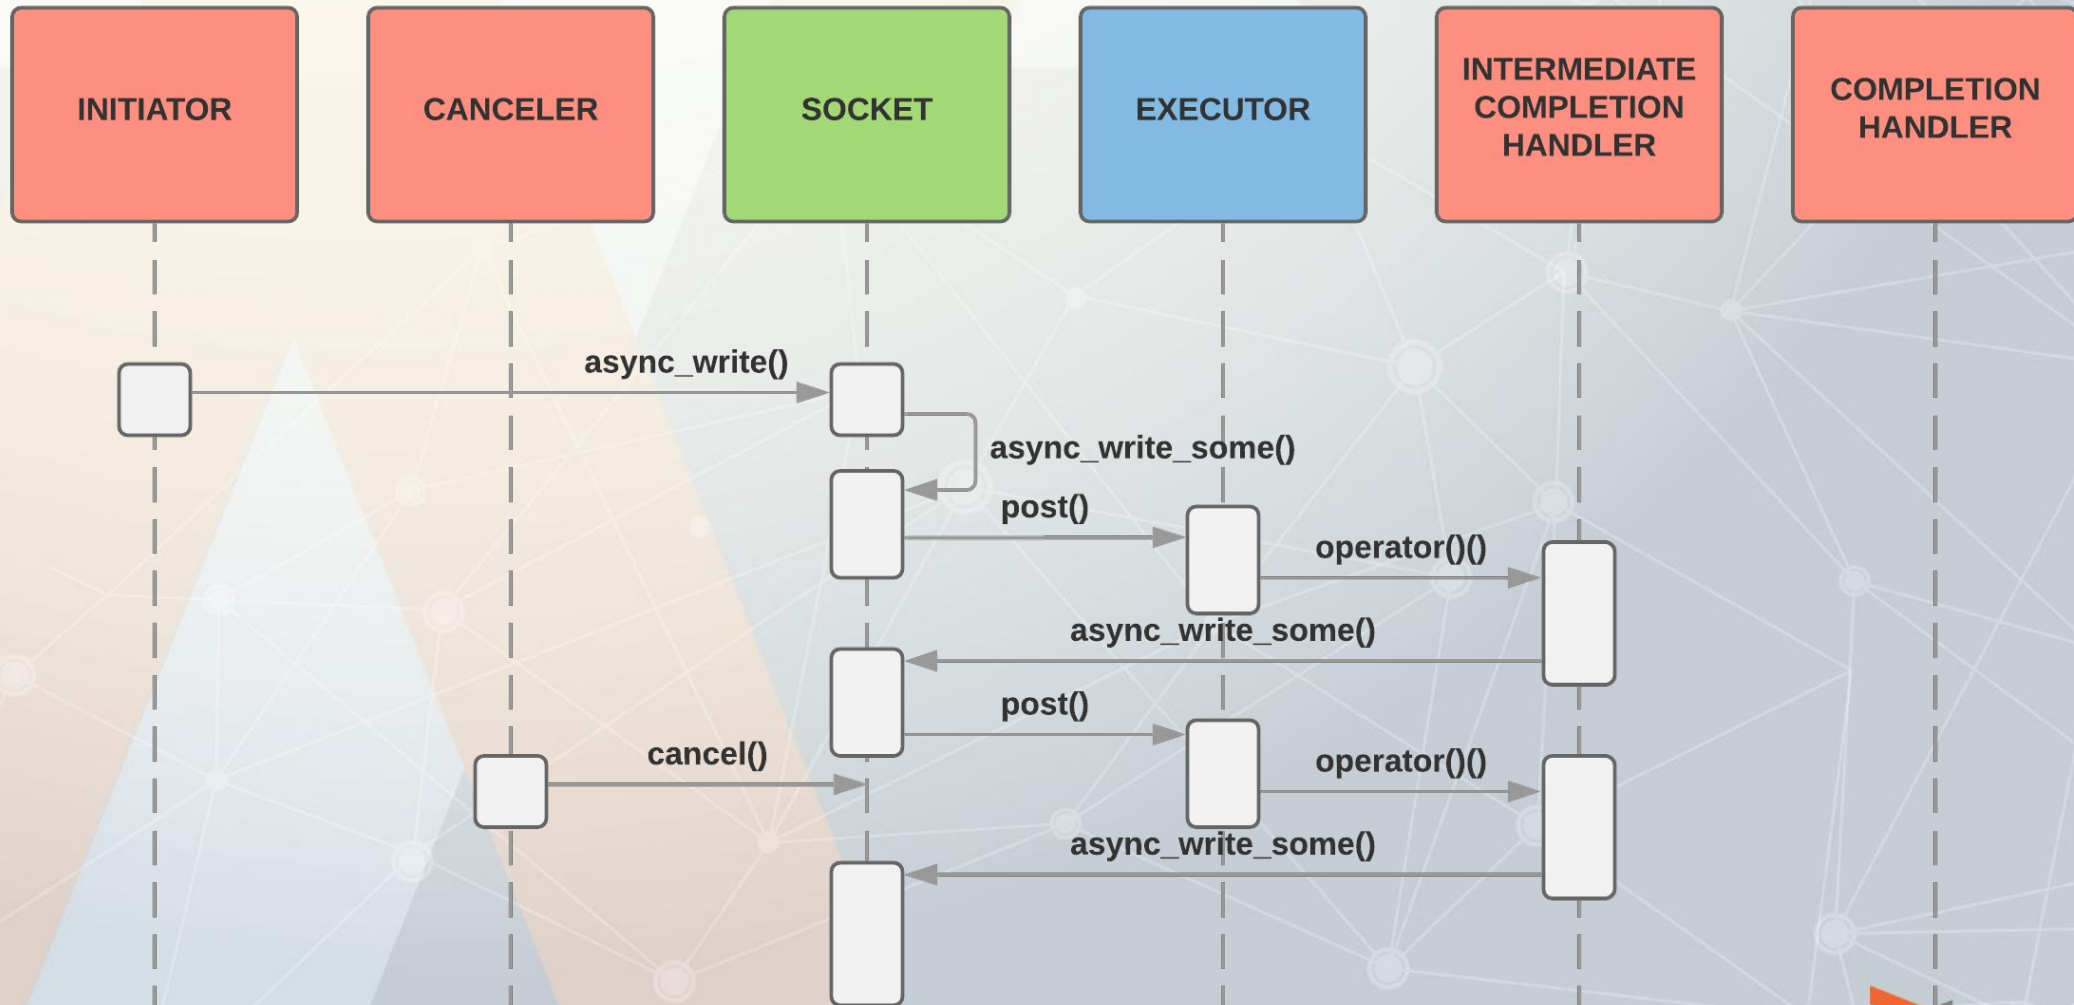
Returns: The number of operations that were canceled.

Remarks: Does not block the calling thread pending completion of the canceled operations.

cancel (Outstanding Operation)



cancel (No Outstanding Operations)



```
1 template<typename AsyncWriteStream>
2 class sticky_cancel_async_write_stream {
3     AsyncWriteStream& stream_;
4     bool canceled_ = false;
5 public:
6     explicit sticky_cancel_async_write_stream(AsyncWriteStream& stream) noexcept
7         : stream_(stream)
8     {}
9     auto get_executor() const noexcept {
10         return stream_.get_executor();
11     }
12     auto cancel() {
13         canceled_ = true;
14         return stream_.cancel();
15     }
16     // Next slide...
31 };
```

```
1 template<typename AsyncWriteStream>
2 class sticky_cancel_async_write_stream {
3     // Previous slide...
16 template<typename ConstBufferSequence, typename Token>
17 auto async_write_some(ConstBufferSequence cb, Token&& token) {
18     if (!canceled_) {
19         return stream_.async_write_some(std::move(cb), std::forward<Token>(token));
20     }
21     /* ... */ completion(token);
22     auto ex = std::net::get_associated_executor(completion.completion_handler,
23         get_executor());
24     auto a = std::net::get_associated_allocator(completion.completion_handler);
25     auto f = [handler = std::move(completion.completion_handler)]() mutable {
26         handler(make_error_code(std::errc::operation_canceled), 0);
27     };
28     ex.post(std::move(f), a);
29     return completion.result.get();
30 }
31 };
```



```
7  struct state : private boost::noncopyable {
8      state(AsyncWriteStream& s, Timer& t, handler_type h) noexcept(/* ... */)
9          : stream(s), timer(t), handler(std::move(h))
10     {}
11     AsyncWriteStream& stream;
11     sticky_cancel_async_write_stream<AsyncWriteStream> stream;
12     Timer& timer;
13     handler_type handler;
14     std::error_code ec;
15     std::size_t bytes_transferred = 0;
16     std::size_t outstanding = 2;
17 };
```

```
18  class write_op {
21      // ...
36      void operator()(std::error_code ec, std::size_t bytes_transferred) {
37          state_>timer.cancel();
38          if (--state_>outstanding) {
39              state_>ec = ec;
40              state_>bytes_transferred = bytes_transferred;
41              state_.reset();
42          } else {
43              auto h = std::move(state_>handler);
44              ec = state_>ec;
45              state_.reset();
46              h(ec, bytes_transferred);
47          }
48      }
49  };
```

```
50 class timeout_op {
51     // ...
68     void operator()(std::error_code) {
69         state_>stream.cancel();
70         if (--state_>outstanding) {
71             state_>ec = make_error_code(std::errc::timed_out);
72             state_.reset();
73         } else {
74             auto h = std::move(state_>handler);
75             auto ec = state_>ec;
76             auto bytes_transferred = state_>bytes_transferred;
77             state_.reset();
78             h(ec, bytes_transferred);
79         }
80     }
81 };
```

Parallelism

- Natural happens-before relationship when composing in series
- Not so with parallel composition
- `std::net::strand`
 - Has runtime cost
 - Guarantee may already be met
 - Initiations and completions can still race
- Put onus on user to maintain strand
 - Boost.Beast does this with `boost::beast::basic_stream`

Managing Parallelism

- Single-threaded I/O objects and assumption of strand make code
 - Easy to write
 - Easy to reason about
 - Efficient by default
- CPUs have many cores so applications ought to leverage parallelism
- Single `std::net::io_context` with many threads using explicit strands
 - Error prone
 - Adds overhead
- Many `std::net::io_context` each running in a single thread
 - Mostly correct by default
 - No overhead
 - Need to distribute work

Round Robin

- Consider TCP server
 - Continually accepts incoming connections
 - Easy to parallelize management of different connections
- Using many `std::net::io_context` objects leverages parallelism implicit in problem
- `std::net::basic_socket_acceptor` allows accepted socket to be associated with arbitrary `std::net::io_context`

```
template<class CompletionToken>  
DEDUCED async_accept(io_context&, CompletionToken&&);
```

```
1 template<class Acceptor, class Iterator, class CompletionHandler>
2 struct async_accept_op {
3     Acceptor& acc_;
4     Iterator begin_, curr_, end_;
5     CompletionHandler h_;
6     using allocator_type = /* ... */;
7     auto get_allocator() const noexcept { /* ... */ }
8     using executor_type = /* ... */;
9     auto get_executor() const noexcept { /* ... */ }
10    template<typename Stream>
11    void operator()(std::error_code ec, Stream s) {
12        ++curr_;
13        if (curr_ == end_) {
14            curr_ = begin_;
15        }
16        h_(ec, std::move(s), curr_);
17    }
18 };
```

```
20 template<typename Acceptor, typename Iterator, typename Token>
21 auto async_accept(Acceptor& acc, Iterator begin, Iterator curr, Iterator end,
22   Token&& token)
23 {
24   using completion_type = /* ... */;
25   completion_type completion(token);
26   using op_type = async_accept_op<Acceptor, Iterator,
27     typename completion_type::completion_handler_type>;
28   op_type op{acc, begin, curr, end, std::move(completion.completion_handler)};
29   acc.async_accept(curr->context(), std::move(op));
30   return completion.result.get();
31 }
```

Operations that Never Succeed

- In networking processing often proceeds until exceptional condition
 - Connection termination
 - Server/client shutdown
- Until now all asynchronous operations complete after performing well-defined amount of work and report success or failure
- Asynchronous operation which never succeeds and proceeds until exceptional condition better models real world
- Repeatedly accept incoming connections and round robin `std::net::io_context` objects until exceptional condition

```
1 template<typename Acceptor, typename Iterator, typename AfterAccept, typename Handler>
2 struct async_accept_loop_op {
3     Acceptor& acc_;
4     Iterator begin_, end_;
5     AfterAccept after_;
6     Handler h_;
7     using allocator_type = /* ... */;
8     auto get_allocator() const noexcept { /* ... */ }
9     using executor_type = /* ... */;
10    auto get_executor() const noexcept { /* ... */ }
11    void initiate(Iterator curr) {
12        auto&& acc = acc_;
13        auto begin = begin_;
14        auto end = end_;
15        ::async_accept(acc, begin, curr, end, std::move(*this));
16    }
17    // Next slide...
26 };
```



```
1 template<typename Acceptor, typename Iterator, typename AfterAccept, typename Handler>
2 struct async_accept_loop_op {
3     // Previous slide...
17    template<typename Stream>
18    void operator()(std::error_code ec, Stream s, Iterator curr) {
19        if (ec) {
20            h_(ec);
21            return;
22        }
23        after_(std::move(s));
24        initiate(curr);
25    }
26 };
```

```
28 template<typename Acceptor, typename Iterator, typename AfterAccept, typename Token>
29 auto async_accept(Acceptor& acc, Iterator begin, Iterator end,
30     AfterAccept after_accept, Token&& token)
31 {
32     using completion_type = /* ... */;
33     completion_type completion(token);
34     using op_type = async_accept_loop_op<Acceptor, Iterator, AfterAccept,
35         typename completion_type::completion_handler_type>;
36     op_type op{acc, begin, end, std::move(after_accept),
37         std::move(completion.completion_handler)};
38     op.initiate(begin);
39     return completion.result.get();
40 }
```

Summary

- Real world problems seem to resist the framework and patterns of the Networking TS but with care they can be reconciled
- Composition of operations in parallel requires special attention to Executor choice to avoid data races
- Multiple `std::net::io_context` objects make parallelism more manageable
- Not all operations need to have success completion condition

Questions



Full examples (including tests)
<https://github.com/RobertLeahy/CppCon2019Talk>